

CSAI 302

Advanced Database Systems

Lec 01

**Advanced DBMS and Database
System Architectures**

Outline

- ◆ Database system essentials
- ◆ DBMS Architecture Fundamentals
- ◆ Evolution of Database System Architectures
- ◆ Data Models
- ◆ OLTP vs. OLAP Workload Characteristics
- ◆ Course Details

Database system essentials

Database System Architecture

- ◆ The overall design and structure of a database system that
 - ★ outlines ***how data is stored***, managed, and accessed.
 - ★ defines ***the components*** of the system
 - how they ***interact***, and how ***data flows*** between them.
- ◆ A well-defined architecture ensures Data:
 - ★ Integrity
 - ★ Efficiency
 - ★ Security
 - ★ Scalability.

Database Architecture Types


Tiered Architectures

1-Tier

2-Tier

3-Tier

Database Architecture Types



Schema

- ◆ The blueprint for how data is organized in a database
- ◆ defining elements like
 - ★ tables, fields, data types, and their relationships
- ◆ but not the actual data itself.
- ◆ ***This defines the structure of data for a data model.***

Three-Schema Architecture

External Schema
View Level

Customized data views



patient medical records and treatment history



patient demographic information and insurance details.

Conceptual Schema
Logical Level


Logical structure of the database



- Tables
- Relationships
- Integrity constraints

Internal Schema
Physical Level


Physical storage structures



- B-tree indexes for frequent lookups
- Hash partitioning for parallel processing
- Columnar storage for analytical queries.

DBMS

- ◆ A software that allows **applications** to **store** and **analyze** information in a database.
- ◆ A general-purpose DBMS is designed to allow the
 - ★ **definition, creation, querying, update, and administration**
 - of databases
 - ★ **in accordance with**
 - some data model.



DBMS Core Components

DBMS Key Components

Database Engine

- Handles **data processing**
- Enables efficient access and manipulation.

Database Schema

- The ***blueprint*** or logical structure of the database.

Query Processor

- Interprets and executes **user queries**.
- Translates them into actions for the database.

Storage Manager


- Manages the **physical storage** of data on various media
- Handles **data access**.

Transaction Manager

- Ensures the **integrity and consistency** of data by
 - Managing **transactions**
 - Enforcing **rules** like ACID properties.


DBMS Key Components

DBMS Key Components



Query Processing Engine

- ◆ Transforms high-level SQL statements into efficient execution plans through several critical phases.



Evolution of Database System Architectures

Historical Progression

- ◆ The evolution of database architectures reflects the changing demands of data processing, from simple record-keeping to complex analytical workloads handling massive datasets.
- ◆ Architectural evolution Drivers
 - ★ Scalability Requirements
 - ★ Data Complexity
 - ★ Performance Expectations
 - ★ Reliability Demands

File-Based Systems (1960s-1970s)

- ◆ Early computing relied on flat file systems
- ◆ Each application maintained its own data files.
- ◆ Integrated Data Store
 - ★ Network data model.
 - ★ Tuple-at-a-time queries.
 - ★ Transaction-oriented
 - ★ Dedicated working storage area for each record type
- ◆ This approach suffered from data redundancy, inconsistency, and lack of standardized access methods.

Hierarchical and Network Models (70s-80s)

- ◆ IBM's Information Management System (IMS) introduced hierarchical structures resembling tree organizations,
- ◆ The Conference on Data Systems Languages (CODASYL) developed network models allowing more complex relationships.
- ◆ These systems provided better data organization but required programmers to navigate complex pointer structures manually.
 - ★ Hierarchical data model.
 - ★ Programmer-defined physical storage format.
 - ★ Tuple-at-a-time queries.

Relational Model (80s-90s)

- ◆ Relational model introduced mathematical foundations based on set theory and first-order predicate logic.
- ◆ Allowed users to focus on what data they wanted rather than how to retrieve it
- ◆ leading to the development of SQL and modern RDBMS systems like Oracle, DB2, and SQL Server.
- ◆ Database abstraction to avoid this maintenance:
 - ★ Store database in simple data structures.
 - ★ Access data through high-level language.
 - ★ Physical storage left up to implementation.



ORACLE

Informix

TANDEM

SYBASE

TERADATA

INGRES

InterBase

Object-Oriented-Systems (2000s)

- ◆ As applications became more complex, the need to store complex data types led to object-oriented databases and object-relational extensions.
- ◆ PostgreSQL exemplifies this evolution, supporting both traditional relational operations and complex data types like arrays, JSON, and user-defined types.



PostgreSQL



2000s – Internet boom

- ◆ All the big players were heavyweight and expensive.
- ◆ Open-source databases were missing important features.
- ◆ Many companies wrote their own custom middleware to scale out database across single-node DBMS instances.

2000s – Data warehouses

- ◆ Rise of the special purpose **OLAP**



DBMSs.

- ★ Distributed / Shared-Nothing
- ★ Relational / SQL
- ★ Usually closed-source.
- ◆ Significant performance benefits from using **columnar data storage model.**



2000s – NoSQL Systems

- ◆ Focus on high-availability & high-scalability:
- ★ Schema-less (i.e., “Schema Last”)
- ★ Non-relational data models (document, key/value, etc)
- ★ No ACID transactions
- ★ Custom APIs instead of SQL
- ★ Usually open-source



2010s – NewSQL

◆ Provide same performance for **OLTP workloads** as **NoSQL DBMSs** without giving up **ACID**:

- ★ Relational / SQL
- ★ Distributed
- ★ Usually closed-source



db Shards



Google
Spanner



Comdb2



2010s – Hybrid systems

- ◆ **H**ybrid **T**ransactional-**A**nalytical **P**rocessing.
- ◆ Execute fast **OLTP** like a **NewSQL** system while also executing complex **OLAP** queries like a **data warehouse** system.
 - ★ Distributed / Shared-Nothing
 - ★ Relational / SQL
 - ★ Mixed open/closed-source.



JustOne_{DB}



HyPer
splice
MACHINE



SNAPPY
Peloton

2010s – Cloud systems

- ◆ First database-as-a-service (**DBaaS**) offerings were "containerized" versions of existing DBMSs.
- ◆ There are new DBMSs that are designed from scratch explicitly for running in a cloud environment.



SlicingDice



FAUNA



amazon
DynamoDB



Amazon
Aurora



Google
Spanner



2010s – Shared-disk engines

- ◆ Instead of writing a custom storage manager, the DBMS relies on third-party distributed storage (object stores).
 - ★ Scale execution layer independently of storage.
 - ★ Favors log-structured approaches.
- ◆ This is what most people think of when they talk about a **data lake**.



presto



databricks



2010s – Stream processing

- ◆ Execute continuous queries on streams of tuples.
- ◆ Extend processing semantics to include notion of windows.
- ◆ Often used in combination of batch-oriented systems in a ***lambda architecture*** deployment.



2010s – Graph systems

- ◆ Systems for storing and querying graph data.
 - ◆ Their main advantage over other data models is to provide a graph-centric query API
- ★ Recent research demonstrated that is unclear whether there is any benefit to using a graph-centric execution engine and storage manager.



graphbase.ai



TerminusDB



TigerGraph




Dgraph



JanusGraph



IndraDB



APACHE
GIRAPH

2010s – Timeseries systems

- ◆ Specialized systems that are designed to store timeseries / event data.
- ◆ The design of these systems make deep assumptions about the distribution of data and workload query patterns.



TIMESCALE



2020s-LAKEHOUSE SYSTEMS

- ◆ Middleware for data lakes that adds support for better schema control / versioning with transactional CRUD operations.
 - ★ Store changes in row-oriented log-structured files with indexes.
 - ★ Periodically compact recently added data into read-only columnar files.



Data Models

*A collection of concepts for
describing the data in database.*

Common Data Models

Relational

- (Most DBMSs)

Key/Value

- (NoSQL)

Graph

- (NoSQL)

Document / XML / Object

- (NoSQL)

Wide-Column / Column-family

- (NoSQL)

Array / Matrix / Vector

- (Machine Learning)

Hierarchical / Network / Multi- Value

- (Obsolete/ Rare)

Relational Model

Relational Model

- ◆ The relational model defines a database abstraction based on relations to avoid maintenance overhead.
- ◆ It has three key ideas:
 - ★ Store database in simple data structures (relations)
 - ★ Physical storage left up to the DBMS implementation
 - ★ Access data through a high-level (declarative) language, where the DBMS figures out best execution strategy

Relational Model concepts

Structure

The definition of relations and their contents independent of their physical representation.

Each relation has a set of attributes. Each attribute has a domain of values.

Integrity

Ensure the database's contents satisfy certain constraints.

An example of a constraint would be that the age of a person cannot be a negative number.

Manipulation

An interface for accessing and modifying a database's contents (e.g. SQL, Dataframes)

Relational Model definitions

- ◆ Relation
- ◆ Tuple
- ◆ Primary key
- ◆ Foreign key
- ◆ Constraint

Document Model

Document Data Model

- ◆ A collection of record documents containing a hierarchy of named field/value pairs.
 - ★ A field's value can be either a scalar type, an array of values, or another document.
 - ★ Modern implementations use JSON.
 - Older systems use XML or custom object representations.
- ◆ Avoid “relational-object impedance mismatch” by tightly coupling objects and database.




Relational vs. Document

SUPPLIER
(sno, sname, scity)

SUPPLY
(sno, pno, qty, price)

PART
(pno, pname)



SUPPLIER
(sno, sname, scity, sstate, parts[])



PART
(pno, pname, qty, price)

JSON file

Î

A A **JAÆEËDÇCÊÉJGAİ**
A A A A İ
A A A A A **AJEEÉJGAEFÂ**
A A A A A **AJEEÄÐCJGAJACCDJG**
A A A A A **AJECÇEÉJGAJÄEEÉÈÄJG**
A A A A A **AJEAÉEÉJGAİ**
A A A A A A A A İ
A A A A A A A A **AJEEÉJGAESG**
A A A A A A A A **AJEEÄÐCJGAJCÇÄÇEJG**
A A A A A A A A **AJEEÉJGASEG**
A A A A A A A A **AJEEÇCÇJUEEE**
A A A A A A A İG
A A A A A A A İ
A A A A A A A **AJEEÉJGAEZEG**
A A A A A A A **AJEEÄÐCJGAJEBDÇJG**
A A A A A A A **AJEEÉJGAEZEG**
A A A A A A A **AJEEÇCÇJUEEE**
A A A A A A A İ
A A A A A İ
A A A A İG

A A A A A İ
A A A A A **AJEEÉJGAEZEG**
A A A A A **AJEEÄÐCJGAJADÉJG**
A A A A A **AJECÇEÉJGAJ3ECAÄCEGJG**
A A A A A **AJEAÉEÉJGAİ**
A A A A A A A A İ
A A A A A A A A **AJEEÉJGAEZEG**
A A A A A A A A **AJEEÄÐCJGAJEDDÇEJG**
A A A A A A A A **AJEEÉJGAEG**
A A A A A A A A **AJEEÇCÇJUEEE**
A A A A A A A İG
A A A A A A A İ
A A A A A A A **AJEEÉJGAEZEG**
A A A A A A A **AJEEÄÐCJGAJEECDÄEECEJG**
A A A A A A A **AJEEÉJGAEZEG**
A A A A A A A **AJEEÇCÇJUEEE**
A A A A A A A İ
A A A A İ
A A İ
İ

JSON table

Suppliers	sno	17		53		
	sname	Adel			Samy	
	scity	Mansoura			6th Oct.	
	parts	pno	124	135	pno	154
		pname	chair	table	pname	wallet
	qty	45	32	qty	5	12
	price	1000	3000	price	2000	1500

XML file

ŞHËĐĐAĚČĒĚҪĒEŞJĘGĘJАČECÉĆҪEҪŞJĄĀĀIĞJHŞ
ŞĒEÉÊŞ

A AŞAĚEĐҪCĒEŞ

A A A AŞEEÉŞEFAǎāEEEŞ

A A A AŞEEÁĐCSÁCĆDǎāEEÁĐCS

A A A AŞECҪEŞAEEEEĚAǎāEĆҪEĚS

A A A AŞEÁEEĚS

A A A A A AŞEEÉŞEĘS aǎāEEEŞ

A A A A A AŞEEÁĐCSCCÁҪE aǎāEEÁĐCS

A A A A A AŞEĘEŞSĘ aǎāEEEŞ

A A A A A AŞEEҪCŞEĘEĚ aǎāEEҪCŞ

A A A AǎāEEEĚS

A A A A AŞEÁEEĚS

A A A A A AŞEEÉŞEĘ aǎāEEEŞ

A A A A A AŞEEÁĐCSĘBĐC aǎāEEÁĐCS

A A A A A AŞEĘEŞSĘ aǎāEEEŞ

A A A A A AŞEEҪCŞSĘEĘ aǎāEEҪCŞ

A A A AǎāEEEĚS

A AǎāaEEEĐҪCĒEŞ

A A AŞaĚEĐҪCĒEŞ

A A A AŞEEÉŞE aǎāEEEŞ

A A A AŞEEÁĐCS aÁĐE aǎāEEÁĐCS

A A A AŞECҪEŞ 3ECAACE G aǎāEĆҪE EŞ

A A A AŞEÁEEĚS

A A A A A AŞEEÉŞE EĘS aǎāEEEŞ

A A A A A AŞEEÁĐCS E ÁĐDČE aǎāEEÁĐCS

A A A A A AŞEĘEŞE aǎāEEEŞ

A A A A A AŞEEҪCŞ E EEE aǎāEEҪCŞ

A A A A aǎāEEEĚS

A A A A AŞEÁEEĚS

A A A A A AŞEEÉŞE 3E aǎāEEEŞ

A A A A A AŞEEÁĐCS E E ČD A E E C E aǎāEEÁĐCS


A A A A A AŞEĘEŞE aǎāEEEŞ

A A A A A AŞEEҪCŞ E E E aǎāEEҪCŞ

A A A A aǎāEEEĚS

A A aǎāaEEEĐҪCĒEŞ


aǎāEEEĚEŞ



Vector data Model

Vector data model

- ◆ One-dimensional arrays used for nearest-neighbor search (exact or approximate).
 - ★ Used for semantic search on embeddings generated by ML-trained transformer models (think ChatGPT).
 - ★ Native integration with modern ML tools and APIs (e.g., LangChain, OpenAI).
- ◆ At their core, these systems use specialized indexes to perform NN searches quickly.




CloseVector




marqo

turbopuffer <(°0°)>

Content-Based Retrieval Workflow Using Vector Space Embeddings





OLTP vs. OLAP Workload Characteristics

Online Transaction Processing (OLTP)

OLTP

- ◆ OLTP systems handle operational business processes:
 - ★ requiring fast, consistent transaction processing
 - ★ with high concurrency levels.
- ◆ A class of systems designed for managing real-time, day-to-day business transactions.
- ◆ Examples
 - ★ ATM transactions
 - ★ online banking
 - ★ credit card processing
 - ★ e-commerce purchases.
- ◆ These systems are characterized by
 - ★ high volume
 - ★ short transactions
 - ★ need for quick response times.

Characteristics and Requirements

- ◆ OLTP workloads involve
 - ★ short-duration transactions
 - ★ accessing small amounts of current data.
- ◆ Response times must be predictable and fast, usually under 100 milliseconds.
- ◆ Consider an e-commerce checkout process
 - ★ Multiple users simultaneously place orders, update inventory, and process payments.
 - ★ Each transaction involves few records but requires immediate consistency to prevent overselling products.

Data Access Patterns

- ◆ Use indexed access to retrieve specific records or small record sets.
- ◆ Queries involve
 - ★ primary key lookups
 - ★ simple range scans.
- ◆ Update operations modify
 - ★ individual records
 - ★ small groups of related records.
- ◆ The workload mix typically includes
 - ★ 70-80% read operations
 - ★ 20-30% write operations
 - distributed across many different tables and records.

System Optimizations

- ◆ **Row-oriented storage** optimizes OLTP performance by storing complete records together, minimizing I/O for transactions accessing multiple attributes of the same entity.
- ◆ **B+ tree indexes** provide efficient access paths for both equality and range queries.
- ◆ Connection pooling and prepared statements reduce overhead for repetitive operations.
- ◆ Partitioning strategies often use range or hash partitioning based on transaction date or customer ID to distribute load evenly.

Online Analytical Processing (OLAP)

OLAP

- ◆ OLAP systems support
 - ★ complex analytical queries and business intelligence applications
 - ★ requiring sophisticated data aggregation and analysis capabilities.
- ◆ OLAP systems enable users to analyze large datasets from multiple perspectives.

OLAP Example

- ◆ Analyzing sales data for a retail company using an OLAP cube.
 - ★ The cube could organize sales by
 - product, time, and location,
 - ★ allowing users to drill down into specific
 - regions, products, or time periods
 - ★ to identify
 - trends, such as which products are most popular in a particular city during a specific month.

Characteristics and Requirements

- ◆ OLAP workloads involve ***long-running*** queries that scan large portions of ***historical data*** to compute ***aggregations, trends, and statistical summaries.***
- ◆ ***Response times*** range from seconds to minutes, but ***query complexity*** and data volume are substantially higher than OLTP systems.
- ◆ Consider a retail analytics query calculating seasonal sales trends across multiple product categories, geographical regions, and customer segments over several years.

Data Access Patterns

- ◆ OLAP queries typically involve **full table scans** or large range scans, often accessing millions or billions of records.
- ◆ **Aggregation operations** like SUM, COUNT, and AVG are common, frequently combined with GROUP BY clauses creating multi-dimensional summaries.
- ◆ **Joins** often involve large fact tables with smaller dimension tables in star or snowflake schema configurations.

System Optimizations

- ◆ **Columnar storage** dramatically improves OLAP performance by storing each column separately, enabling efficient compression and reducing I/O for queries accessing few columns.
- ◆ **Bitmap indexes** support complex Boolean queries common in analytical workloads.
- ◆ **Materialized views** pre-compute common aggregations, trading storage space for query performance.
- ◆ **Parallel processing** distributes query execution across multiple processors or machines.

Hybrid Approaches

Hybrid Approaches

- ◆ HTAP (Hybrid Transactional/Analytical Processing)
 - ★ Modern systems increasingly support mixed workloads combining OLTP and OLAP requirements.
 - ★ SAP HANA exemplifies this approach using in-memory storage and columnar organization to support both transaction processing and real-time analytics on the same data.
- ◆ Data Replication Strategies
 - ★ Many organizations maintain separate OLTP and OLAP systems connected through data replication mechanisms.
 - ★ Change Data Capture (CDC) identifies and propagates OLTP system changes to analytical systems with minimal impact on operational performance.



System Optimizations

Physical Design Optimizations

- ◆ Index Selection and Design
- ◆ Partitioning Strategies

Query Optimization Techniques

- ◆ Cost-Based Optimization
- ◆ Advanced Optimization Techniques

Memory and Caching Strategies

- ◆ Buffer Pool Management
- ◆ Query Result Caching

Course Details

Course plan

W	LECTURE	LAB	Coursework
1	Advanced DBMS and Database System Architectures	<ul style="list-style-type: none"> Environment setup 	
2	Storage Management	<ul style="list-style-type: none"> Implementing a buffer pool manager. Benchmarking different buffer pool strategies under various workloads 	Assignment 1 <ul style="list-style-type: none"> Storage Manager and Buffer Pool Implementation with Different Replacement Policies
3	Storage Models - Row-Oriented and Columnar Systems	<ul style="list-style-type: none"> Implement storage engines. Convert row-store to columnar 	
4	Indexing Structures and Implementation	<ul style="list-style-type: none"> Build a B+ tree index; Implement extendible hashing. 	Assignment 2 <ul style="list-style-type: none"> Indexing structures (b+ trees and lsm trees)
5	Advanced Programming and Query Processing	<ul style="list-style-type: none"> Build inventory-auditing triggers. Atomic SPs with error handling Implementing join algorithms and vectorized execution operators 	Quiz 1
6	Query Optimization	<ul style="list-style-type: none"> Implementing a cost-based query optimizer. Tune query plans. 	Assignment 3 <ul style="list-style-type: none"> Query Processing and Optimization
7	Transaction Processing and Concurrency Control	<ul style="list-style-type: none"> Deadlock simulation; lock contention analysis. Analyzing isolation level bugs in production systems 	
8			Mid Term Exam

Course plan

W	LECTURE	LAB	Coursework
9	Crash Recovery	<ul style="list-style-type: none"> Implementing WAL and recovery procedure 	Assignment 4 <ul style="list-style-type: none"> Transaction Processing with Concurrency Control and Recovery
10	Parallel and Distributed Database Architectures	<ul style="list-style-type: none"> Implement cross-shard ACID transactions Distributed transaction coordination 	
11	NoSQL Databases and Big Data Storage Systems	<ul style="list-style-type: none"> Use MongoDB's query language to explore document-based databases. 	Assignment 5 <ul style="list-style-type: none"> Implement NoSQL queries
12	Advanced NoSQL Models.	<ul style="list-style-type: none"> Implement core components of different NoSQL models 	
13	Caching Strategies, ORM Integration, and Modern Database Trends	<ul style="list-style-type: none"> Session data caching API caching E-commerce platforms 	Assignment 6 <ul style="list-style-type: none"> Build a Simple Web-site Integrating an ORM Framework
14	Database Security	<ul style="list-style-type: none"> Implement column-level encryption; audit log triggers 	
15	Course Wrap up	Assignments Presentation as a Whole Project	

Assessment Strategy

Assessment Method	Degrees
Quizzes (best 2 of 3) (week 5, 10, 14)	10
Midterm Exam (week 8)	20
Lecture participation	8
Lab participation	7
Term Project	15
Final Exam	40
Total	100





THANK
YOU

