

CSAI 302

Advanced Database Systems


Lec 02

Storage Management

Storage hierarchy

Volatile
Random Access
Byte-Addressable

Non-Volatile
Sequential Access
Block-Addressable



Disk-based Architecture

Disk-based Architecture

- ◆ The DBMS assumes that the primary storage location of the database is on non-volatile disk.
- ◆ The DBMS's components manage the movement of data between non-volatile and volatile storage.
- ◆ Lowest layer of DBMS
 - ★ Physical details hidden from higher levels of system
- ◆ Purpose
 - ★ Map pages to locations on disk
 - ★ Load pages from disk to memory
 - ★ Save pages back to disk & ensuring write

Terminologies

- ◆ **Block**

- ★ unit of transfer for disk read/write

- ◆ **Page**

- ★ a block-sized chunk of RAM

Files of Pages of Records

- ◆ Each **table** is stored in one or more **OS files**
- ◆ Each **file** contains many **pages**
- ◆ Each **page** contains many **records**
- ◆ **Pages** are understood by multiple layers
 - ★ Managed on disk by the **disk space manager**:
 - pages read from/written to physical disk/files
 - ★ Managed in memory by the **buffer manager**:
 - higher levels of DBMS only operate in memory


Representing Data Elements

	Data Element	Record	Collection
DBMS	Attribute	Tuple	Relation
File System	Field	Record	File

System design goals

- ◆ Allow the DBMS to manage databases that exceed the amount of memory available.
- ◆ Reading/writing to disk is expensive
 - ★ It must be managed carefully to avoid large stalls and performance degradation.
- ◆ Random access on disk is usually much slower than sequential access
 - ★ the DBMS needs to maximize sequential access.

Disk-oriented DBMS



Database Storage Questions

- ◆ How the DBMS represents the database in files on disk?
- ◆ How the DBMS manages its memory and moves data back-and-forth from disk?

File storage

- ◆ The DBMS stores a database as one or more files on disk typically in a proprietary format.
 - ★ The OS does not know anything about the contents of these files.
- ◆ Early systems in the 1980s used custom filesystems on raw block storage.
 - ★ Some enterprise DBMSs still support this.
 - ★ Most newer DBMSs do not do this.

Storage manager

- ◆ Responsible for maintaining a database's files.
- ◆ Organizes the files as a collection of **pages**.
 - ★ Tracks data read/written to pages.
 - ★ Tracks the available space.
- ◆ A DBMS typically does **not** maintain multiple copies of a page on disk.

Database pages

- ◆ A **page** is a fixed-size block of data.
 - ★ It can contain tuples, meta-data, indexes, log records...
 - ★ Most systems do not mix page types.
 - ★ Some systems require a page to be self-contained.
- ◆ Each page is given a unique identifier (**page ID**).
 - ★ A page ID could be unique per DBMS instance, per database, or per table.
 - ★ The DBMS uses an indirection layer to map page IDs to physical locations.

Database pages

- ◆ There are three different notions of "pages" in a DBMS:
 - ★ Hardware Page (usually 4KB)
 - ★ OS Page (usually 4KB, x64 2MB/1GB)
 - ★ Database Page (512B-32KB)
- ◆ A hardware page is the largest block of data that the storage device can guarantee failsafe writes.
- ◆ DBMSs that specialize in read-only workloads have larger page sizes.

★ **4KB**



★ **8KB**



★ **16KB**



Page storage architecture

- ◆ Different DBMSs manage pages in files on disk in different ways.

★ **Heap** File Organization

★ **Tree** File Organization

★ **Sequential** / Sorted File Organization (ISAM)

★ **Hashing** File Organization

- ◆ At this point in the hierarchy, we do **not** need to know anything about what is inside of the pages.

DB File Structures

- ◆ Unordered heap files
 - ★ Records placed arbitrarily across pages
 - ★ As file shrinks/grows, pages (de)allocated
 - ★ Suitable when typical access is a full scan of all records
- ◆ Clustered heap files
 - ★ Group data into blocks to enable fast lookup and efficient modifications
- ◆ Sorted files
 - ★ Pages and records are in strict sorted order
 - ★ Best for retrieval in order, or when a range of records is needed
- ◆ Index files
 - ★ B+ trees, linear hashing, extensible hashing,
 - ★ May contain records or point to records in other files


Heap File organization

Heap file

- ◆ A heap file is an unordered collection of pages with tuples that are stored in random order.
- ◆ Need ***additional meta-data*** to track location of files and free space availability.
- ◆ Heap files has one special ***header page***
 - ★ Two pointers for **free space** and **data**
- ◆ Location of the heap file and the header page are saved in ***DB catalog***

Heap file

Get Page #2




Heap file: Page directory


- ◆ The DBMS maintains special pages that tracks the ***location of*** data ***pages in*** the database ***files***.
 - ★ One entry per database object.
 - ★ Must make sure that the directory pages are in sync with the data pages.
- ◆ DBMS also keeps meta-data about pages' contents:
 - ★ Amount of **free space** per page.
 - ★ List of **free / empty** pages.
 - ★ Page **type** (data vs. meta-data).

Page directory

File Location \Rightarrow Page# \times PageSize




Heap file: Page directory



Unordered Heap Files With Page Directories

- ◆ Page directories, with multiple header pages, each encoding
 - ★ A pointer to page
 - ★ # free bytes on the page
- ◆ Header pages are accessed often, so likely deployed in cache
- ◆ Finding a page to fit a record can be done efficiently
 - ★ One header page load reveals free space of MANY pages




Page Layout

Page header

- ◆ Every page contains a header of meta-data about the page's contents.

- ★ Page Size
- ★ Checksum
- ★ DBMS Version
- ★ Transaction Visibility
- ★ Compression / Encoding Meta-data
- ★ Schema Information
- ★ Data Summary / Sketches



- ◆ Some systems require pages to be **self-contained** (e.g., Oracle).

- ★ Data and its associated Metadata should reside together on the same physical storage unit.

Page layout

- ◆ For any page storage architecture, we now need to decide how to organize the data inside of the page.
 - ★ assuming that we are only storing ***tuples*** in a **row-oriented storage** model.

Tuple-oriented Storage

Log-structured Storage


Index-organized Storage

Tuple-oriented storage

- ◆ How to store tuples in a page?
 - ★ Keep track of the number of tuples in a page and then just append a new tuple to the end.
 - What happens if we delete a tuple?
 - What happens if we have a variable-length attribute?

Tuples in a page


Page




Page



Page




Page



Variable-length attribute




- ◆ The most common layout scheme is called **slotted pages**.
- ◆ The slot array maps "slots" to the tuples' starting position offsets.
- ◆ The header keeps track of:
 - ★ The # of used slots
 - ★ The offset of the starting location of the last slot used.

Slotted pages



Fixed- and Var-length Tuple Data

Delete tuple



Record IDs

- ◆ The DBMS assigns each logical tuple a unique **record identifier** that represents its physical location in the database.
 - ★ File Id, Page Id, Slot #
 - ★ Most DBMSs do not store ids in tuple.
 - ★ SQLite uses **ROWID** as the true primary key and stores them as a hidden attribute.
- ◆ Applications should **never** rely on these IDs to mean anything.

Record IDs



CTID
(6-bytes)



ROWID
(8-bytes)




%%physloc%%
(8-bytes)



ROWID
(10-bytes)

Tuple Layout


Tuple layout



- ◆ A tuple is essentially a sequence of bytes.
 - ★ These bytes do not have to be contiguous.
- ◆ It is the job of the DBMS to interpret those bytes into attribute types and values.
- ◆ Each tuple is prefixed with a **header** that contains meta-data about it.
 - ★ Visibility info (concurrency control)
 - ★ Bit Map for **NULL** values.

Tuple data

- ◆ Attributes are typically stored in the order that you specify them when you create the table.
- ◆ This is done for software engineering reasons (i.e., simplicity).
- ◆ However, it might be more efficient to lay them out differently.



```
CREATE TABLE foo (
    a INT PRIMARY KEY,
    b INT NOT NULL,
    c INT,
    d DOUBLE,
    e FLOAT
);
```

Denormalized tuple data

- ◆ DBMS can physically denormalize (e.g., "pre-join") related tuples and store them together in the same page.

★ Potentially **reduces the amount of I/O** for common workload patterns.

★ Can make **updates more expensive**.

```
CREATE TABLE foo (
  a INT PRIMARY KEY,
  b INT NOT NULL,
);
```

```
CREATE TABLE bar (
  c INT PRIMARY KEY,
  a INT REFERENCES foo (a),
);
```

Denormalized tuple data

foo

Header	a	b
--------	---	---

```
SELECT * FROM foo JOIN bar  
ON foo.a = bar.a;
```

bar

Header	c	a
Header	c	a
Header	c	a



foo

Header	a	b	c	c	c	...
foo	foo	bar	bar	bar	bar	...

Denormalized tuple data

- ◆ Not a new idea.

- ★ IBM System R did this in the 1970s.

- ★ Several NoSQL DBMSs do this without calling it physical denormalization.



RethinkDB



CouchDB



Other Memory Pools

Other memory pools

- ◆ The DBMS needs memory for things other than just tuples and indexes.
- ◆ These other memory pools may not always backed by disk, depends on implementation.
 - ★ Sorting + Join Buffers
 - ★ Query Caches
 - ★ Maintenance Buffers
 - ★ Log Buffers
 - ★ Dictionary Caches

Buffer pool organization


- ◆ Memory region organized as an array of **fixed-size pages**.

*Buffer
Pool*

- ◆ An array entry is called a **frame**.

- ◆ When the DBMS requests a page, an exact copy is placed into one of these frames.


- ◆ Dirty pages are buffered and **not** written to disk immediately



★ Write-Back Cache

Buffer pool meta-data

- ◆ The ***page table*** keeps track of pages
 - ◆ that are currently in memory.
 - ★ Usually a fixed-size hash table protected with latches to ensure thread-safe access.
- ◆ Additional meta-data per page:
 - ★ Dirty Flag
 - ★ Pin/Reference Counter
 - ★ Access Tracking Information



Page table vs. Page directory

Page directory

- ◆ The mapping from page ids to page locations in the database files.
 - ★ All changes must be recorded on disk to allow the DBMS to find on restart.

Page table

- ◆ The mapping from page ids to a copy of the page in buffer pool frames.
 - ★ This is an in-memory data structure that does not need to be stored on disk.

when a Page is Requested ...

1. If requested page is NOT in the buffer pool:
 1. Choose an **un-pinned** (pin count = 0) frame for replacement
 2. If frame is **dirty**, write current page to disk, mark as "**clean**"
 3. Read requested page into frame
 2. Pin the page and return its address
- ◆ If requests can be predicted (e.g., sequential scans), pages can be **prefetched**
 - ★ Several pages at a time
 - ◆ When the buffer pool is full ...
 - ★ Need to evict an existing page
 - ★ Need a **page replacement policy**


Buffer Replacement

Buffer Replacement Policies

- ◆ When the DBMS needs to free up a frame to make room for a new page, it must decide which page to **evict** from the buffer pool.
- ◆ Goals:
 - ★ Correctness
 - ★ Accuracy
 - ★ Speed
 - ★ Meta-data overhead

LEAST-RECENTLY USED

- ◆ Maintain a single timestamp of when each page was last accessed.
 - ◆ When the DBMS needs to evict a page, select the one with the oldest timestamp.
- ★ Keep the pages in sorted order to reduce the search time on eviction.



LRU Replacement Policy

◆ Least Recently Used (LRU)

- ★ **Pinned frame**: not available to replace
- ★ Track time for each frame last unpinned (end of use)
- ★ Replace the frame which was least recently used
 - Need to “find min” on the last used attributed (priority queue)

- ★ Very common policy: intuitive and simple
 - Good for repeated access to popular pages (temporal locality)

◆ Sequential Scan + LRU

- ★ Lead to **sequential flooding**
 - 0% hit rate in cache
- ★ Repeated sequential scan is very common in DB workload
 - Nested-loop join

Better Policies: LRU-K

- ◆ Track the history of last K references to each page as timestamps and compute the interval between subsequent accesses.
 - ★ Can distinguish between reference types
- ◆ Use this history to estimate the next time that page is going to be accessed.
 - ★ Replace the page with the oldest "K-th" access.
 - ★ Balances recency vs. frequency of access.
 - ★ Maintain an ephemeral in-memory cache for recently evicted pages to prevent them from always being evicted.




Example: LRU-3

- ◆ Page A
 - ★ Access at T=10 , T=40 , T=80
- ◆ Page B
 - ★ Access at T=5 , T=30 , T=90


MySQL approximate LRU-K

- ◆ Single LRU linked list but with two entry points ("old" vs "young").
 - ★ New pages are always inserted to the head of the old list.
 - ★ If pages in the old list is accessed again,
 - then insert into the head of the young list.


MySQL approximate LRU-K




Disk Pages




MySQL approximate LRU-K




Disk Pages




MySQL approximate LRU-K



Disk Pages



MySQL approximate LRU-K



Better policies: Localization

- ◆ The DBMS chooses which pages to evict on a per query basis.
- ◆ This minimizes the pollution of the buffer pool from each query.
 - ★ Keep track of the pages that a query has accessed.
- ◆ Example
 - ★ Postgres assigns a limited number of buffer pool pages to a query and uses it as a ***circular ring buffer***.

Better policies: Priority hints

- ◆ The DBMS knows about the context of each page during query execution.
- ◆ It can provide hints to the buffer pool on whether a page is important or not.

Dirty Pages

- ◆ Fast Path
 - ★ If a page in the buffer pool is ***not*** dirty, then the DBMS can simply "drop" it.
- ◆ Slow Path
 - ★ If a page is ***dirty***, then the DBMS must write back to disk to ensure that its changes are persisted.
- ◆ Trade-off between ***fast evictions*** versus dirty writing pages that will ***not be read again*** in the future.

Background writing

- ◆ The DBMS can periodically walk through the page table and write dirty pages to disk.
- ◆ When a dirty page is safely written, the DBMS can either evict the page or just unset the dirty flag.
- ◆ Need to be careful that the system doesn't write dirty pages before their log records are written.

Buffer Pool Optimizations

Multiple Buffer Pools
Pre-Fetching
Scan Sharing
Buffer Pool Bypass

Multiple buffer pools

- ◆ The DBMS does not always have a single buffer pool for the entire system.
 - ★ Multiple buffer pool instances
 - ★ Per-database buffer pool
 - ★ Per-page type buffer pool
- ◆ Partitioning memory across multiple pools helps reduce latch contention and improve locality.
 - ★ Avoids contention on LRU tracking meta-data.




Multiple buffer pools

- ◆ Object Id

- ★ Embed an object identifier in record ids and then maintain a mapping from objects to specific buffer pools.

- ◆ Hashing

- ★ Hash the page id to select which buffer pool to access.



Pre-fetching

- ◆ The DBMS can also prefetch pages based on a query plan.
 - ★ Examples: Sequential vs. Index Scans
- ◆ Some DBMS prefetch to fill in empty frames upon start-up.



THANK
YOU ☺

