

**CSAI 302**

# **Advanced Database Systems**

**Lec 03**

**Storage Management [2]**

# Tuple-oriented storage

## ◆ Insert a new tuple:

- ★ Check page directory to find a page with a free slot.
  - Retrieve the page from disk (if not in memory).
- ★ Check slot array to find empty space in page that will fit.

## ◆ Update an existing tuple using its record id:

- ★ Check page directory to find location of page.
  - Retrieve the page from disk (if not in memory).
- ★ Find offset in page using slot array.
  - If new data fits, overwrite existing data.
  - Otherwise, mark existing tuple as deleted and insert new version in a different page.

# Problems

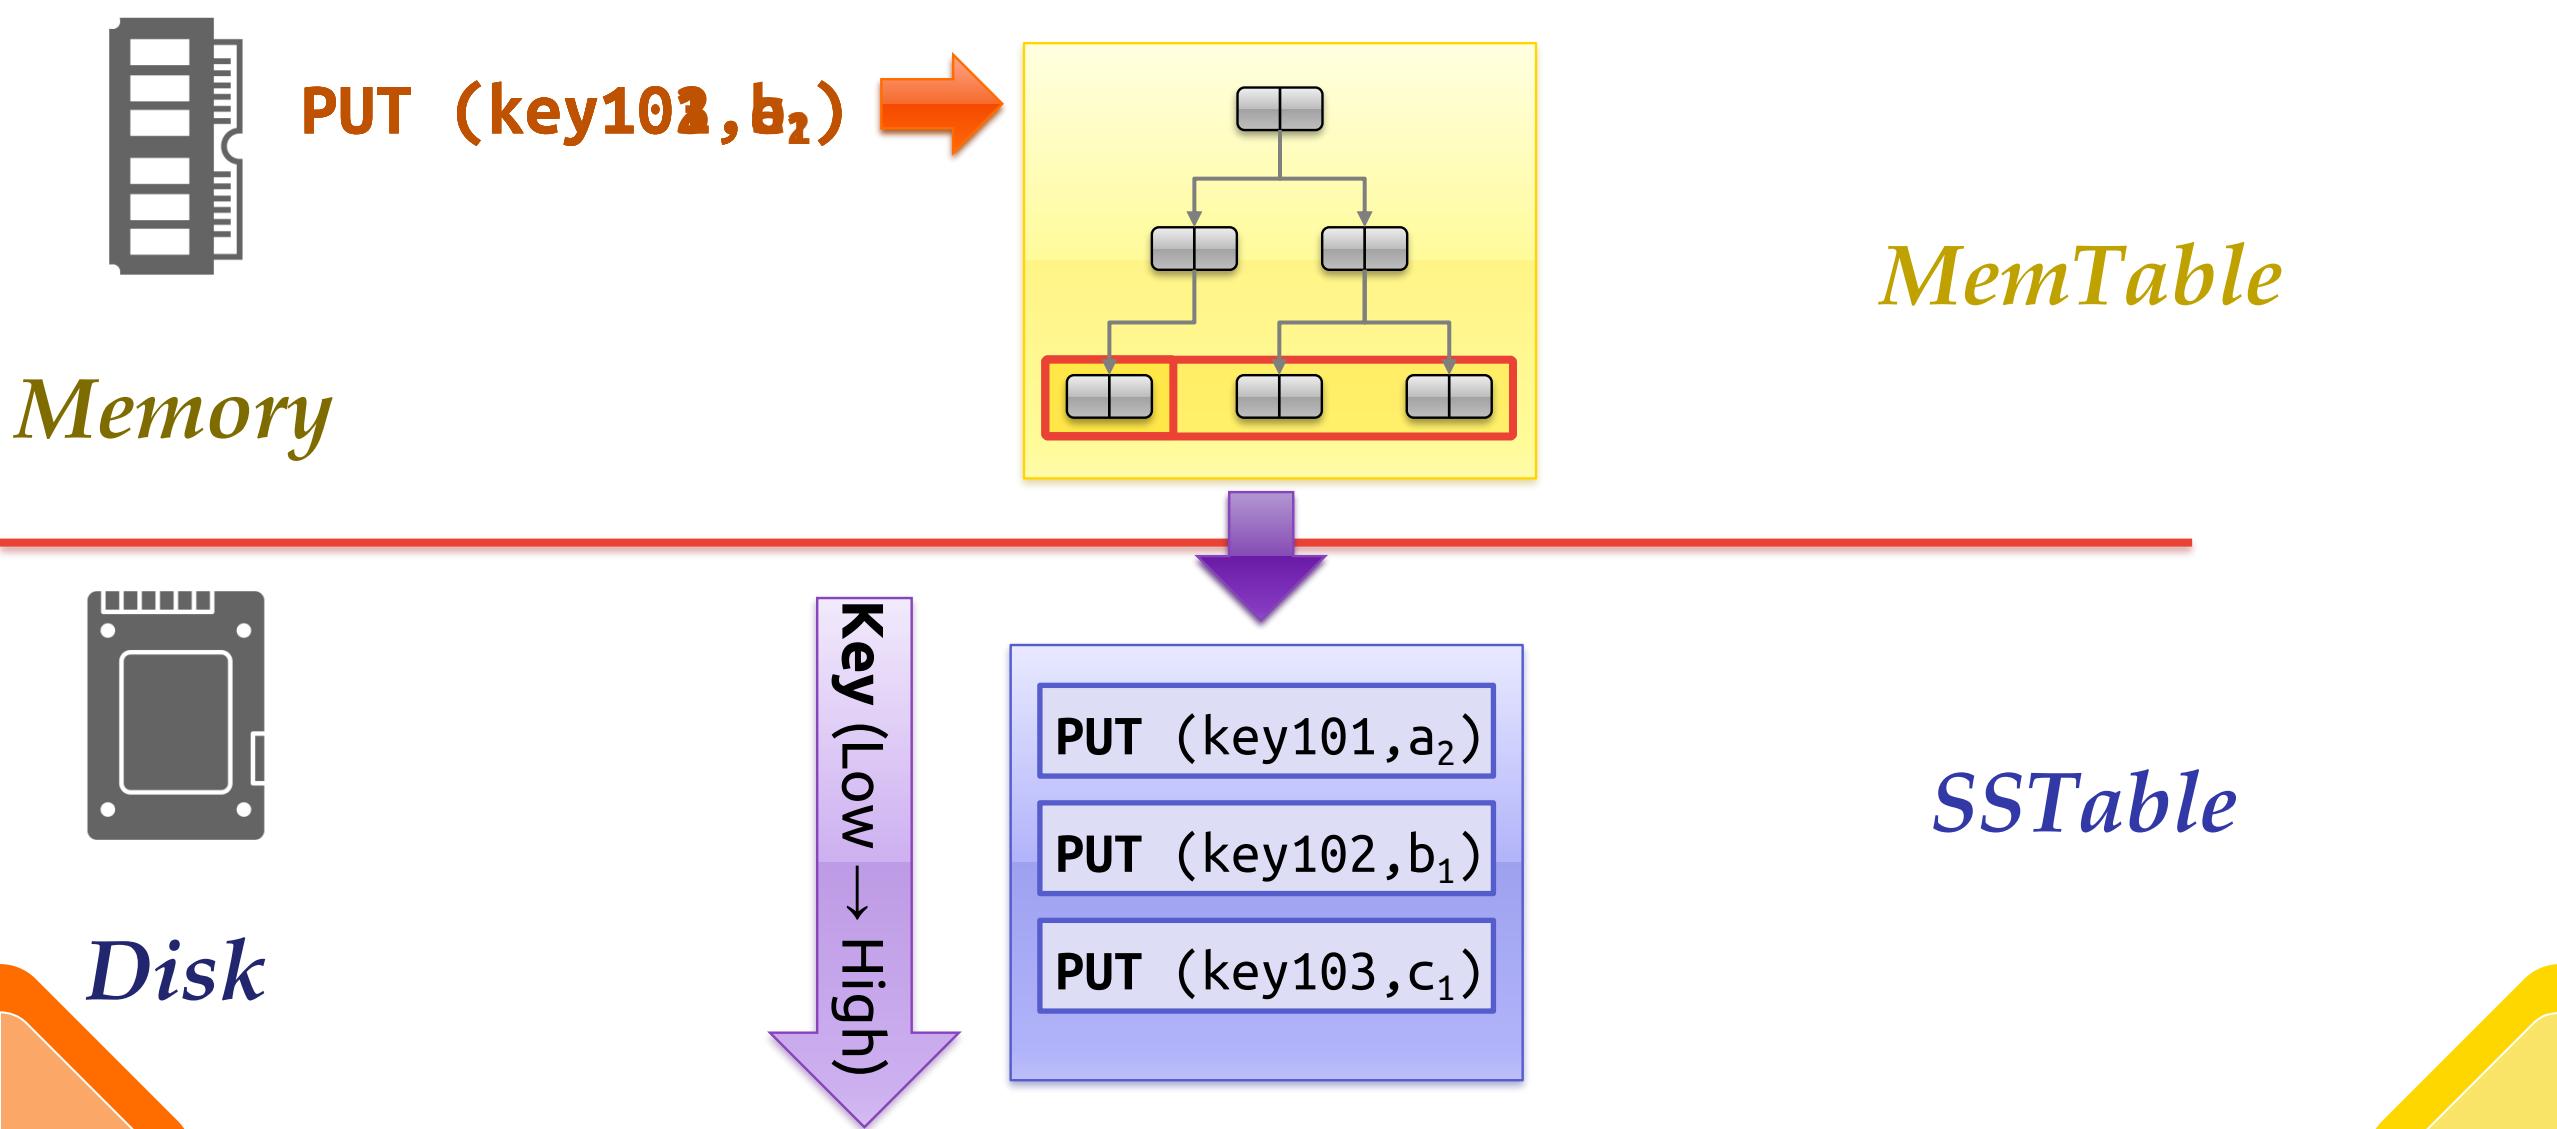
- ◆ Fragmentation
  - ★ Pages are not fully utilized (unusable space, empty slots).
- ◆ Useless Disk I/O
  - ★ DBMS must fetch entire page to update one tuple.
- ◆ Random Disk I/O
  - ★ Worse case scenario when updating multiple tuples is that each tuple is on a separate page.
- ◆ What if the DBMS cannot overwrite data in pages and could only create new pages?
  - ★ Examples: Some object stores, **HDFS, Google Colossus**

# **Log-structured storage**

# Log-structured Storage

- ◆ Instead of storing tuples in pages and updating them in-place, the DBMS maintains a log that records changes to tuples.
  - ★ Each log entry represents a tuple **PUT/DELETE** operation.
  - ★ Originally proposed as **log-structure merge** trees (LSM Trees) in 1996.
- ◆ The DBMS applies changes to an in-memory data structure (**MemTable**) and then writes out the changes sequentially to disk (**SSTable**).

# Log-structured storage

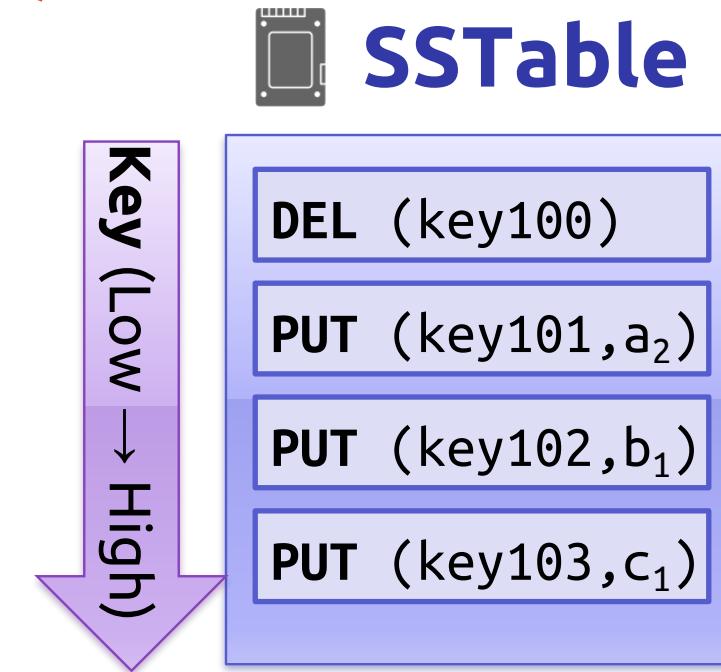


# Log-structured storage

- ◆ Key-value storage that appends log records on disk to represent changes to tuples (**PUT**, **DELETE**).

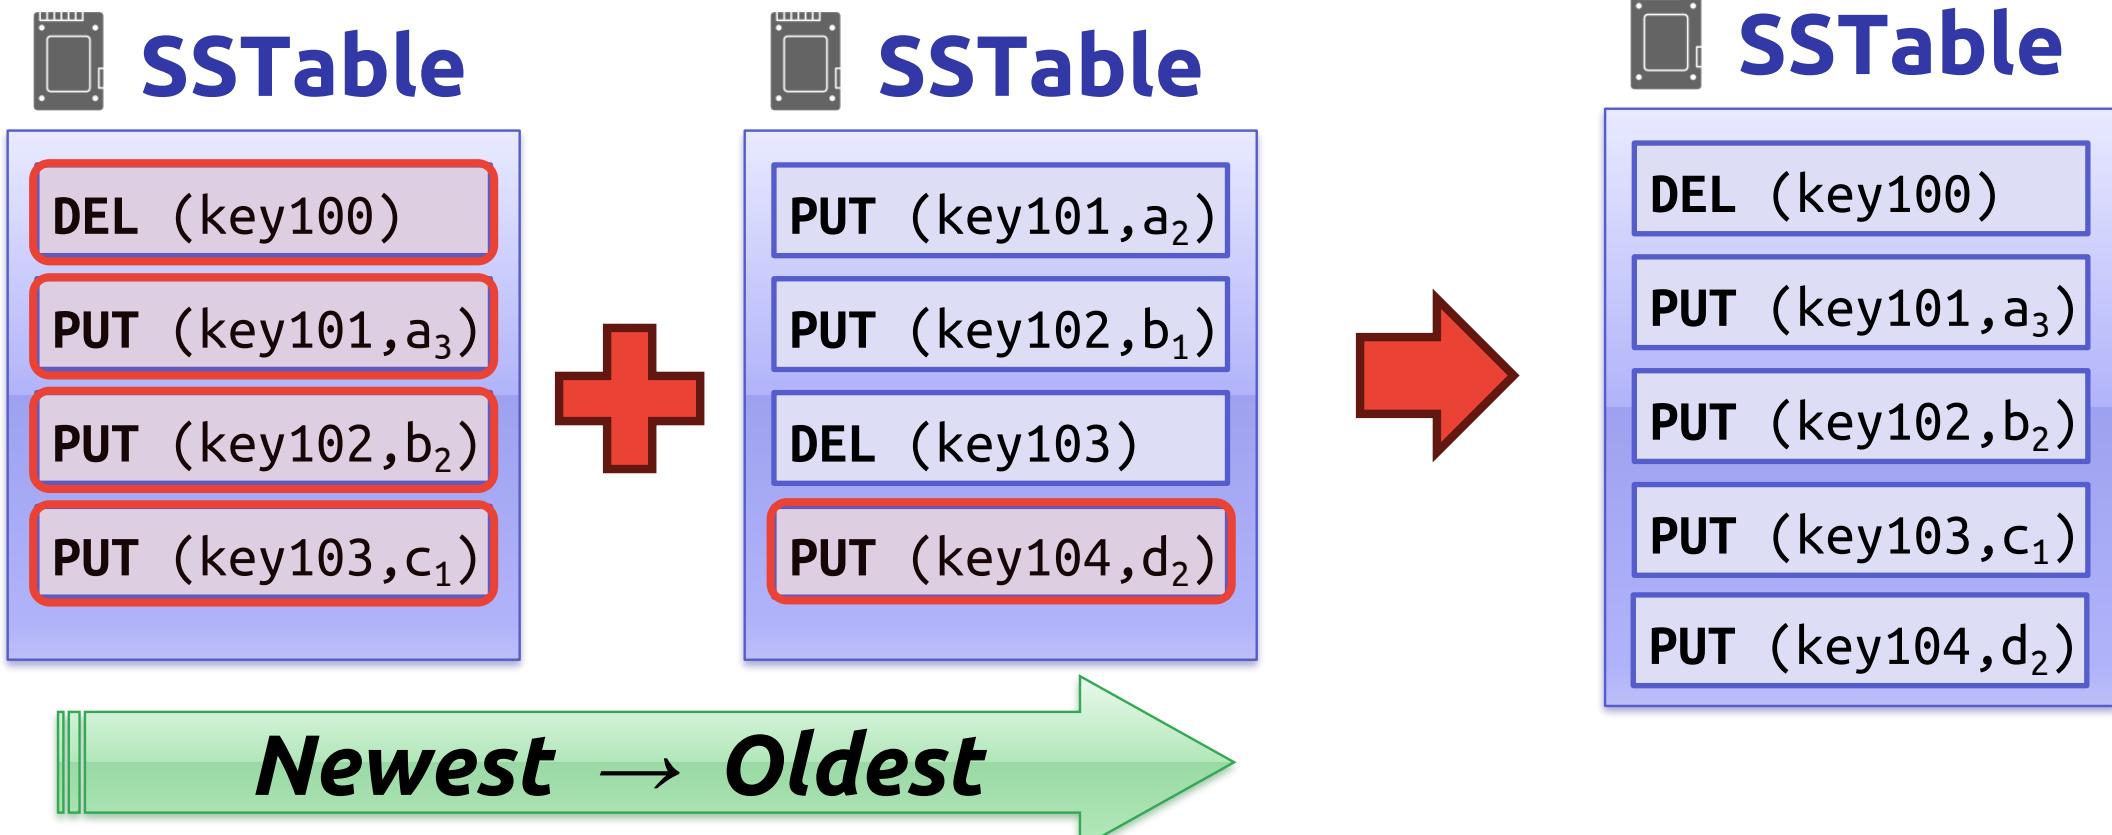
- ★ Each log record must contain the tuple's unique identifier.
- ★ Put records contain the tuple contents.
- ★ Deletes marks the tuple as deleted.

- ◆ As the application makes changes to the database, the DBMS appends log records to the end of the file without checking previous log records.



# Log-structured compaction

- ◆ Periodically compact SSTables to reduce wasted space and speed up reads.
  - ★ Only keep the "latest" values for each key using a sort-merge algorithm.



# Discussion

- ◆ Log-structured storage managers are more common today than in previous decades.
  - ★ This is partly due to the proliferation of RocksDB.



RocksDB

WIREDTIGER



yugabyteDB  
*cassandra*

APACHE  
HBASE



- ◆ What are some downsides of this approach?
  - ★ Write-Amplification.
  - ★ Compaction is expensive.

# **Index-organized Storage**

# Observation

- ◆ The two-table storage approach relies on indexes to find individual tuples.
  - ★ Such indexes are necessary because the tables are inherently unsorted.
- ◆ But what if the DBMS could keep tuples sorted automatically using an index?

# Index-organized Storage

- ◆ DBMS stores a table's tuples as the value of an index data structure.



- ★ Still use a page layout that looks like a slotted page.



- ★ Tuples are typically sorted in page based on key.

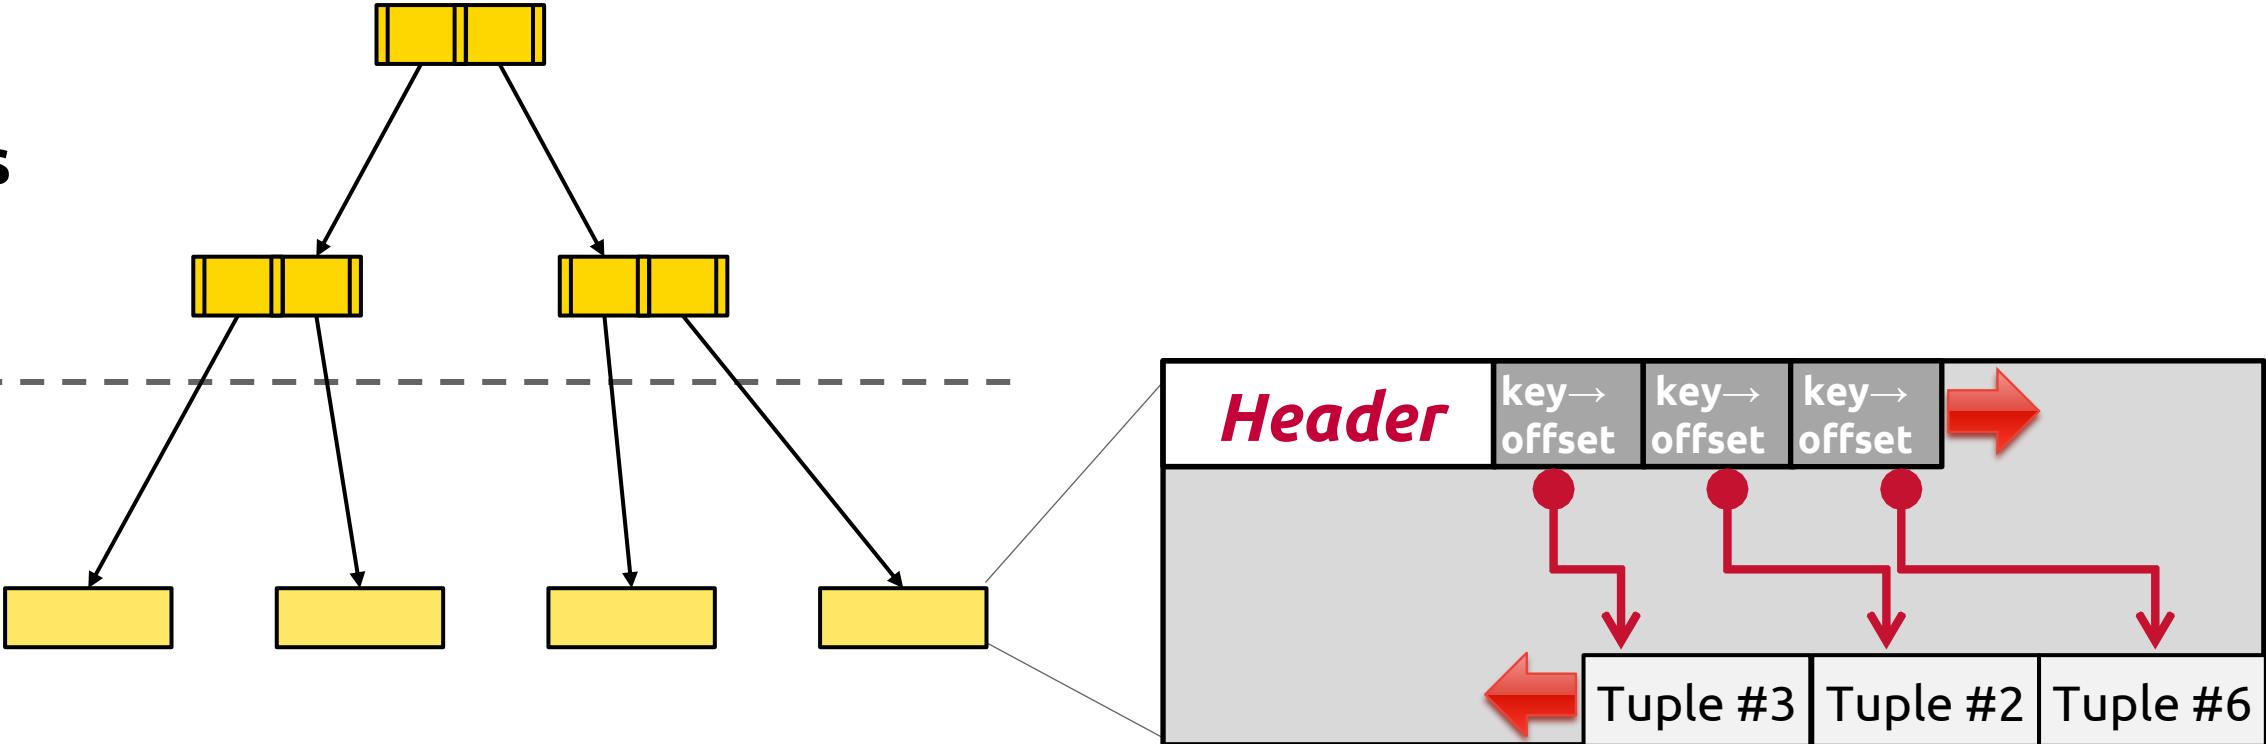


- ◆ B+Tree pays maintenance costs upfront, whereas LSMs pay for it later.



# Index-organized Storage

Inner  
Nodes



# Tuple storage

- ◆ A tuple is essentially a sequence of bytes prefixed with a header that contains meta-data about it.
- ◆ It is the job of the DBMS to interpret those bytes into attribute types and values.
- ◆ The DBMS's catalogs contain the schema information about tables that the system uses to figure out the tuple's layout.

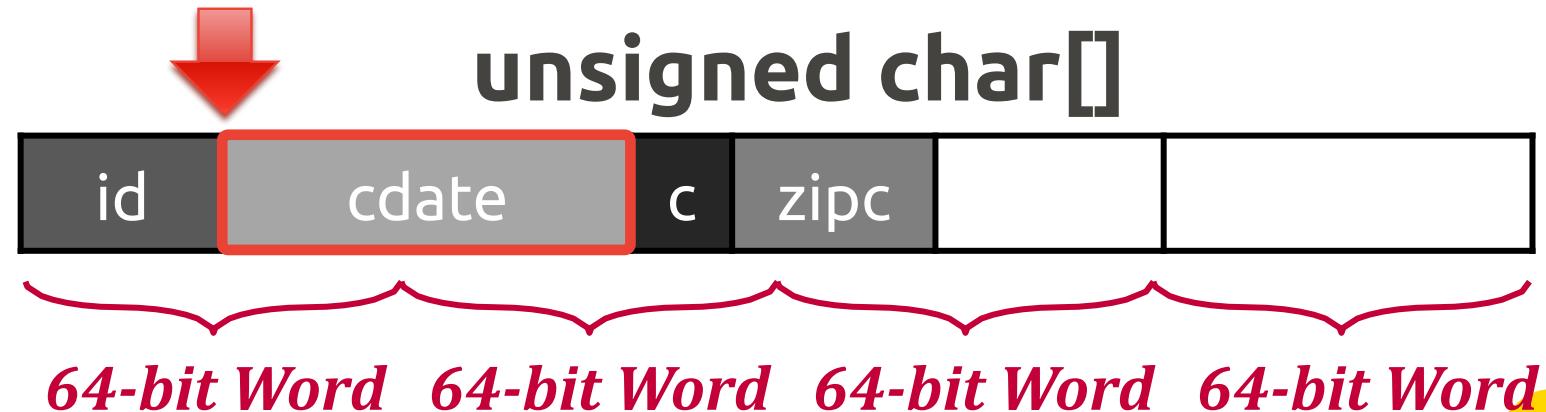
```
CREATE TABLE foo (  
    id INT PRIMARY KEY,  
    value BIGINT  
)
```



# Word-aligned tuples

- ◆ All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

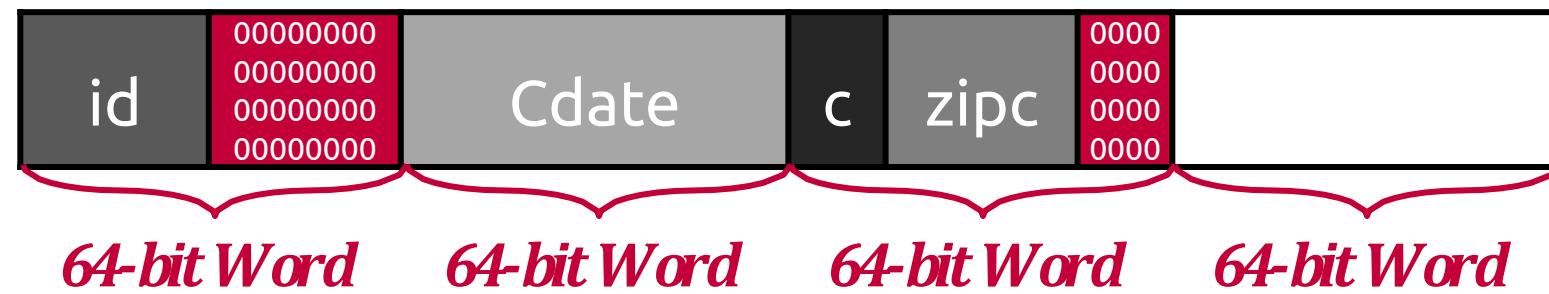
```
CREATE TABLE foo (  
    id INT PRIMARY KEY,  
    cdate TIMESTAMP,  
    color CHAR(2),  
    zipcode INT  
);
```



# Word-alignment: Padding

- ◆ Add empty bits after attributes to ensure that tuple is word aligned.
- ◆ Essentially round up the storage size of types to the next largest word size.

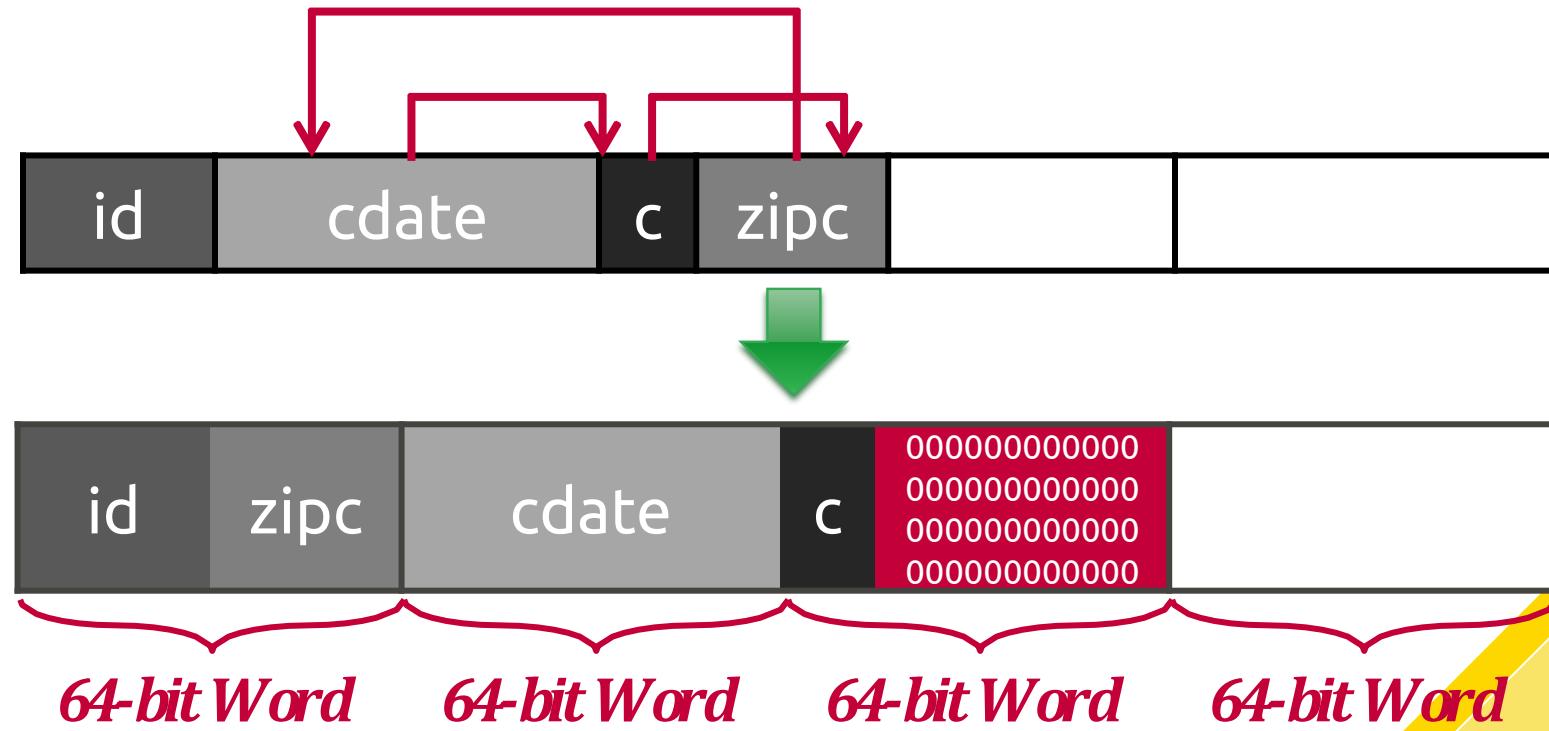
```
CREATE TABLE foo (  
    id INT PRIMARY KEY,  
    cdate TIMESTAMP,  
    color CHAR(2),  
    zipcode INT  
);
```



# Word-alignment: Reordering

- ◆ Switch the order of attributes in the tuples' physical layout to make sure they are aligned.
  - ★ May still have to use padding to fill remaining space.

```
CREATE TABLE foo (  
 32-bits id INT PRIMARY KEY,  
 64-bits cdate TIMESTAMP,  
 16-bits color CHAR(2),  
 32-bits zipcode INT  
)
```



# **Storage models**

# Storage models

- ◆ A DBMS's storage model specifies how it physically organizes tuples on disk and in memory.
  - ★ N-ary Storage Model (**NSM**)
  - ★ Decomposition Storage Model (**DSM**)
  - ★ Hybrid Storage Model (**PAX**)
- ◆ Can have different performance characteristics based on the target workload (OLTP vs. OLAP).
- ◆ Influences the design choices of the rest of the DBMS.

# **N-ary Storage Model (NSM)**

# N-ary Storage Model (NSM)

- ◆ The DBMS stores (almost) all attributes for a single tuple contiguously in a single page.
  - ★ Also commonly known as a **row store**
- ◆ Ideal for OLTP workloads where queries are more likely to access individual entities and execute write-heavy workloads.
- ◆ NSM database page sizes are typically some constant multiple of **4 KB** hardware pages.

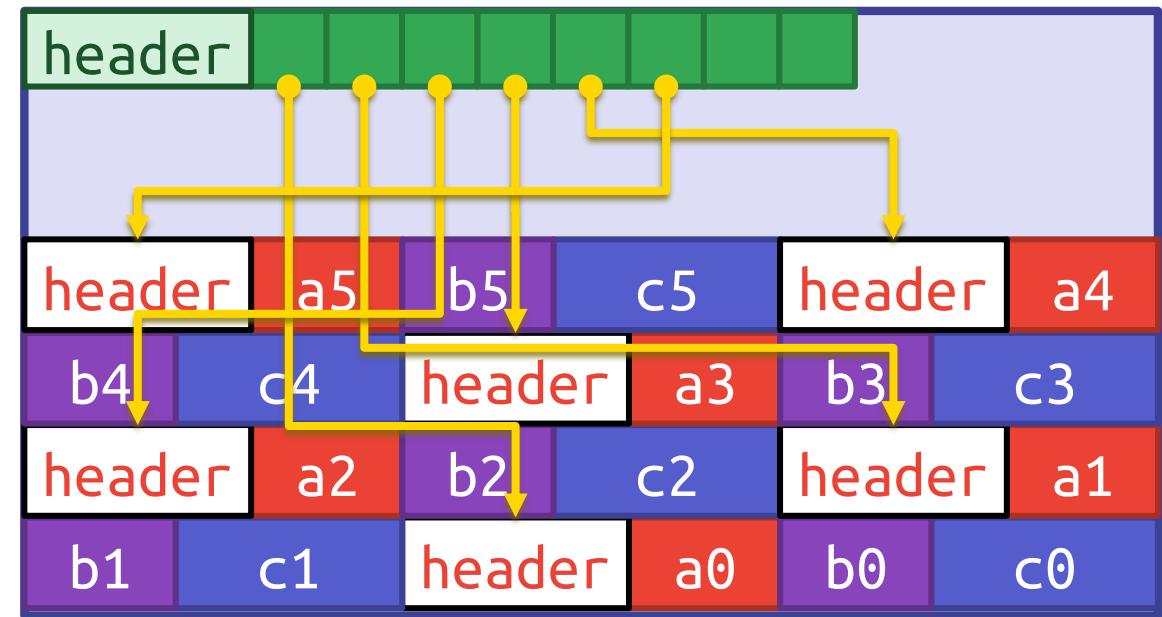
# NSM: Physical Organization

- ◆ A disk-oriented NSM system stores a tuple's fixed-length and variable-length attributes contiguously in a single slotted page.
- ◆ The tuple's **record id** (page#, slot#) is how the DBMS uniquely identifies a physical tuple.

# NSM: Physical Organization

	Col A	Col B	Col C
Row#0	a0	b0	c0
Row#1	a1	b1	c1
Row#2	a2	b2	c2
Row#3	a3	b3	c3
Row#4	a4	b4	c4
Row#5	a5	b5	c5

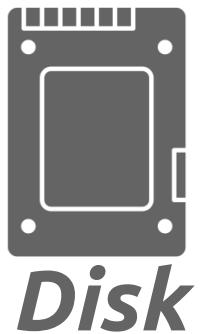
Database Page



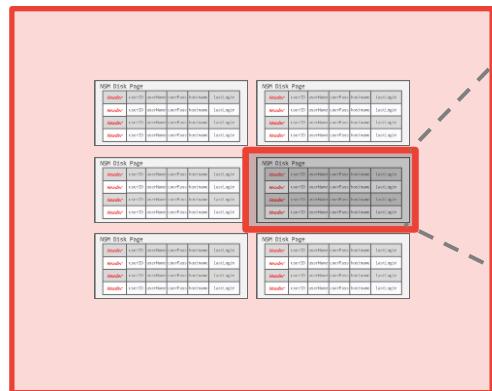
# NSM: OLTP example

**SELECT \* FROM useracct  
WHERE userName = ?  
AND userPass = ?**

**INSERT INTO useracct  
VALUES (?,?,...?)**

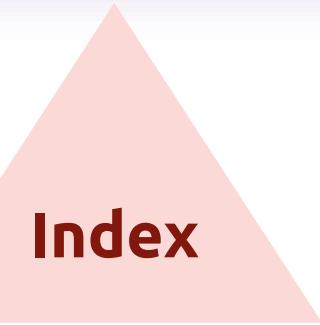


Database File



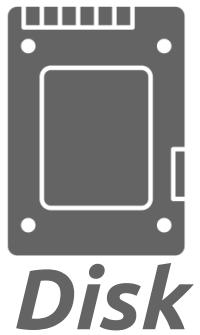
**NSM Disk Page**

header	userID	userName	userPass	hostname	lastLogin
header	userID	userName	userPass	hostname	lastLogin
header	userID	userName	userPass	hostname	lastLogin
header	-	-	-	-	-

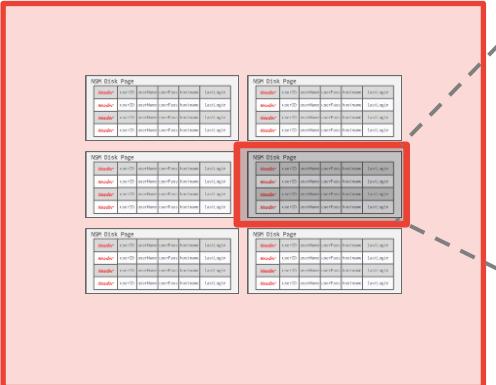


# NSM: OLAP example

```
SELECT COUNT(U.lastLogin),  
       EXTRACT(month FROM U.lastLogin) AS month  
  FROM useracct AS U  
 WHERE U.hostname LIKE '%.gov'  
 GROUP BY EXTRACT(month FROM U.lastLogin)
```



Database File



## NSM Disk Page

header	userID	userName	userPass	hostname	lastLogin
header	userID	userName	userPass	hostname	lastLogin
header	userID	userName	userPass	hostname	lastLogin
header	userID	userName	userPass	hostname	lastLogin



Useless Data!

# NSM: summary

## ◆ Advantages

- ★ Fast inserts, updates, and deletes.
- ★ Good for queries that need the entire tuple (OLTP).
- ★ Can use index-oriented physical storage for clustering.

## ◆ Disadvantages

- ★ Not good for scanning large portions of the table and/or a subset of the attributes.
- ★ Terrible memory locality in access patterns.
- ★ Not ideal for compression because of multiple value domains within a single page.

# **Decomposition Storage Model (DSM)**

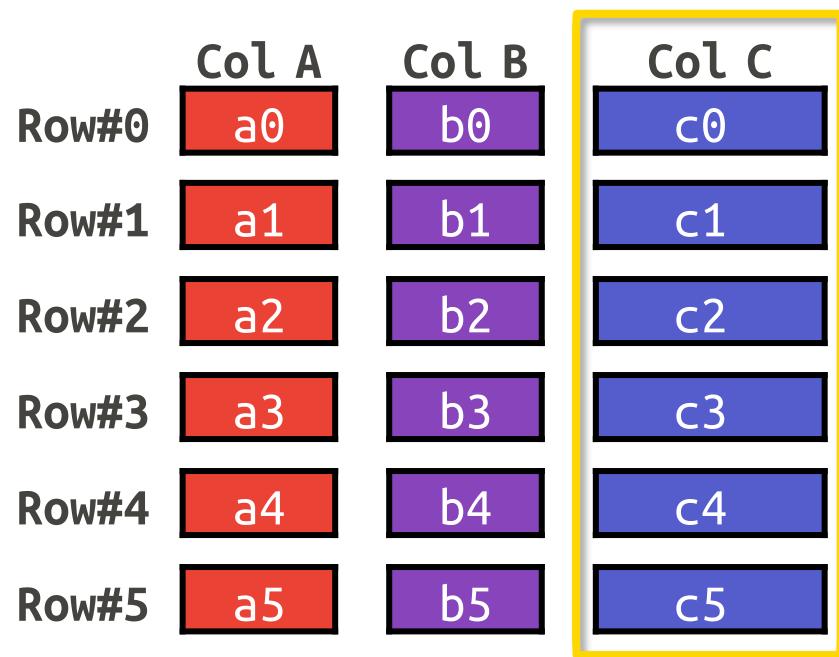
# Decomposition Storage Model (DSM)

- ◆ Store a single attribute for all tuples contiguously in a block of data.
  - ★ Also known as a "***column store***"
- ◆ Ideal for OLAP workloads where read-only queries perform large scans over a subset of the table's attributes.
- ◆ DBMS is responsible for combining/splitting a tuple's attributes when reading/writing.

# DSM: Physical Organization

- ◆ Store attributes and meta-data (e.g., nulls) in separate arrays of fixed-length values.
  - ★ Most systems identify unique physical tuples using offsets into these arrays.
  - ★ Need to handle variable-length values.
- ◆ Maintain separate pages per attribute with a dedicated header area for meta-data about entire column.

# DSM: Physical Organization



Page #1

header	Null bitmap					
a0	a1	a2	a3	a4	a5	

Page #2

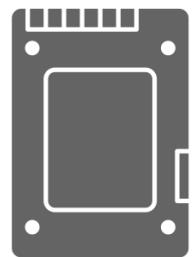
header	Null bitmap					
b0	b1	b2	b3	b4	b5	

Page #3

header	Null bitmap					
c0	c1	c2	c3	c4	c5	

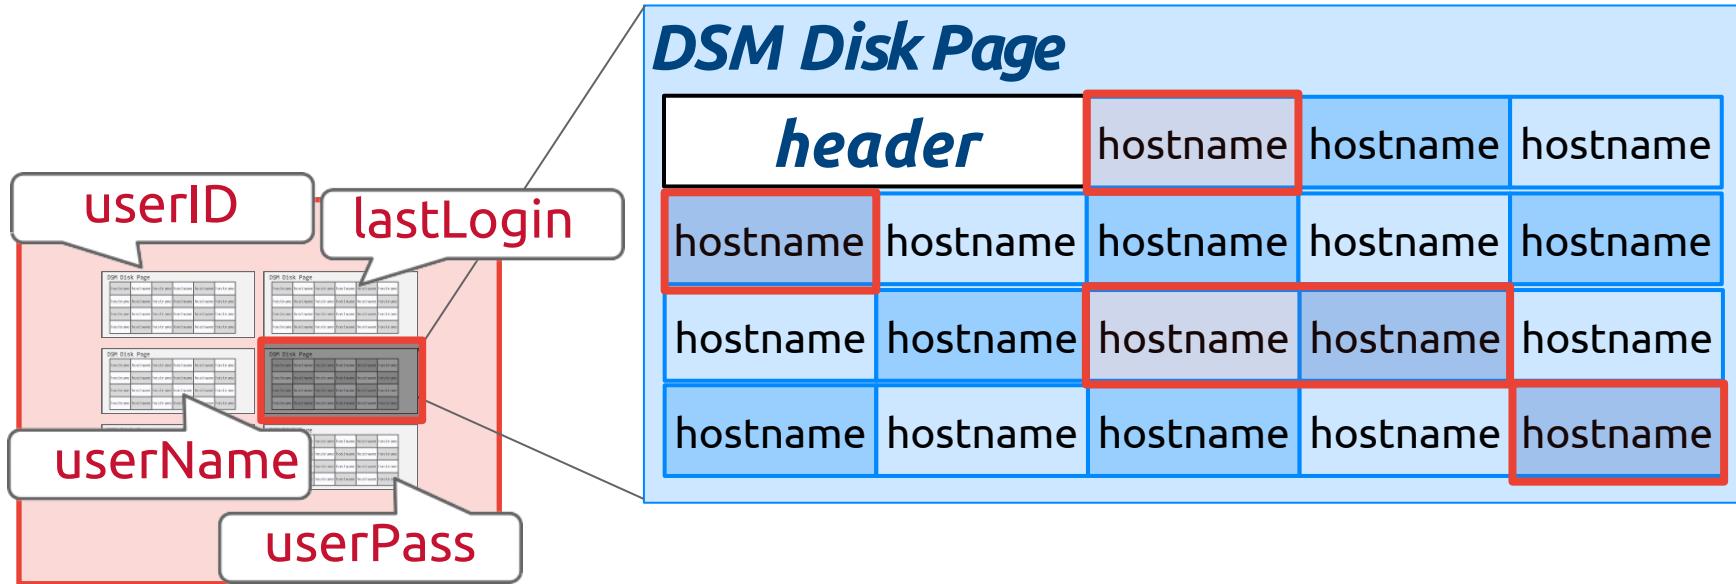
# DSM: OLAP example

```
SELECT COUNT(U.lastLogin),  
       EXTRACT(month FROM U.lastLogin) AS month  
  FROM useracct AS U  
 WHERE U.hostname LIKE '%.gov'  
 GROUP BY EXTRACT(month FROM U.lastLogin)
```



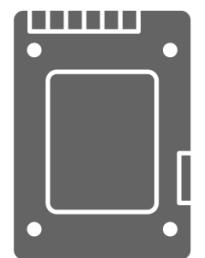
Disk

Database File



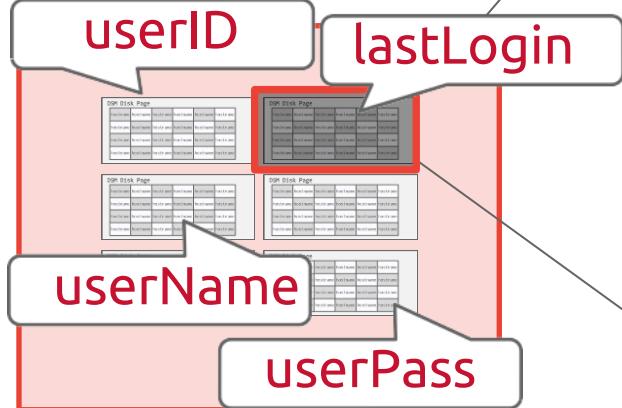
# DSM: OLAP example

```
SELECT COUNT(U.lastLogin),  
       EXTRACT(month FROM U.lastLogin) AS month  
  FROM useracct AS U  
 WHERE U.hostname LIKE '%.gov'  
 GROUP BY EXTRACT(month FROM U.lastLogin)
```



Disk

Database File



DSM Disk Page

header	lastlogin	lastlogin	lastlogin
lastlogin	lastlogin	lastlogin	lastlogin
lastlogin	lastlogin	lastlogin	lastlogin
lastlogin	lastlogin	lastlogin	lastlogin

# DSM: Variable-length Data

- ◆ Padding variable-length fields to ensure they are fixed-length is wasteful, especially for large attributes.
- ◆ A better approach is to use ***dictionary compression*** to convert repetitive variable-length data into fixed-length values (typically 32-bit integers).

# DSM: Summary

## ◆ Advantages

- ★ Reduces the amount wasted I/O per query because the DBMS only reads the data that it needs.
- ★ Faster query processing because of increased locality and cached data reuse.
- ★ Better data compression.

## ◆ Disadvantages

- ★ Slow for point queries, inserts, updates, and deletes because of tuple splitting/stitching/reorganization.

# OBSERVATION

- ◆ OLAP queries almost never access a single column in a table by itself.
  - ★ At some point during query execution, the DBMS must get other columns and stitch the original tuple back together.
- ◆ But we still need to store data in a columnar format to get the storage + execution benefits.
- ◆ We need columnar scheme that still stores attributes separately but keeps the data for each tuple physically close to each other.



# **Partition Attributes Across (PAX)**

# PAX Storage Model

- ◆ Partition Attributes Across (PAX) is a hybrid storage model that vertically partitions attributes within a database page.
  - ★ Examples: Parquet, ORC, and Arrow.
- ◆ The goal is to get the benefit of faster processing on columnar storage while retaining the spatial locality benefits of row storage.

# PAX: Physical Organization

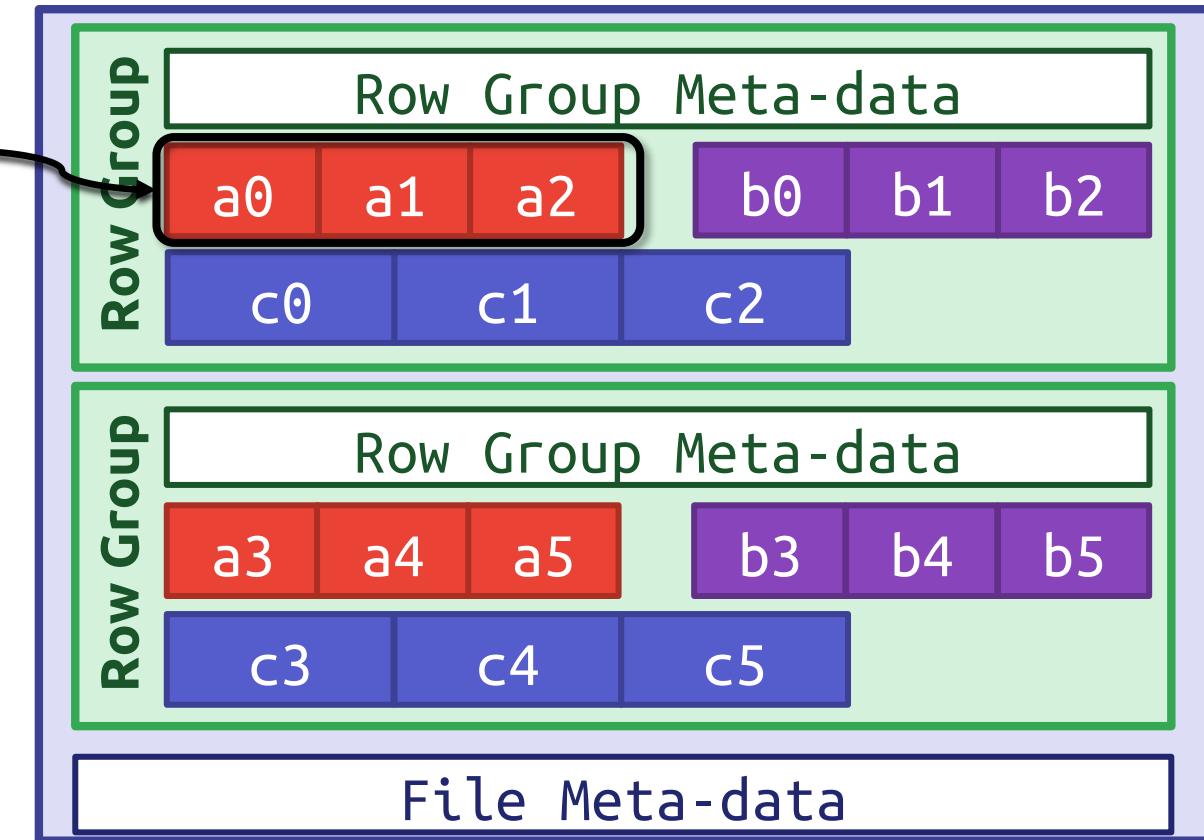
- ◆ Horizontally partition data into **row groups**.
- ◆ Then vertically partition their attributes into **column chunks**.
- ◆ Global meta-data directory contains offsets to the file's row groups.
  - ★ This is stored in the footer if the file is immutable (Parquet, Orc).
- ◆ Each row group contains its own meta-data header about its contents.

# PAX: Physical Organization

	Col A	Col B	Col C
Row#0	a0	b0	c0
Row#1	a1	b1	c1
Row#2	a2	b2	c2
Row#3	a3	b3	c3
Row#4	a4	b4	c4
Row#5	a5	b5	c5

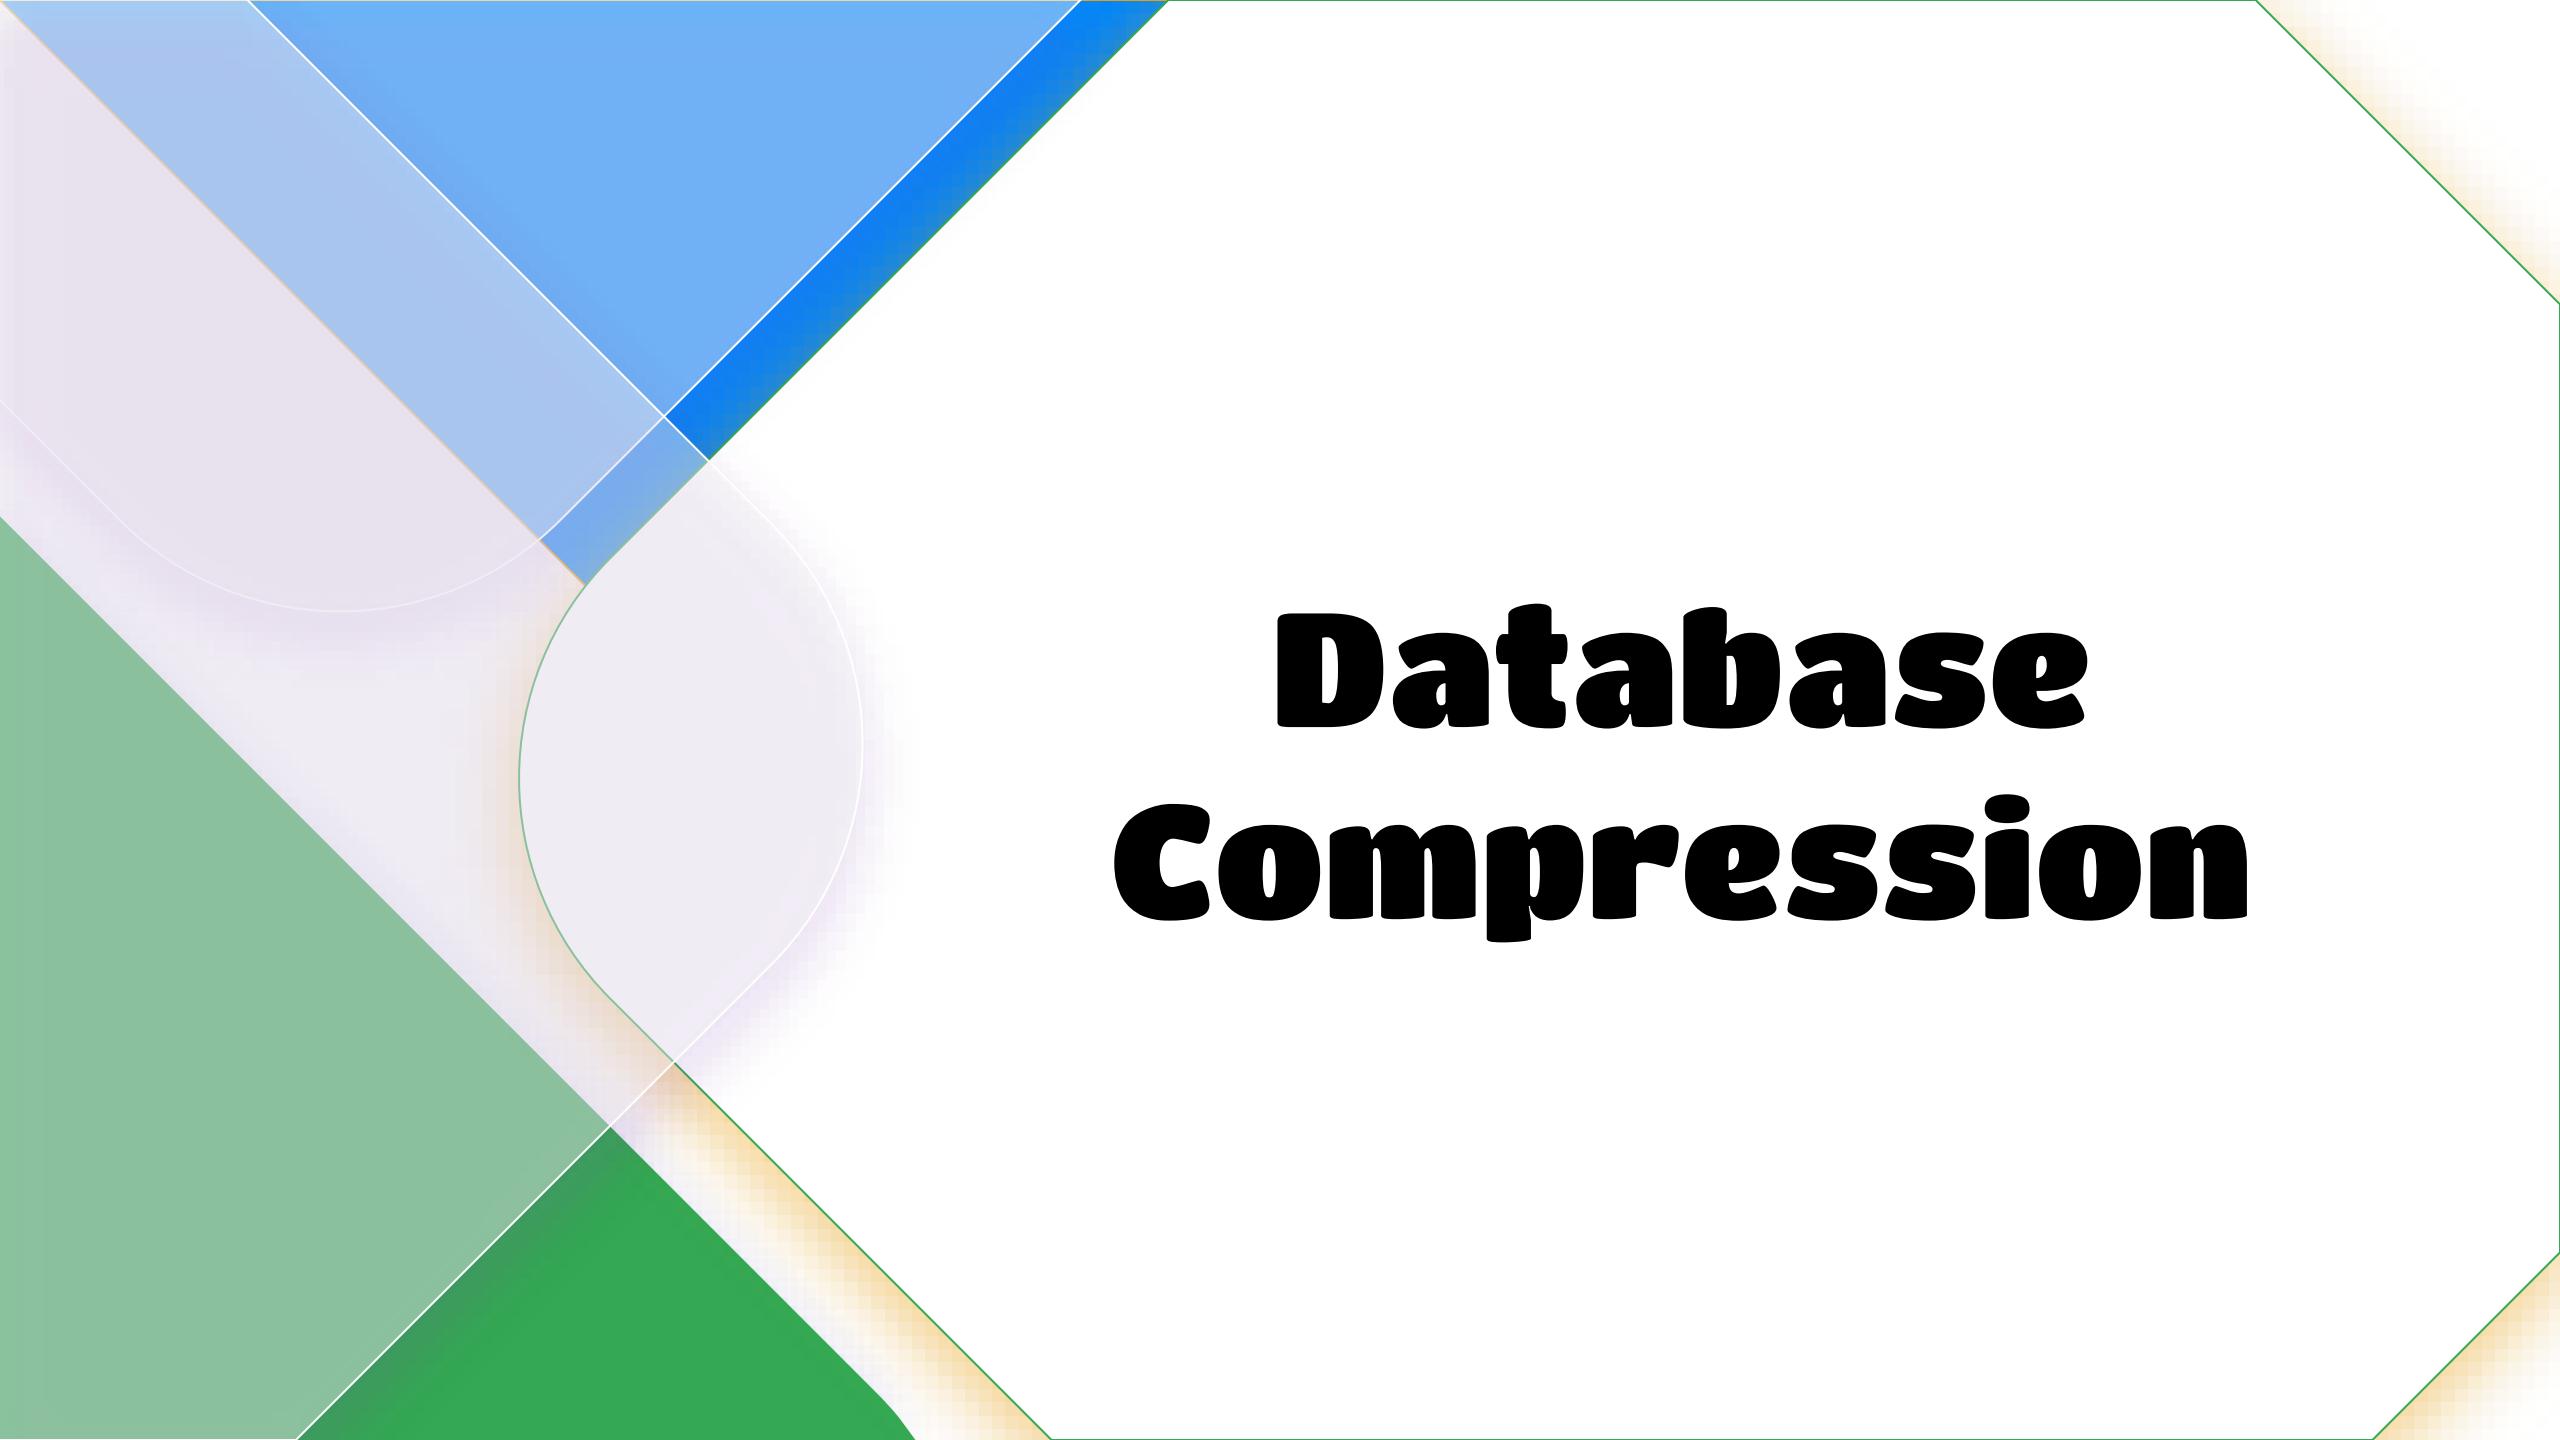
Column  
Chunk

PAX File



# Observation

- ◆ I/O is the main bottleneck if the DBMS fetches data from disk during query execution.
- ◆ The DBMS can **compress** pages to increase the utility of the data moved per I/O operation.
- ◆ Key trade-off is **speed** vs. **compression ratio**
  - ★ Compressing the database reduces DRAM requirements.
  - ★ It may decrease CPU costs during query execution.



# **Database Compression**

# Database Compression

- ◆ Goal #1: Must produce fixed-length values.
  - ★ Only exception is var-length data stored in separate pool.
- ◆ Goal #2: Postpone decompression for as long as possible during query execution.
  - ★ Also known as ***late materialization***.
- ◆ Goal #3: Must be a **lossless** scheme.
  - ★ People (typically) don't like losing data.
  - ★ Any **lossy** compression must be performed by application.

# Compression Granularity

- ◆ Block-level

- ★ Compress a block of tuples for the same table.

- ◆ Tuple-level

- ★ Compress the contents of the entire tuple (NSM-only).

- ◆ Attribute-level

- ★ Compress a single attribute within one tuple (overflow).

- ★ Can target multiple attributes for the same tuple.

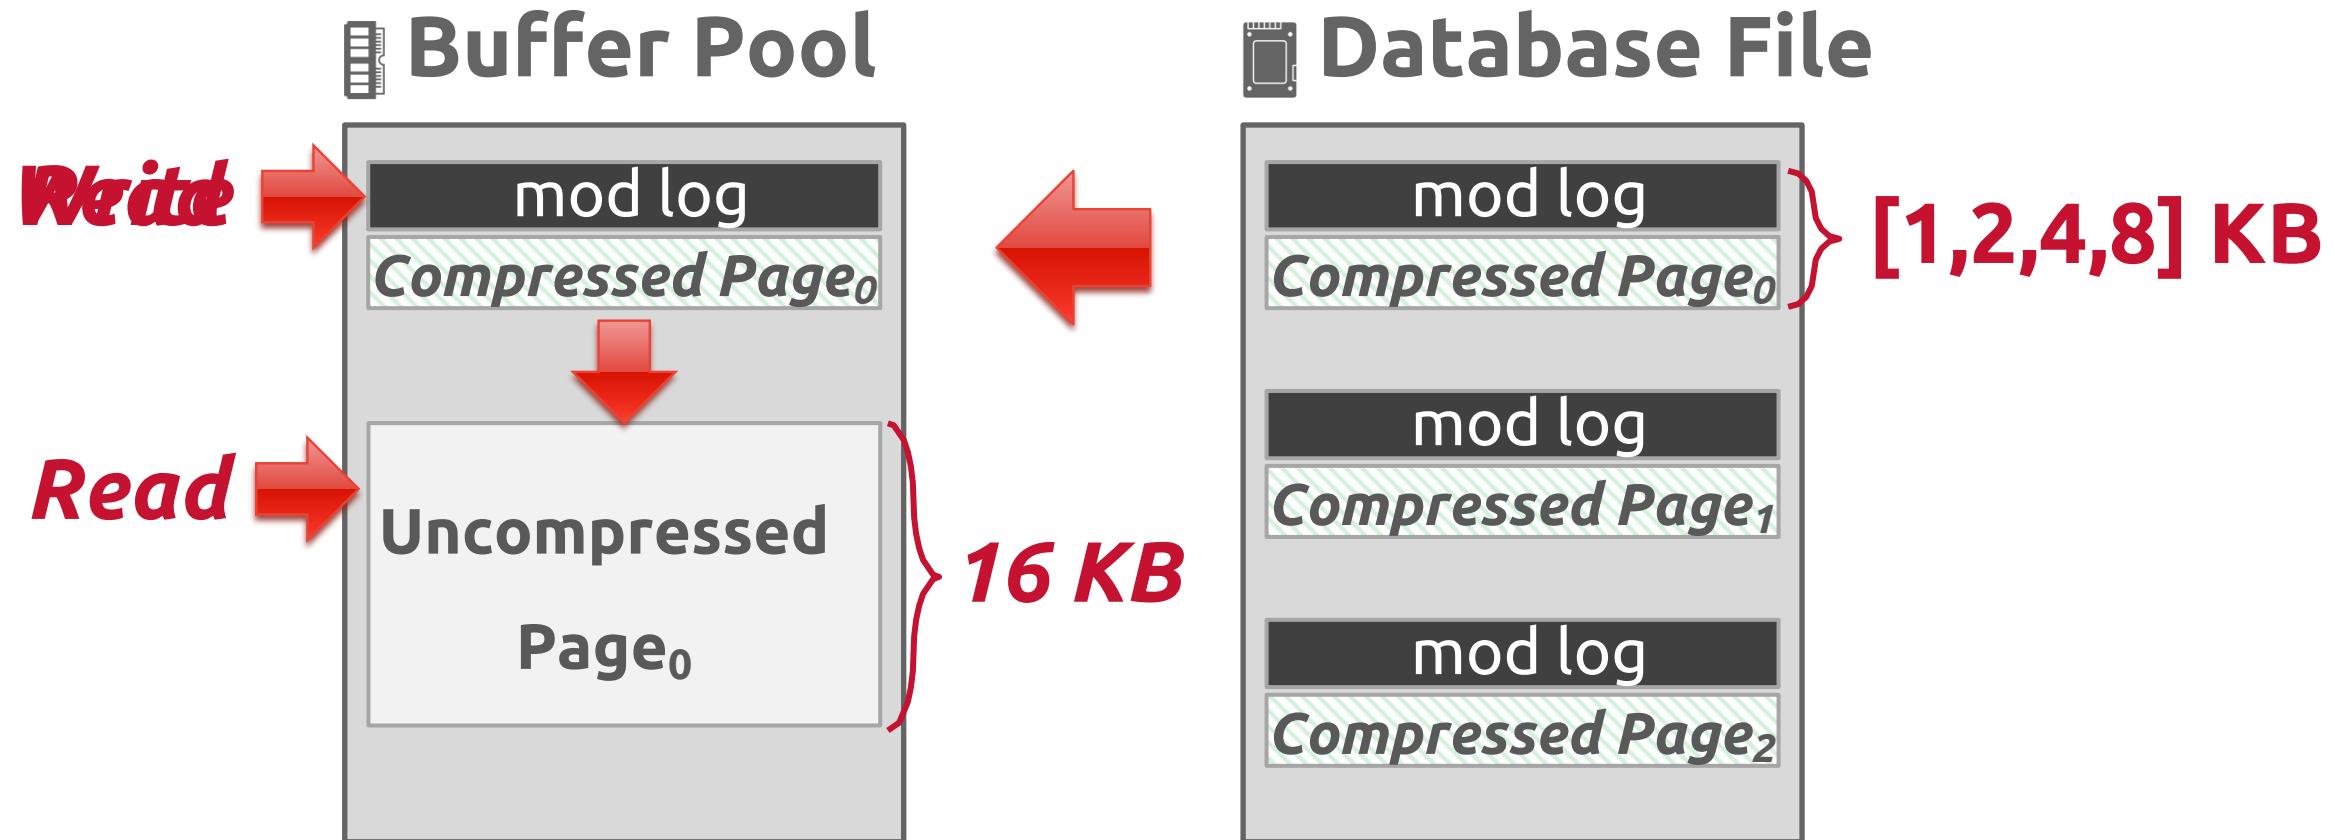
- ◆ Column-level

- ★ Compress multiple values for one or more attributes stored for multiple tuples (DSM-only).

# NAÏVE Compression

- ◆ Compress data using a general-purpose algorithm.
- ◆ Scope of compression is only based on the data provided as input.
  - ★ LZO (1996), LZ4 (2011), Snappy (2011), Oracle OZIP (2014), Zstd (2015)
- ◆ Considerations
  - ★ Computational overhead
  - ★ Compress vs. decompress speed.

# MySQL InnoDB Compression



# NAÏVE Compression

- ◆ The DBMS must decompress data first before it can be read and (potentially) modified.
  - ★ This limits the "scope" of the compression scheme.
- ◆ These schemes also do not consider the high-level meaning or semantics of the data.

# Observation

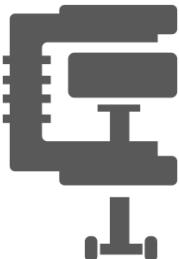
- ◆ Ideally, we want the DBMS to operate on compressed data without decompressing it first.

**SELECT \* FROM users  
WHERE name = 'Andy'**

*Database  
Magic!*

**SELECT \* FROM users  
WHERE name = XX**

NAME	SALARY
Andy	99999
Jignesh	88888



NAME	SALARY
XX	AA
YY	BB



# Compression Granularity

- ◆ Block-level

- ★ Compress a block of tuples for the same table.

- ◆ Tuple-level

- ★ Compress the contents of the entire tuple (NSM-only).

- ◆ Attribute-level

- ★ Compress a single attribute within one tuple (overflow).

- ★ Can target multiple attributes for the same tuple.

- ◆ Column-level

- ★ Compress multiple values for one or more attributes stored for multiple tuples (DSM-only).

# **Columnar Compression**

# Columnar Compression

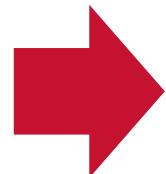
- ◆ Run-length Encoding
- ◆ Bit-Packing Encoding
- ◆ Bitmap Encoding
- ◆ Delta / Frame-of-Reference Encoding
- ◆ Incremental Encoding
- ◆ Dictionary Encoding

# Run-length encoding

- ◆ Compress runs of the same value in a single column into triplets:
  - ★ The value of the attribute.
  - ★ The start position in the column segment.
  - ★ The # of elements in the run.
- ◆ Requires the columns to be sorted intelligently to maximize compression opportunities.

# Run-length Encoding

```
SELECT isDead, COUNT(*)  
FROM users  
GROUP BY isDead
```



Compressed Data

id	isDead
1	(Y,0,3)
2	(N,3,1)
3	(Y,4,1)
4	(N,5,1)
6	(Y,6,2)
7	
8	
9	

*RLE Triplet*  
- Value  
- Offset  
- Length

Original Data

id	isDead
1	Y
2	Y
3	Y
4	N
6	Y
7	N
8	Y
9	Y

# Run-length Encoding

Sorted Data

id	isDead
1	Y
2	Y
3	Y
6	Y
8	Y
9	Y
4	N
7	N

Compressed Data

id	isDead
1	(Y, 0, 6)
2	(N, 7, 2)
3	
6	
8	
9	
4	
7	

# Bit Packing

- ◆ If the values for an integer attribute is **smaller** than the range of its given data type size, then reduce the number of bits to represent each value.
- ◆ Use bit-shifting tricks to operate on multiple values in a single word.

Original  
Data

int32
13
191
56
92
81
120
231
172

# Bit Packing

## Original Data

int32

13

00000000 00000000 00000000 00001101

191

00000000 00000000 00000000 10111111

56

00000000 00000000 00000000 00111000

92

00000000 00000000 00000000 01011100

81

00000000 00000000 00000000 01010001

120

00000000 00000000 00000000 01111000

231

00000000 00000000 00000000 11100111

172

00000000 00000000 00000000 10101100

Original:

$8 \times 32\text{-bits} = 256\text{ bits}$

## Original Data

int32

13

00001101

191

10111111

56

00111000

92

01011100

81

01010001

120

01111000

231

11100111

172

10101100

Compressed:

$8 \times 8\text{-bits} = 64\text{ bits}$

# Patching / mostly Encoding

- ◆ A variation of bit packing for when an attribute's values are "mostly" less than the largest size, store them with smaller data type.
  - ★ The remaining values that cannot be compressed are stored in their raw form.

# Patching / mostly Encoding

## Original Data

int32

13

191

999999999

92

81

120

231

172



## Compressed Data

mostly8

13

191

XXX

92

81

120

231

172

<i>offset</i>	<i>value</i>
3	99999999

Compressed:  
 $(8 \times 8\text{-bits}) + 16\text{-bits} + 32\text{-bits}$   
 $= 112 \text{ bits}$

Original:

$8 \times 32\text{-bits} = 256 \text{ bits}$

# Bitmap Encoding

- ◆ Store a separate bitmap for each unique value for an attribute where an offset in the vector corresponds to a tuple.
  - ★ The  $i^{\text{th}}$  position in the Bitmap corresponds to the  $i^{\text{th}}$  tuple in the table.
  - ★ Typically segmented into chunks to avoid allocating large blocks of contiguous memory.
- ◆ Only practical if the value cardinality is low.
- ◆ Some DBMSs provide **bitmap indexes**.

# Bitmap encoding

## Original Data

id	isDead
1	Y
2	Y
3	Y
4	N
6	Y
7	N
8	Y
9	Y

Original:  
 $8 \times 8\text{-bits} = 64 \text{ bits}$



## Compressed Data

id	isDead	Y	N
1		1	0
2		1	0
3		1	0
4		0	1
6		1	0
7		0	1
8		1	0
9		1	0

Compressed:  
 $16\text{-bits} + 16\text{-bits} = 32\text{-bits}$

$2 \times 8\text{-bits}$   
 $= 16\text{-bits}$

$8 \times 2\text{-bits}$   
 $= 16\text{-bits}$

# Bitmap encoding: Example

- ◆ Assume we have 10 million tuples.  
43,000 zip codes in the US.
  - ★  $10000000 \times 32\text{-bits} = 40 \text{ MB}$
  - ★  $10000000 \times 43000 = 53.75 \text{ GB}$
- ◆ Every time the application inserts a new tuple, the DBMS must extend 43,000 different bitmaps.
- ◆ There are compressed data structures for sparse data sets:
  - ★ Roaring Bitmaps

```
CREATE TABLE customer (
    id INT PRIMARY KEY,
    name VARCHAR(32),
    email VARCHAR(64),
    address VARCHAR(64),
    zip_code INT
);
```



# Delta Encoding

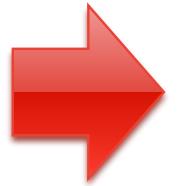
- ◆ Recording the difference between values that follow each other in the same column.
  - ★ Store base value in-line or in a separate look-up table.
  - ★ Combine with RLE to get even better compression ratios.
- ◆ Frame-of-Reference Variant: Use global min value.

# Delta Encoding

Original Data	
time64	temp
12:00	99.5
12:01	99.4
12:02	99.5
12:03	99.6
12:04	99.4



Compressed Data	
time64	temp
12:00	99.5
+1	-0.1
+1	+0.1
+1	+0.1
+1	-0.2



Compressed Data	
temp	time64
99.5	12:00
-0.1	(+1, 4)
+0.1	
+0.1	
-0.2	

**5 × 64bits  
= 320 bits**

**64bits + (4 × 16bits)  
= 128 bits**

**64bits + (2 × 16bits)  
= 96 bits**

# Dictionary Compression

- ◆ Replace frequent values with smaller fixed-length codes and then maintain a mapping (dictionary) from the codes to the original values
  - ★ Typically, one code per attribute value.
  - ★ Most widely used native compression scheme in DBMSs.
- ◆ The ideal dictionary scheme supports fast encoding and decoding for both point and range queries.
  - ★ **Encode/Locate:** For a given uncompressed value, convert it into its compressed form.
  - ★ **Decode/Extract:** For a given compressed value, convert it back into its original form.

# Dictionary: Order-preserving

- ◆ The encoded values need to support the same collation as the original values.

```
SELECT * FROM users  
WHERE name LIKE 'And%'
```



```
SELECT * FROM users  
WHERE name BETWEEN 10 AND 20
```

*Original Data*

name
Andrea
Mr. Pickles
Andy
Jignesh
Mr. Pickles



*Compressed Data*

name
10
40
20
30
40

value	code
Andrea	10
Andy	20
Jignesh	30
Mr. Pickles	40

*Sorted Dictionary*

# Order-preserving encoding

SELECT name FROM users

WHERE name LIKE 'And%'

SELECT DISTINCT name FROM users

WHERE name LIKE 'And%'



Still must perform  
scan on column



Only need to  
access dictionary

*Original Data*

name
Andrea
Mr. Pickles
Andy
Jignesh
Mr. Pickles



*Compressed Data*

name
10
40
20
30
40

value	code
Andrea	10
Andy	20
Jignesh	30
Mr. Pickles	40

*Sorted Dictionary*

# Conclusion

- ◆ It is important to choose the right storage model for the target workload:
  - ★ OLTP → Row Store
  - ★ OLAP → Column Store
- ◆ DBMSs can combine different approaches for even better compression.
- ◆ Dictionary encoding is probably the most useful scheme because it does not require pre-sorting.



THANK  
YOU ☺

