

**CSAI 302**

# **Advanced Database Systems**

**Lec 04**

**Indexing Structures and  
Implementation**

# Index

- ◆ A data structure that improves the speed of data retrieval operations on a database table.
  - ◆ Built on **one or more columns** of a table and store a sorted copy of the indexed data along with pointers to the corresponding rows in the main table.
- ★ Example: B+Tree

# Types of indexes

## Sparse

- ◆ One entry per data block
- ◆ Requires data to be sorted
- ◆ Identifies the first record of the block
- ★ Faster access

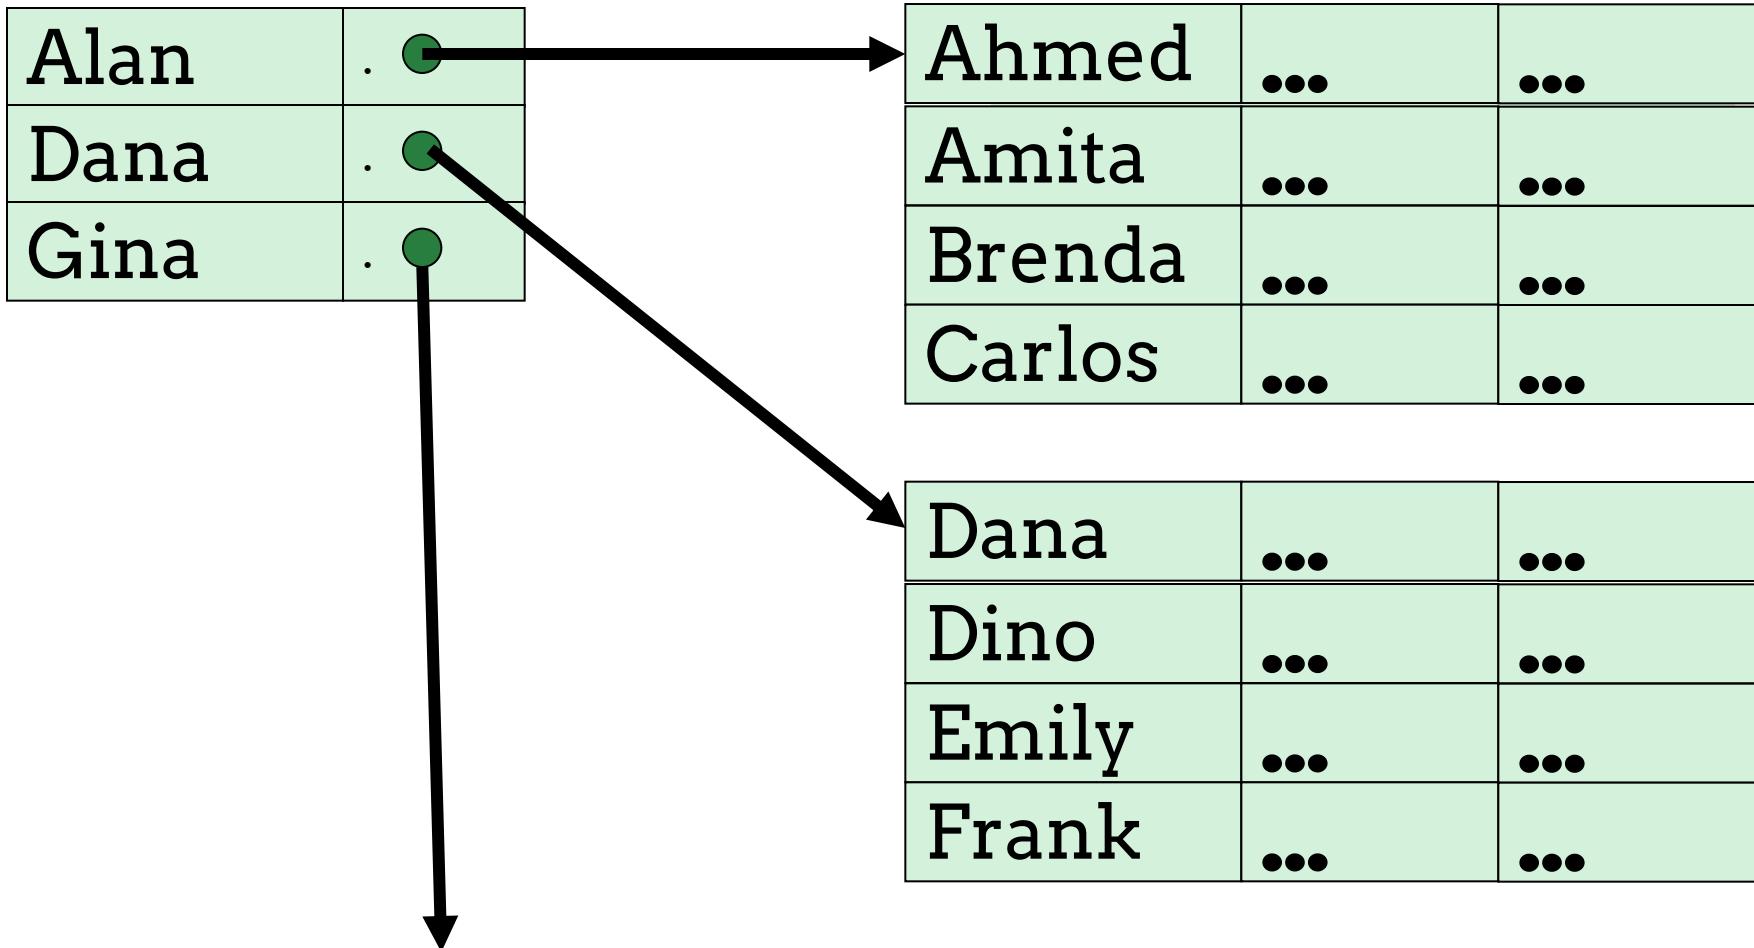
## Dense

- ◆ One entry per record
- ◆ Data do not have to be sorted
- ◆ Can tell if a given record exists without accessing the file

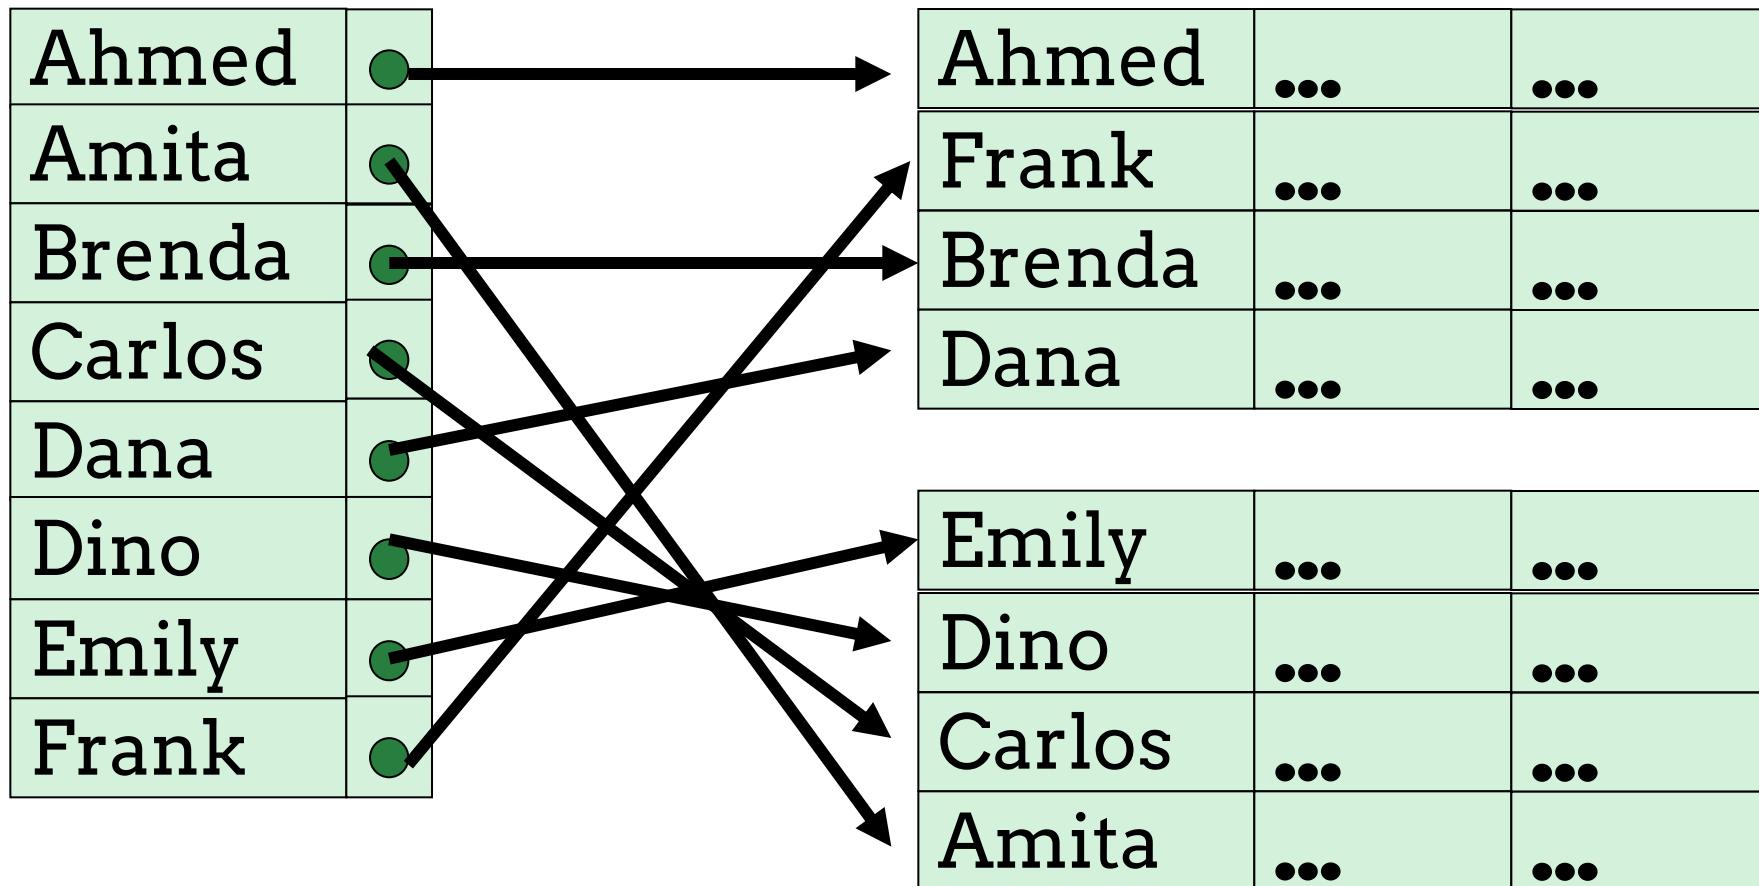
# Indexes based on primary keys

- ◆ Each key value corresponds to a specific record
- ◆ Two cases to consider:
  - ★ Table is sorted on its primary key
    - Can use a sparse index
  - ★ Table is either non-sorted or sorted on another field
    - Must use a dense index

# Sparse Index



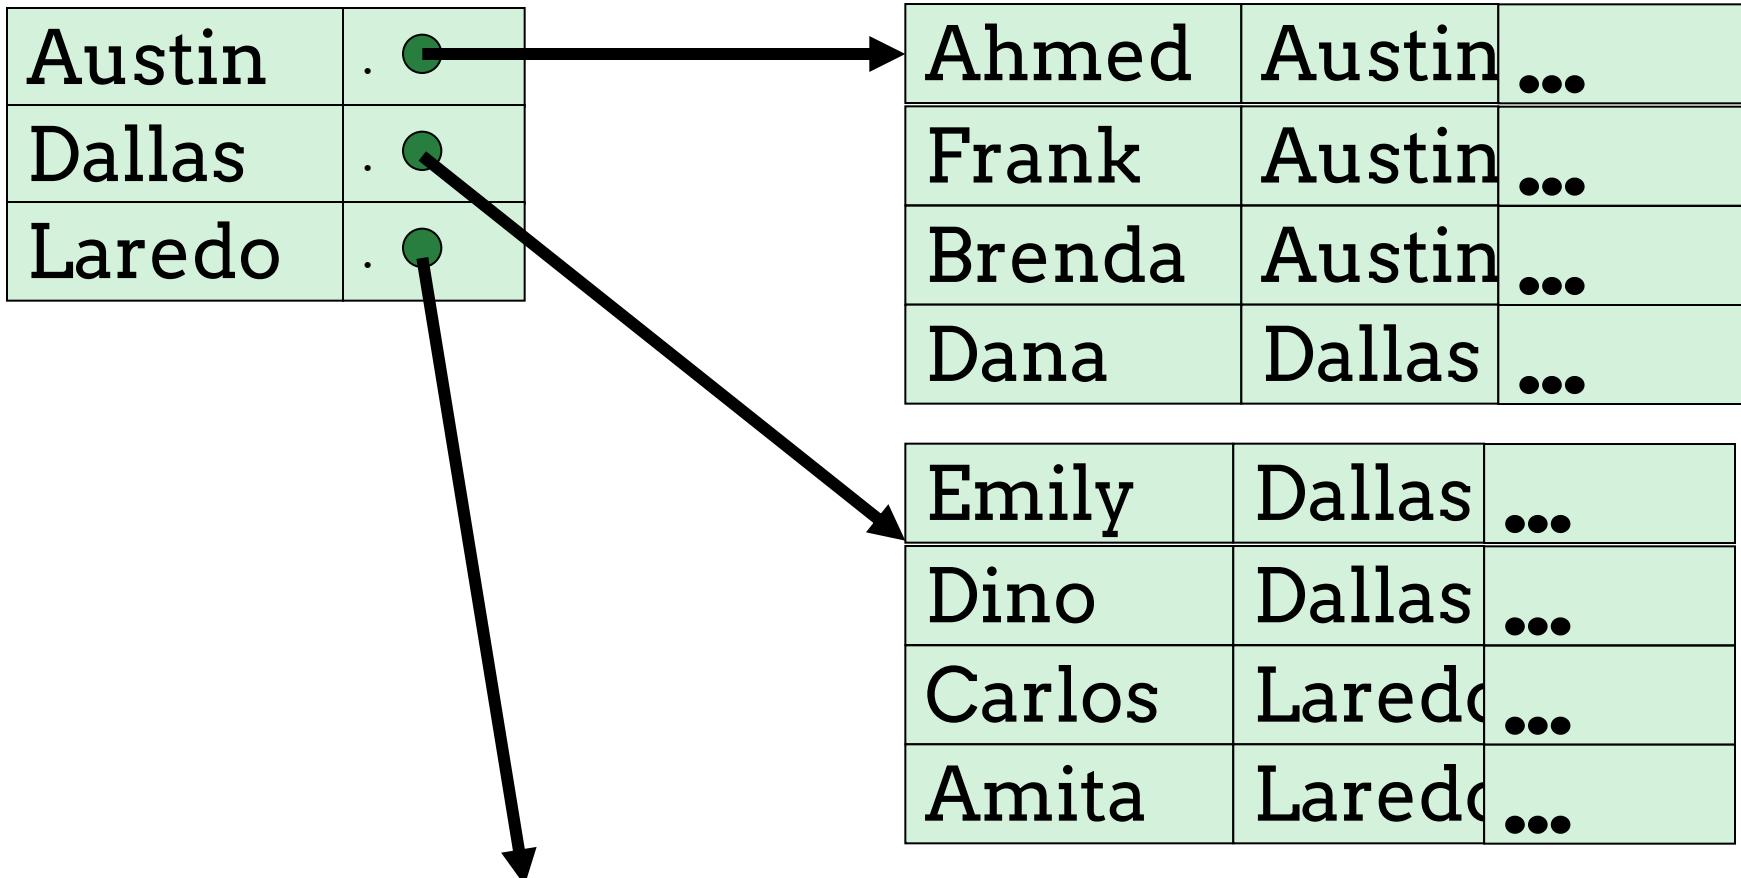
# Dense Index



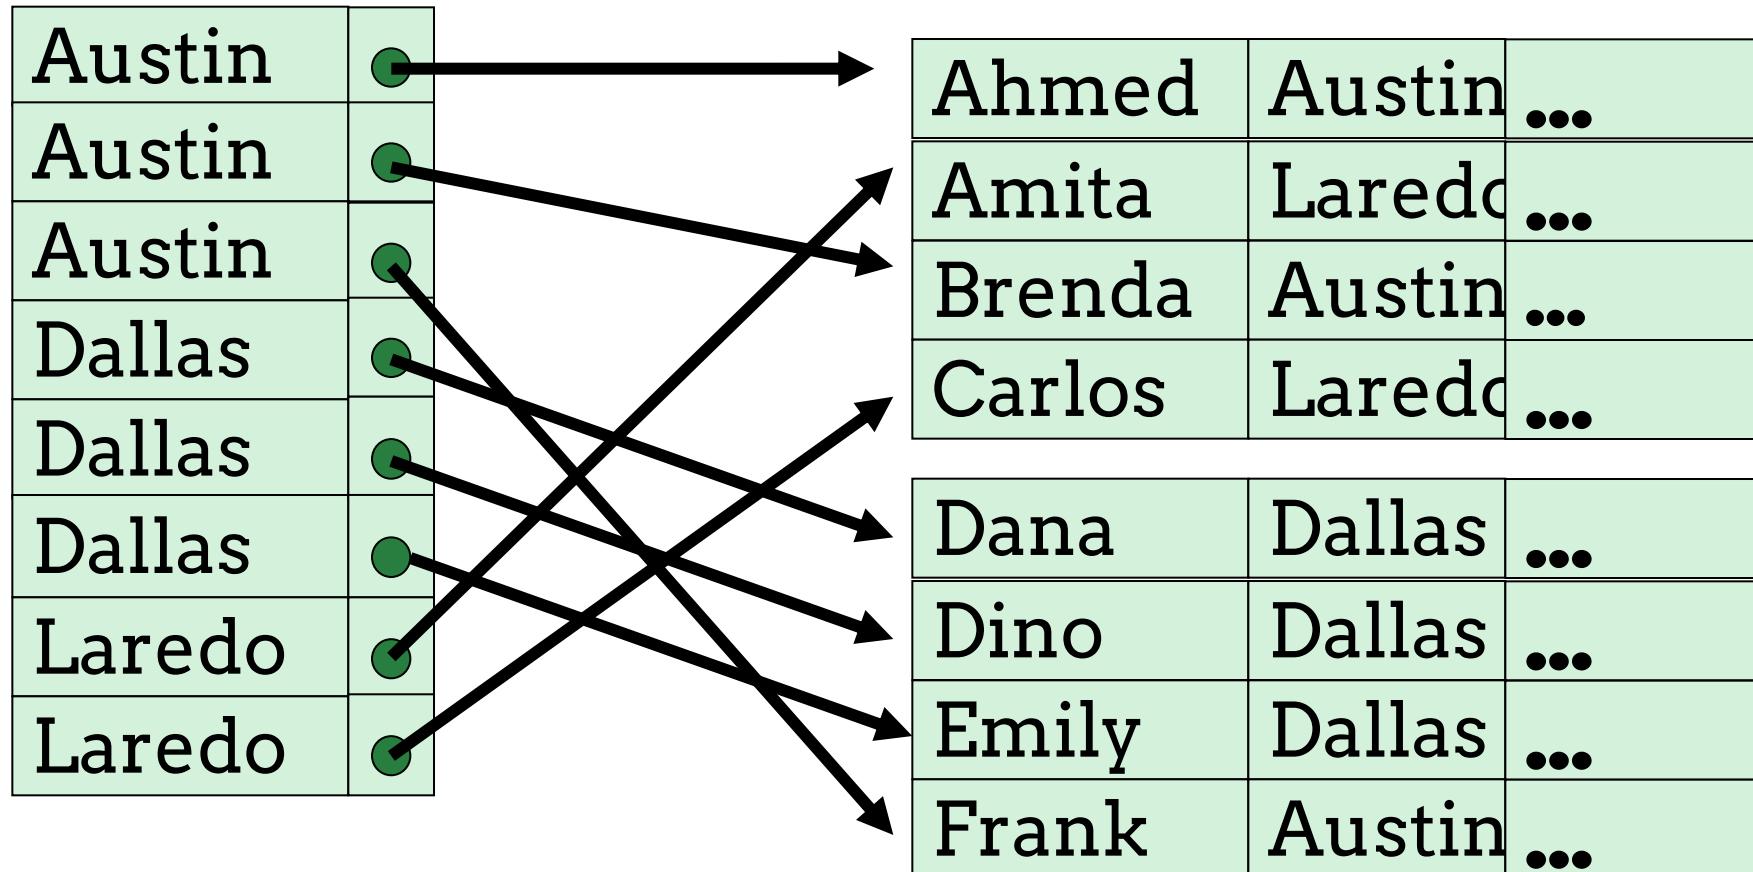
# Indexes based on other fields

- ◆ Each key value may correspond to more than one record
  - ★ ***clustering index***
- ◆ Two cases to consider:
  - ★ Table is sorted on the field
    - Can use a sparse index
  - ★ Table is either non-sorted or sorted on another field
    - Must use a dense index

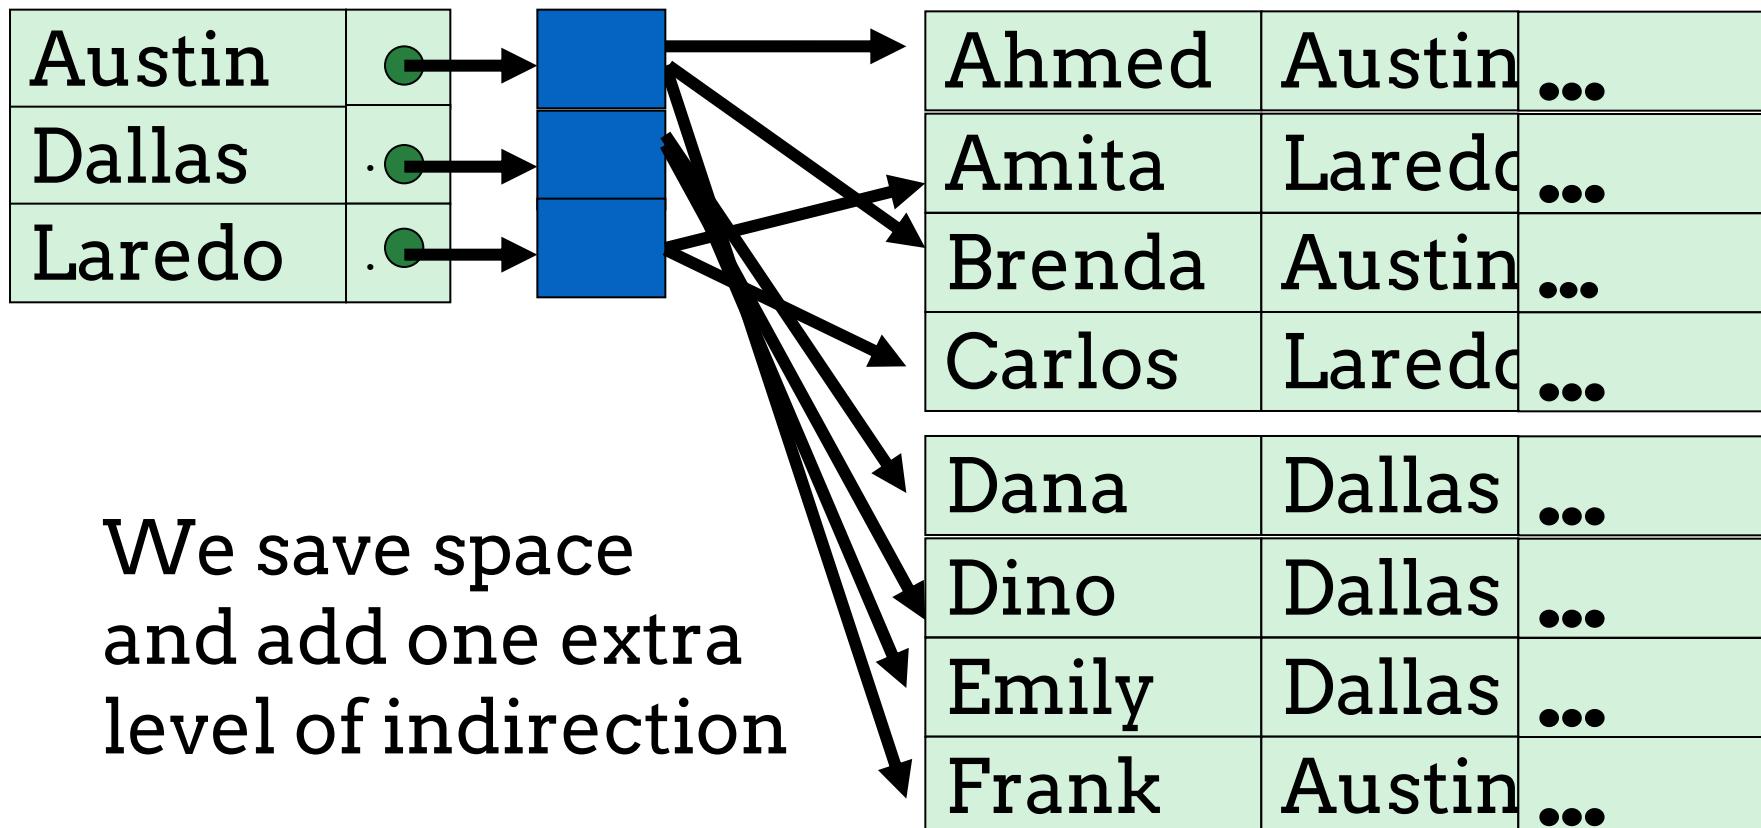
# Sparse clustering index



# Dense clustering index



# Another realization





# **B-trees and B+ trees**

# Motivation

- ◆ To have dynamic indexing structures that can evolve when records are **added** and **deleted**
  - ★ Static indexes are completely rebuilt
- ◆ Optimized for searches on ***block devices***
- ◆ Both B trees and B+ trees are ***not binary***
  - ★ Objective is to increase ***branching factor*** to reduce the number of device accesses

# B-Tree family

- ◆ B-Tree (1970)
- ◆ B+Tree (1973)
- ◆ B\* Tree (1977)
- ◆ B<sup>link</sup>Tree (1981)
- ◆ B<sup>ε</sup>-Tree (2003)
- ◆ B<sup>w</sup>-Tree (2013)

# Binary vs. higher-order tree

## ◆ **Binary trees:**

- ★ Designed for in-memory searches
- ★ Try to minimize the number of memory accesses

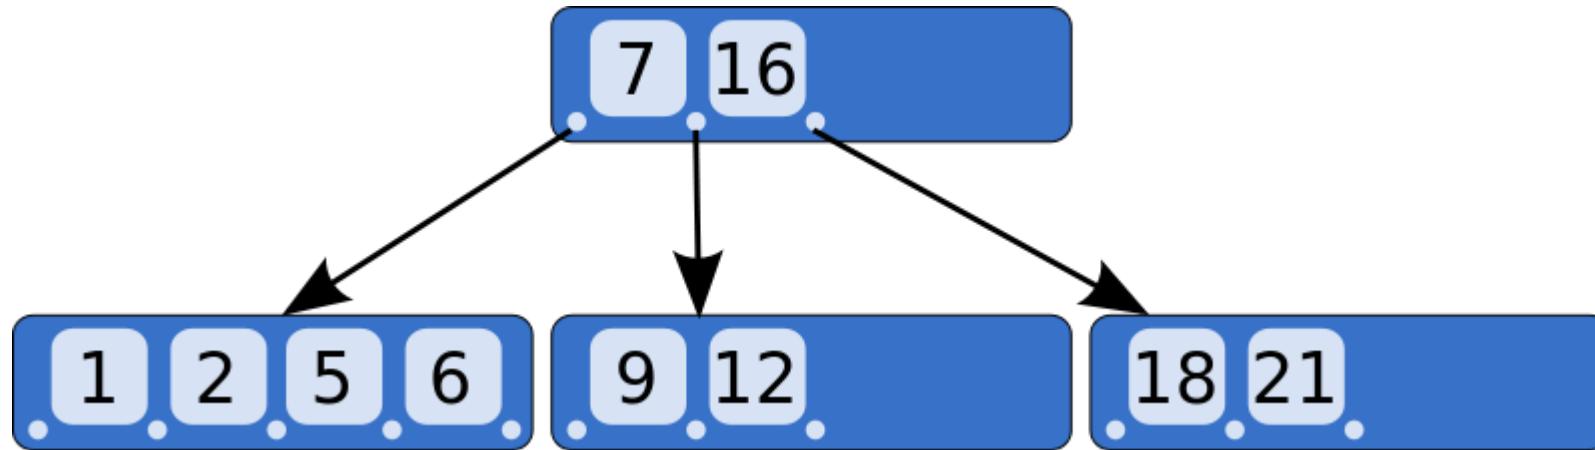
## ◆ **Higher-order trees:**

- ★ Designed for searching data on block devices
- ★ Try to minimize the number of device accesses
  - Searching within a block is cheap!

# B trees

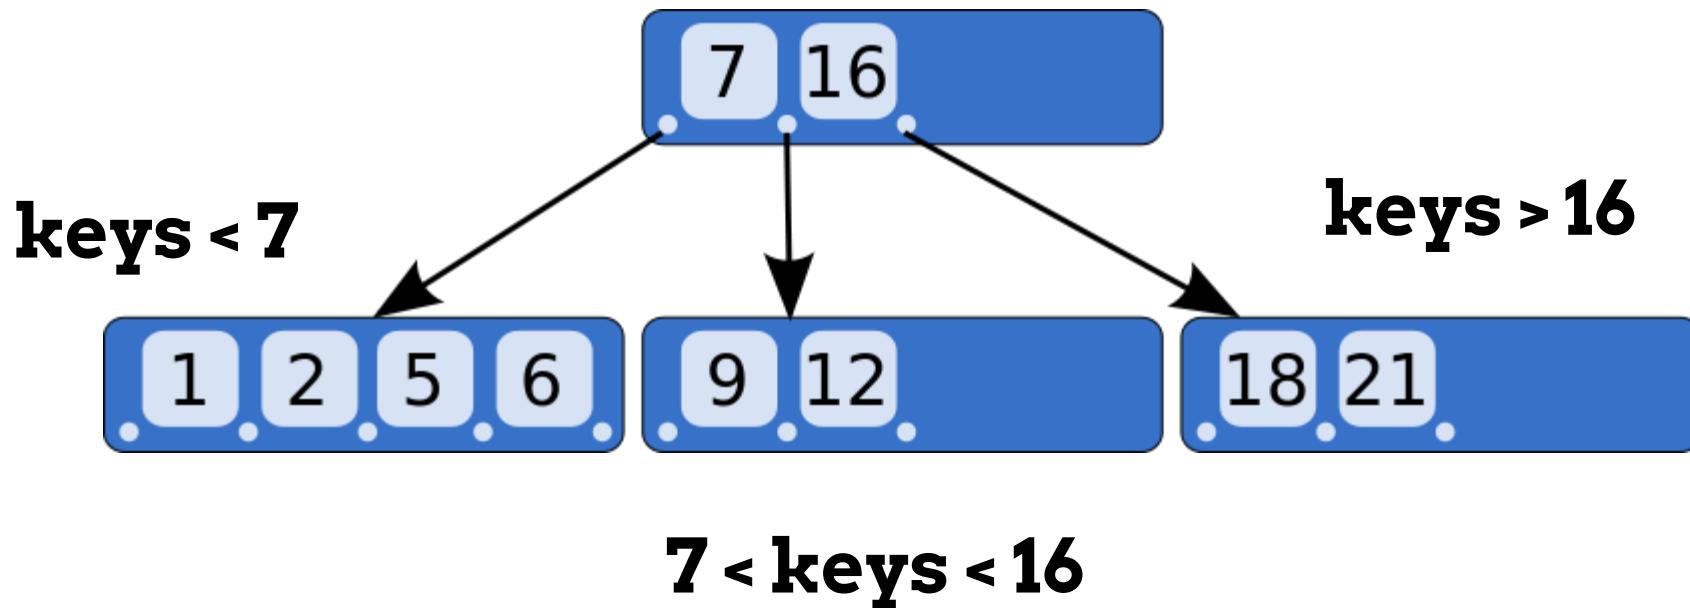
- ◆ Generalization of binary search trees
  - ★ Not binary trees
  - ★ The B stands for Bayer (or Boeing)
- ◆ Designed for searching data stored on block-oriented devices

# A very small B tree



- ◆ Bottom nodes are leaf nodes: all their pointers are NULL

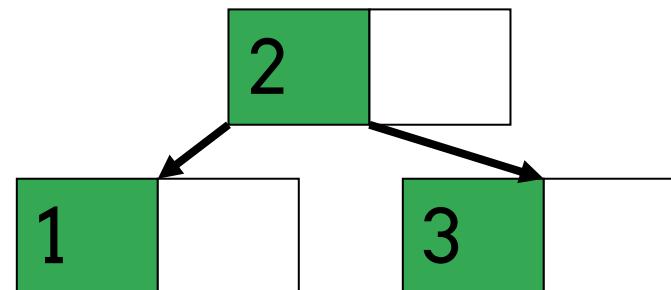
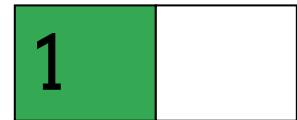
# Searching the tree



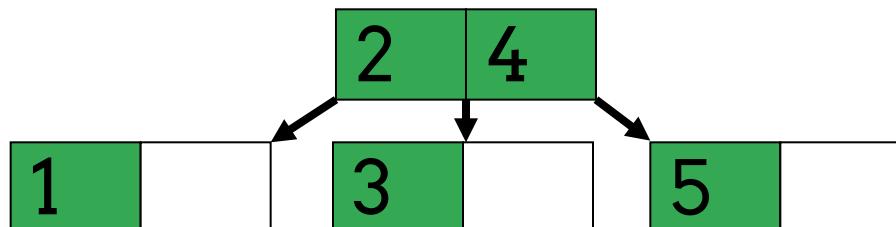
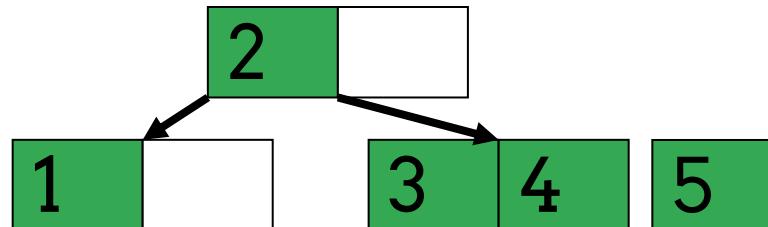
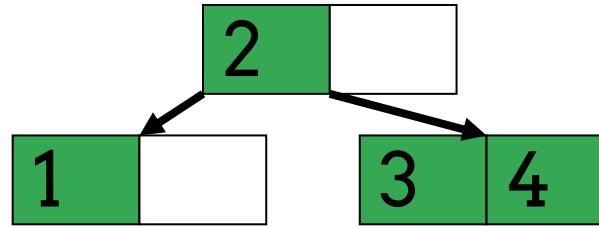
# Balancing B trees

- ◆ Objective is to ensure that all terminals nodes be at the same depth

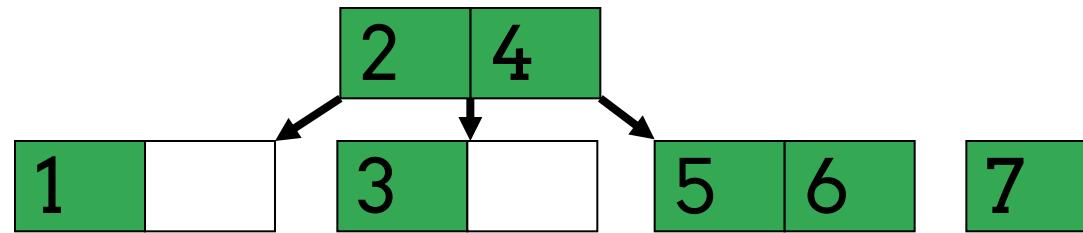
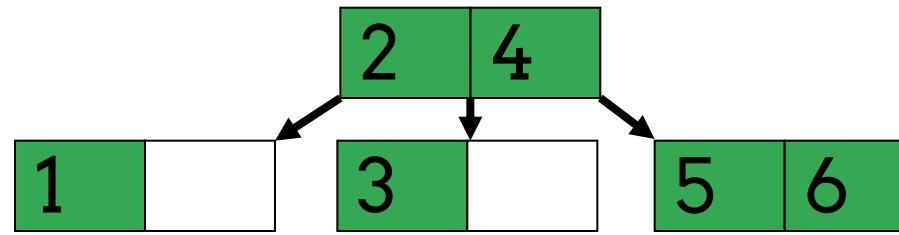
# Insertions



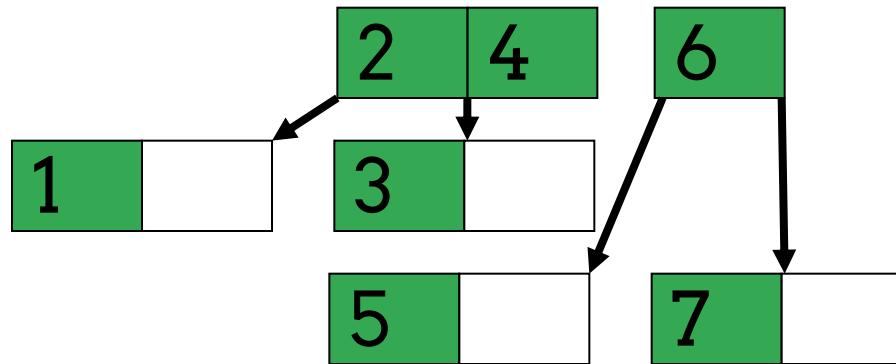
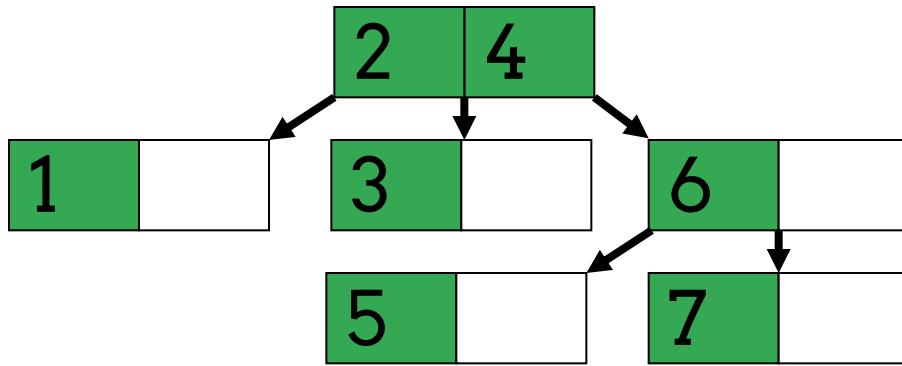
# Insertions



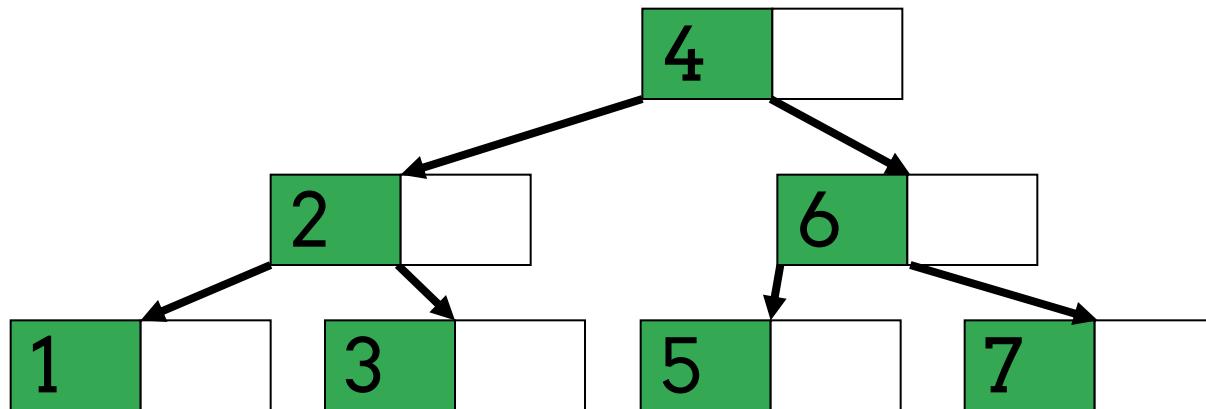
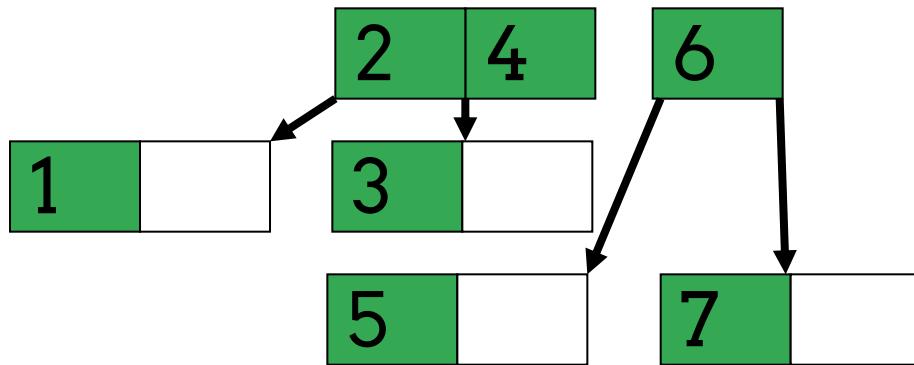
# Insertions



# Insertions



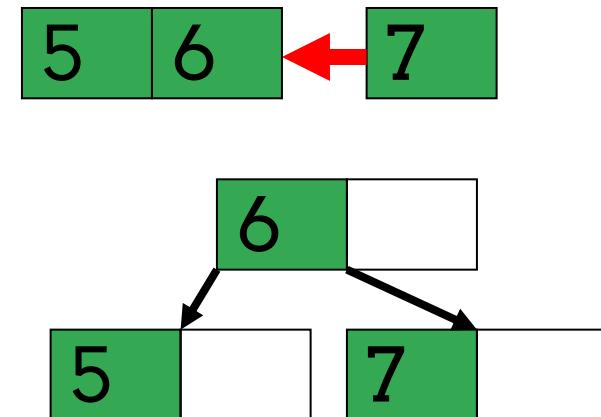
# Insertions



# Two basic operations

## ◆ Split:

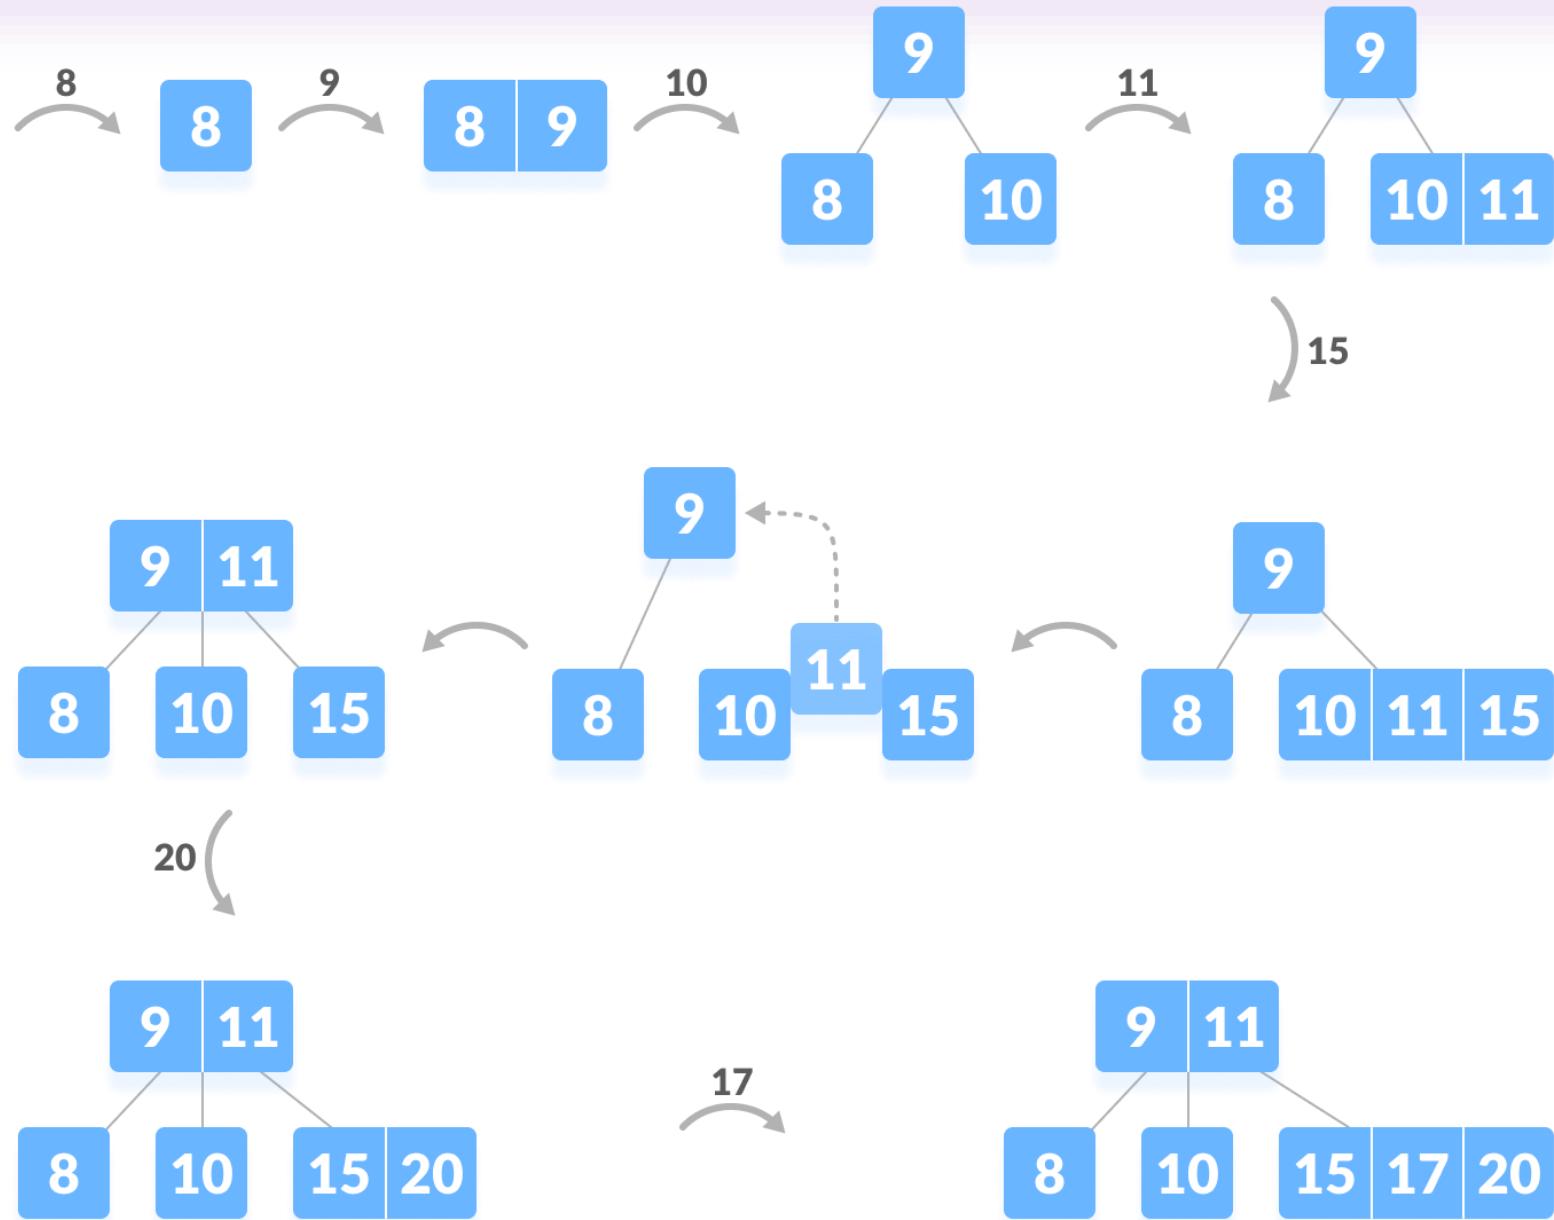
- ★ When trying to add to a full node
- ★ Split node at central value



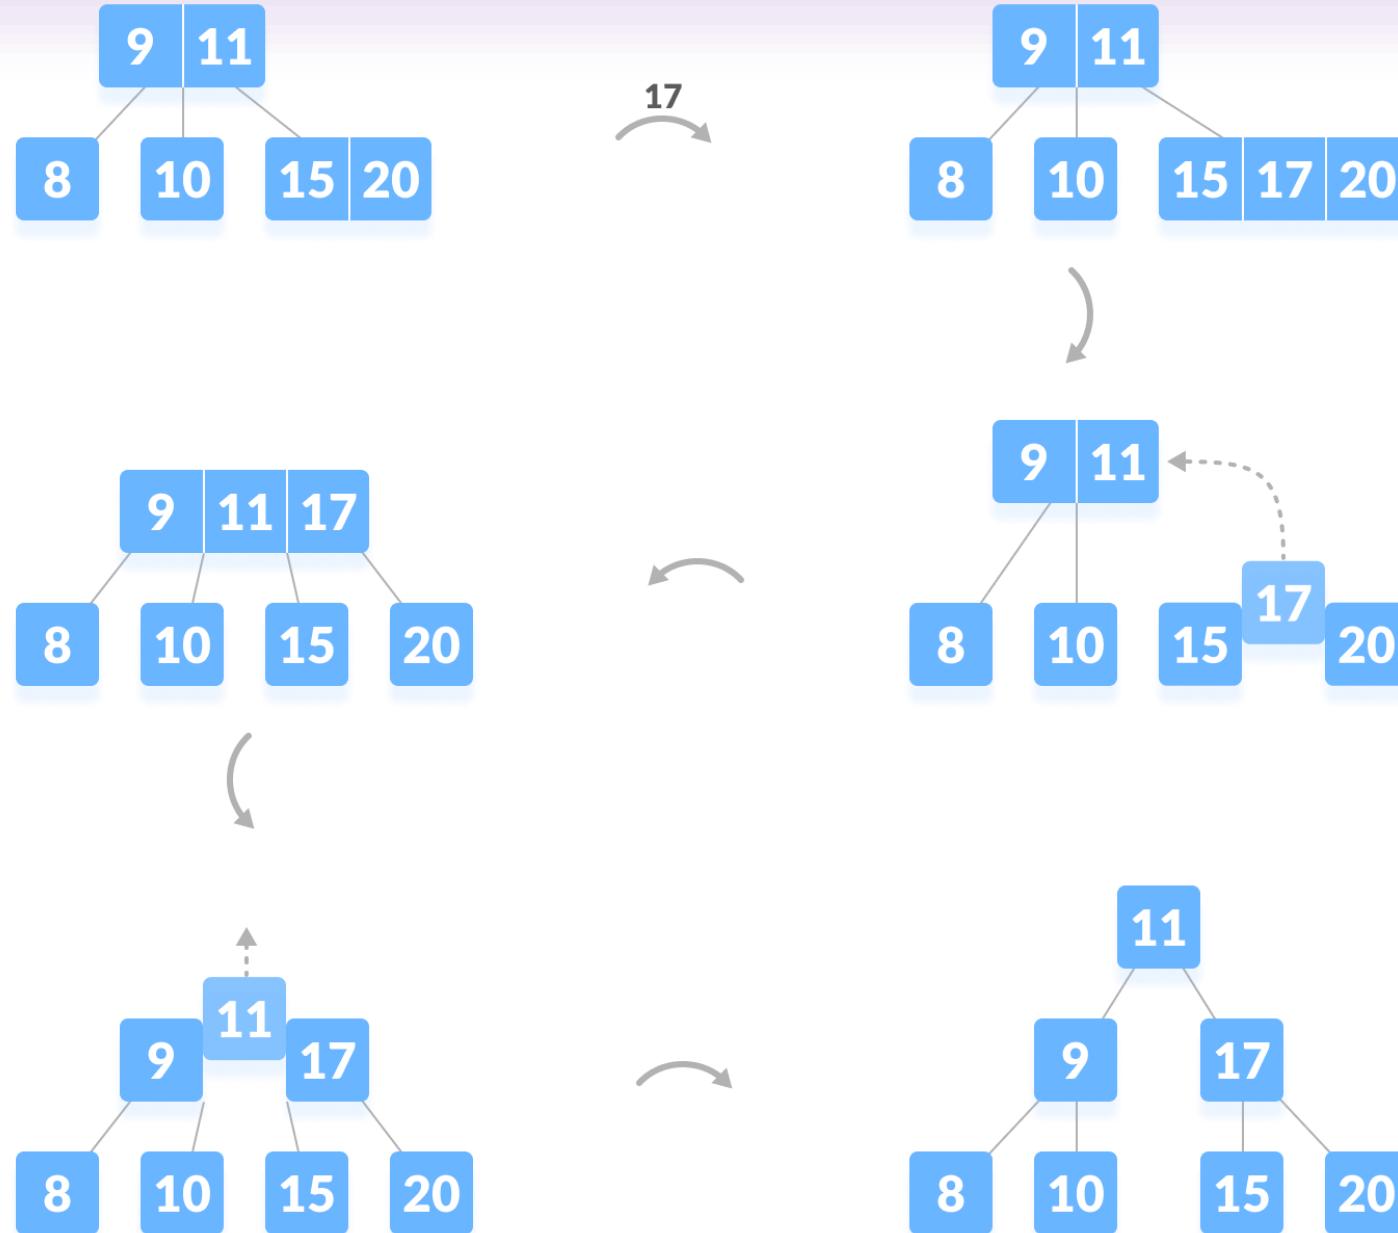
## ◆ Promote:

- ★ Must insert root of split node higher up
- ★ May require a new split

# Insertion in B-Trees



# Insertion in B-Trees



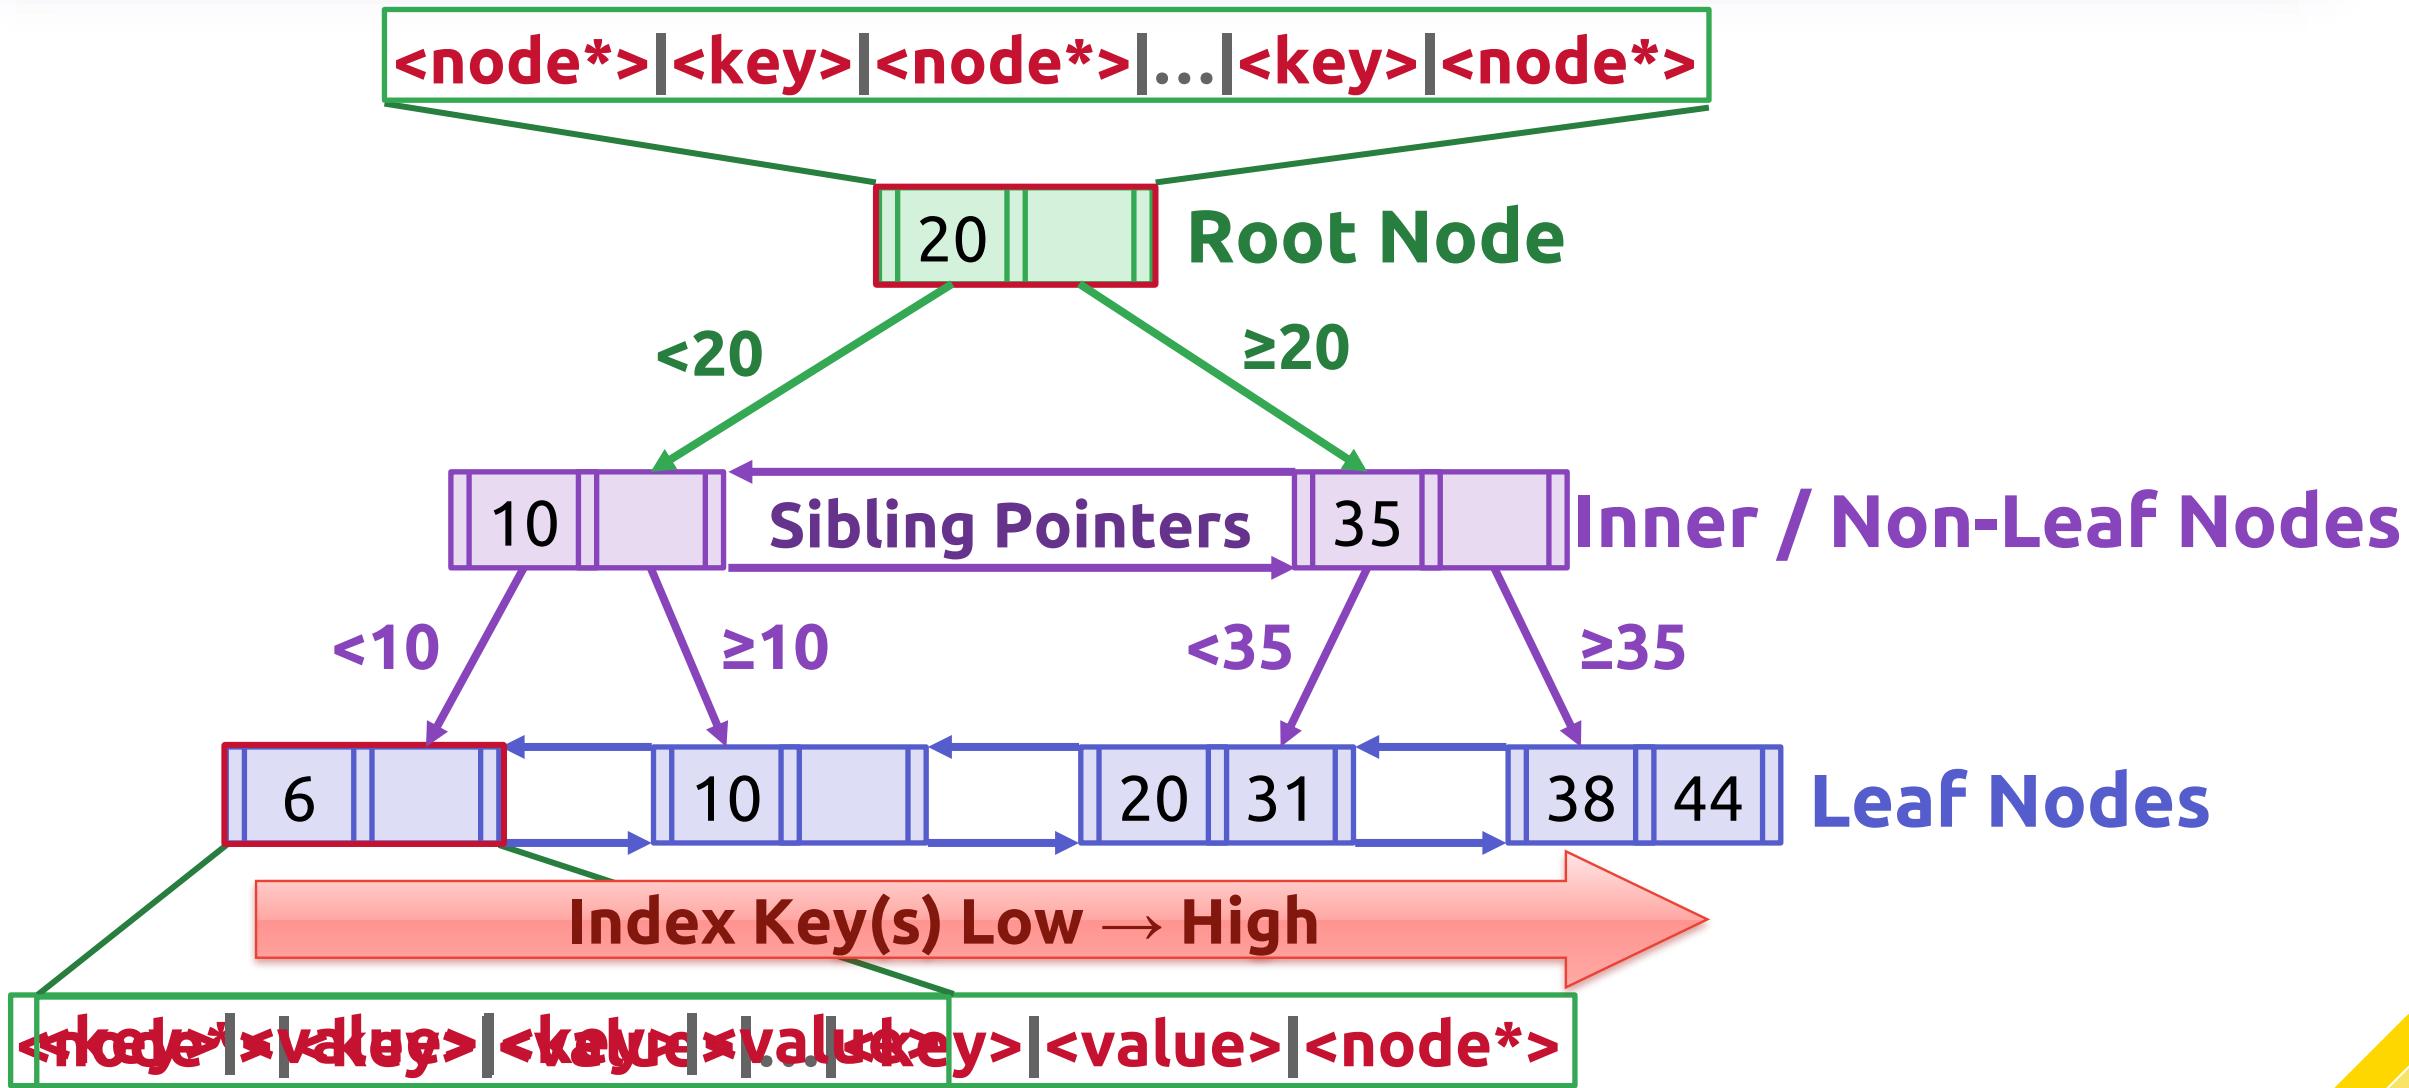
# B+ trees

- ◆ Variant of B trees
- ◆ Two types of nodes
  - ★ Internal nodes have no data pointers
  - ★ Leaf nodes have no in-tree pointers

# B+Tree

- ◆ A self-balancing, ordered  $m$ -way tree
- ◆ for searches, sequential access, insertions, and deletions in  $O(\log_m n)$  where  $m$  is the tree fanout.
  - ★ It is perfectly balanced (i.e., every leaf node is at the same depth in the tree)
  - ★ Every node other than the root is at least half-full
    - $m/2-1 \leq \#keys \leq m-1$
  - ★ Every inner node with  $k$  keys has  $k+1$  non-null children.
  - ★ Optimized for reading/writing large data blocks.

# B+Tree Example



# Nodes

- ◆ Every B+Tree node is comprised of an array of key/value pairs.
  - ★ The keys are derived from the index's target attribute(s).
  - ★ The values will differ based on whether the node is classified as an **inner node** or a **leaf node**.
- ◆ The arrays are (usually) kept in sorted key order.
- ◆ Store all **NULL** keys at either first or last leaf nodes.



# Searches

```
def tree_search(k, node):
    if node is a leaf:
        return node
    elif k < k_0:
        return tree_search(k, p_0)
    ...
    elif k_i ≤ k < k_{i+1}
        return tree_search(k, p_{i+1})
    ...
    elif k_d ≤ k
        return tree_search(k, p_{d+1});
```

# Insertion in B+ Trees

- ◆ 1. Navigate to the correct leaf node
- ◆ 2. Insert the new key in sorted order
- ◆ 3. If overflow occurs:
  - ★ Split the leaf node
  - ★ Push the middle key to the parent
  - ★ This process may propagate up to the root

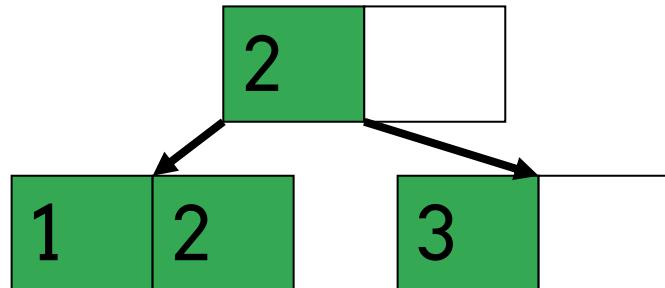
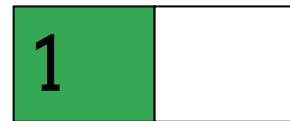
# Insertions

- ◆ **def insert (entry) :**
  - ★ Find target leaf  $L$
  - ★ **if**  $L$  has less than  $m - 2$  entries :
    - add the entry
  - else** :
    - Allocate new leaf  $L'$
    - Pick the  $m/2$  highest keys of  $L$  and move them to  $L'$
    - Insert **highest key** of  $L$  and corresponding address leaf into the parent node
    - If the parent is full :
      - Split it and add the middle key to its parent node
    - Repeat until a parent is found that is not full

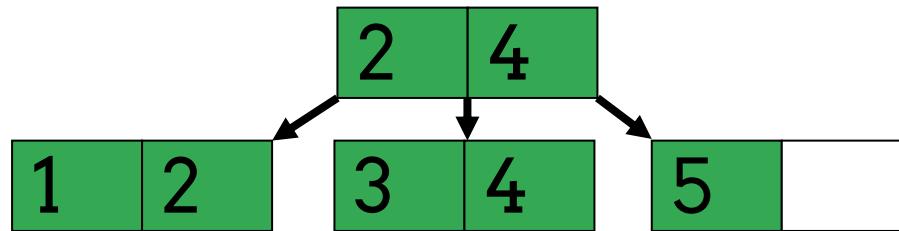
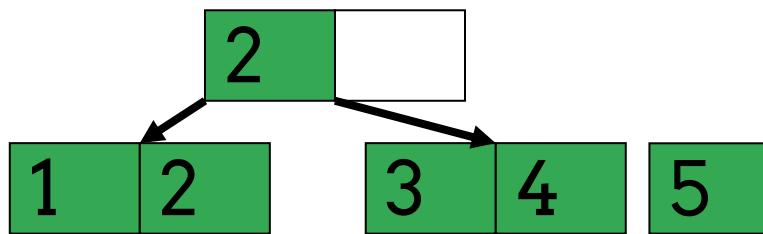
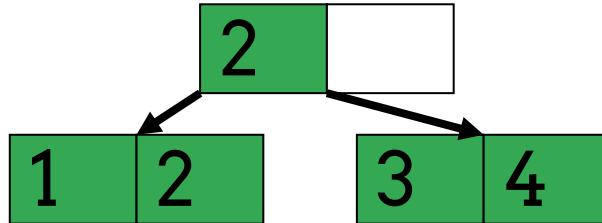
# B+TREE – INSERT

- ◆ Find correct leaf node  $L$ .
- ◆ Insert data entry into  $L$  in sorted order.
- ◆ If  $L$  has enough space, done!
- ◆ Otherwise, split  $L$  keys into  $L$  and a new node  $L_2$ 
  - ★ Redistribute entries evenly, copy up middle key.
  - ★ Insert index entry pointing to  $L_2$  into parent of  $L$ .
- ◆ To split inner node, redistribute entries evenly, but push up middle key.

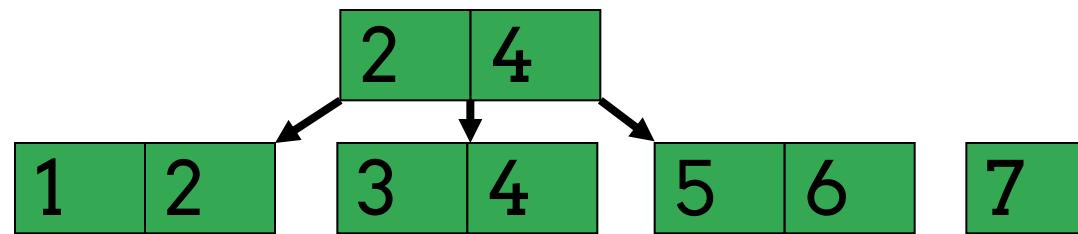
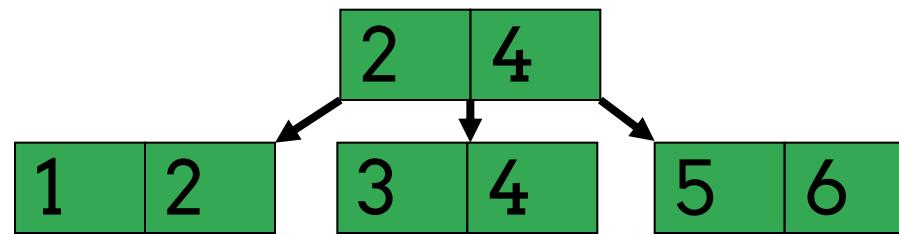
# Insertions



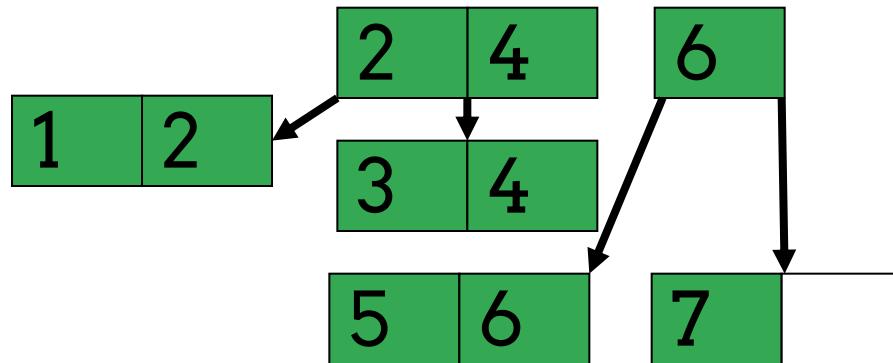
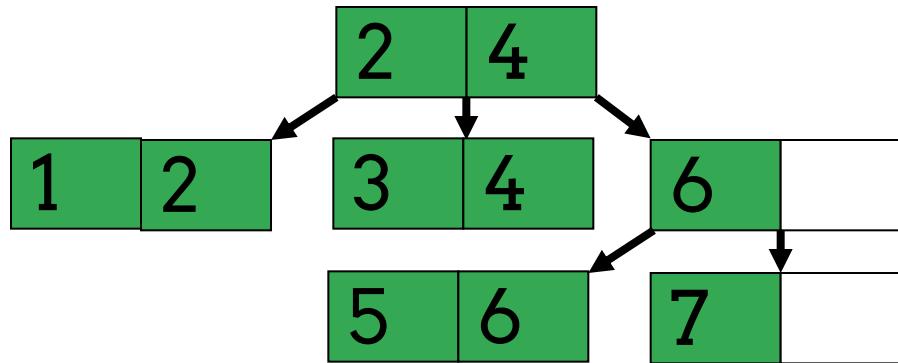
# Insertions



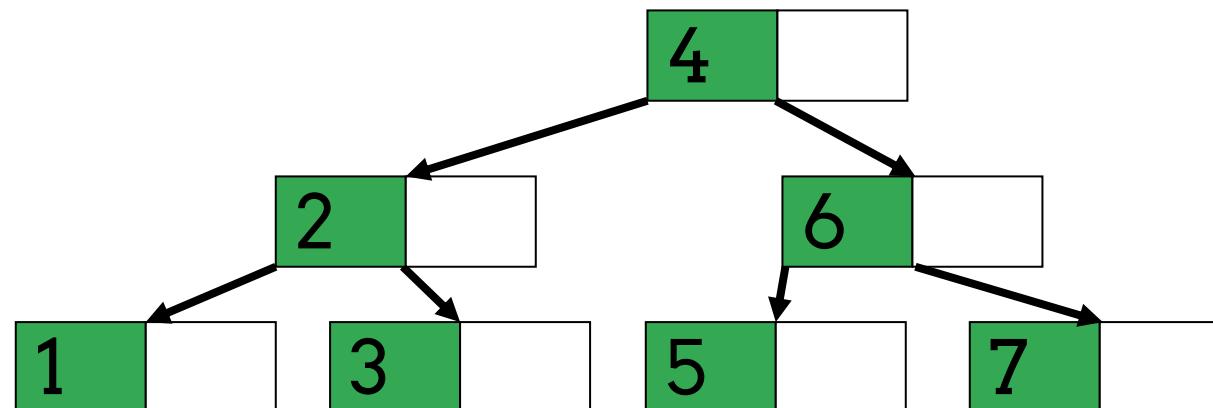
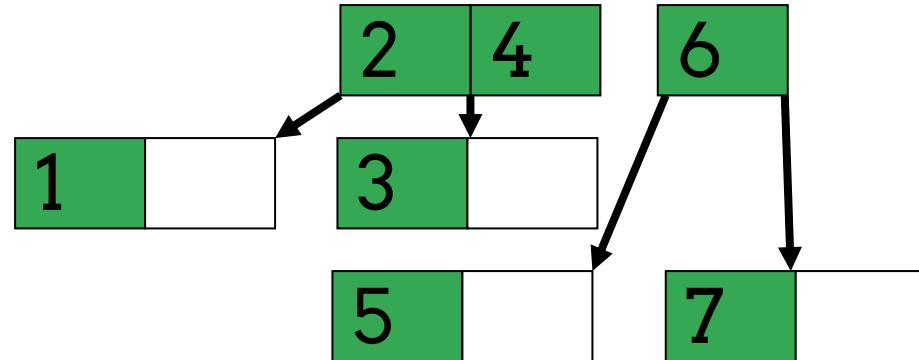
# Insertions



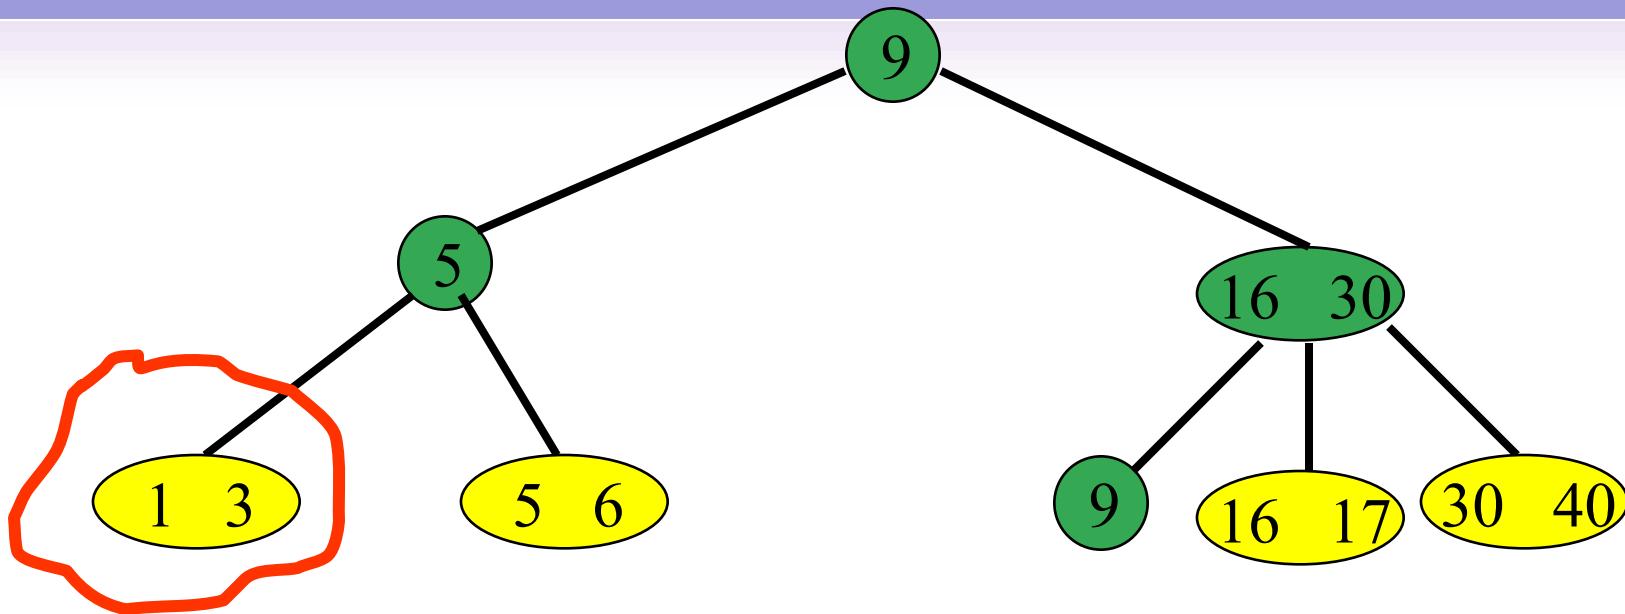
# Insertions



# Insertions



# Insert



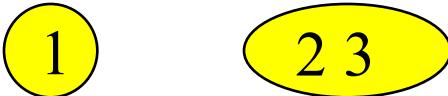
- Insert a pair with key = 2.
- New pair goes into a full node.

# Insert Into A Full Node

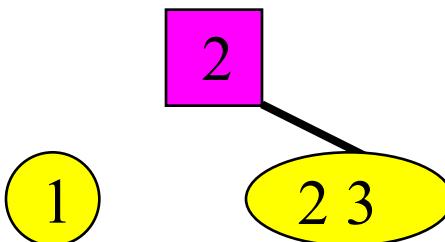
- ◆ Insert new pair so that the keys are in ascending order.



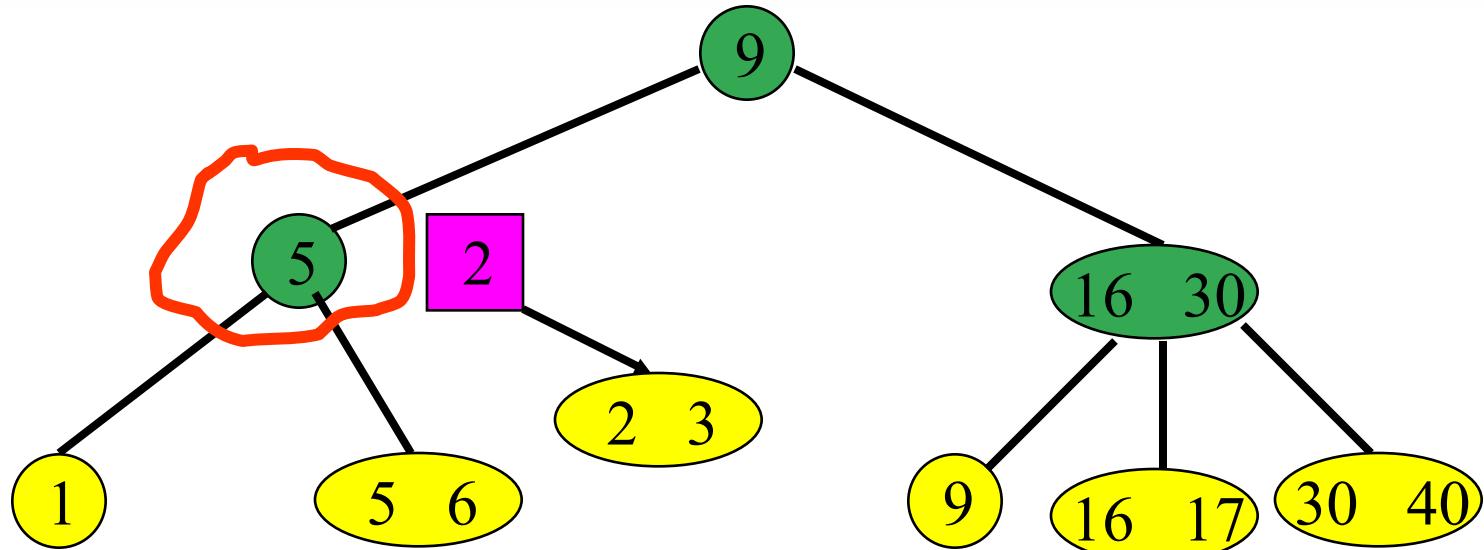
- Split into two nodes.



- Insert smallest key in new node and pointer to this new node into parent.

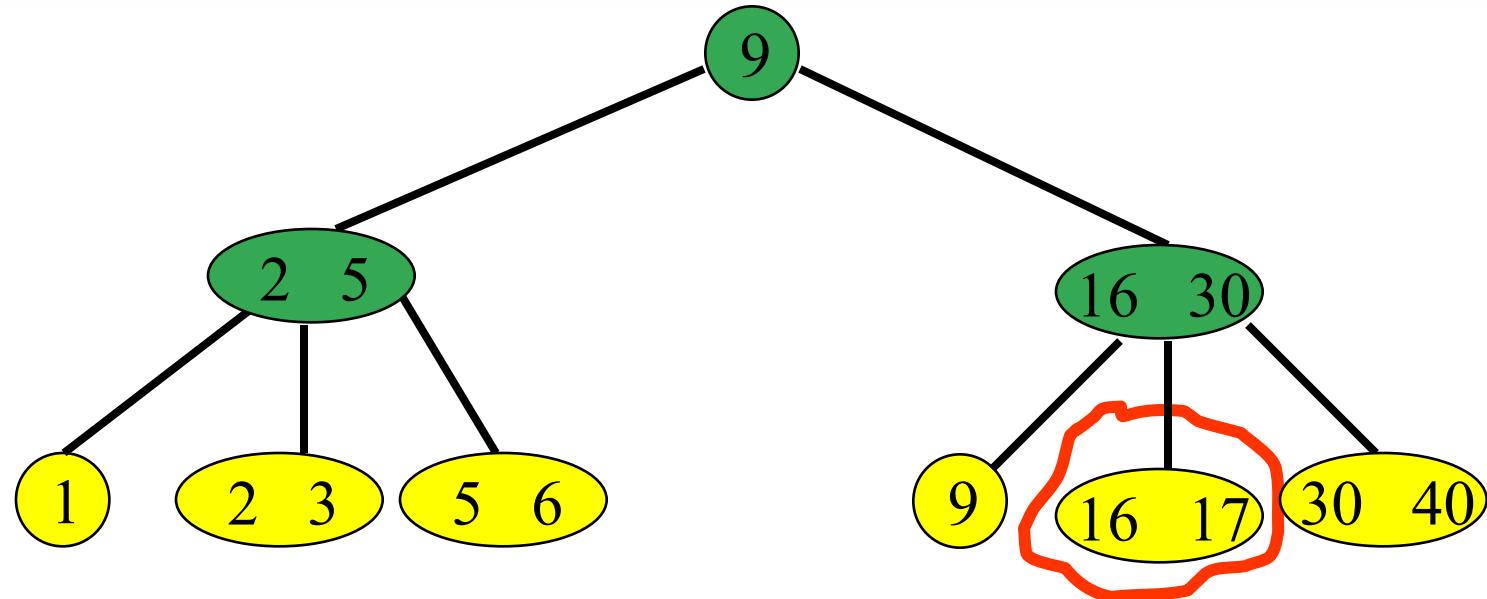


# Insert



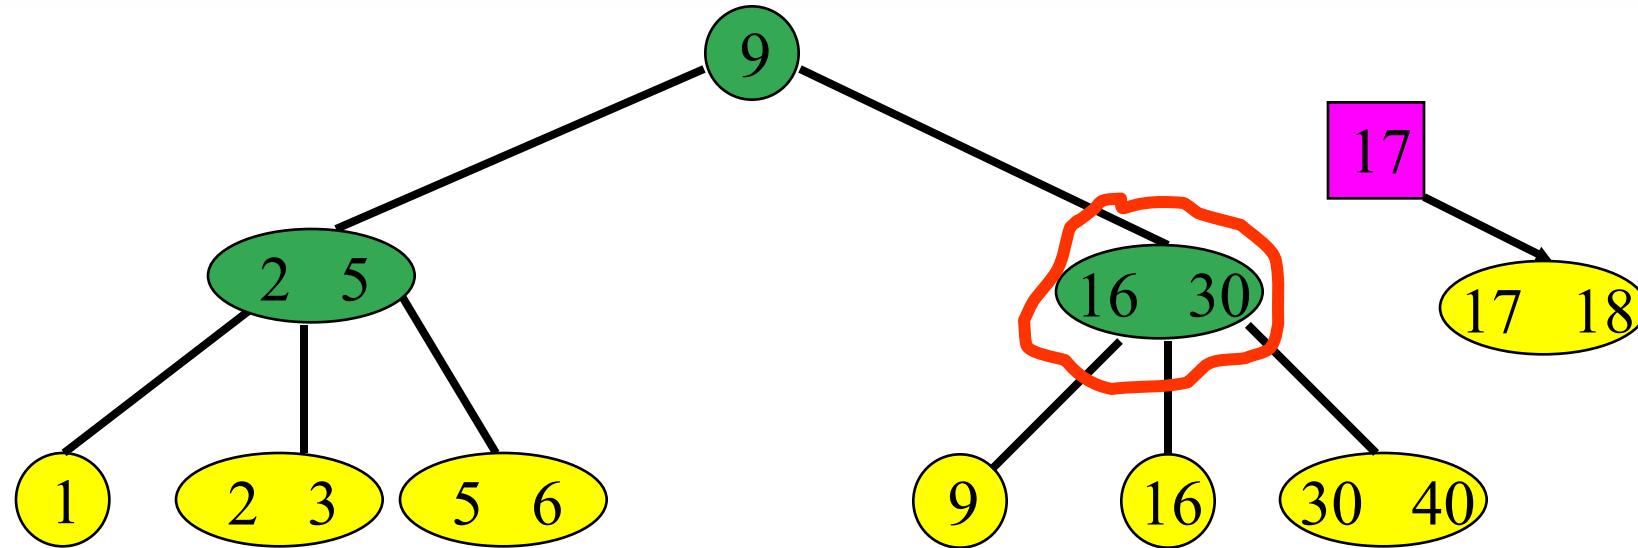
- Insert an index entry **2** plus a pointer into parent.

# Insert



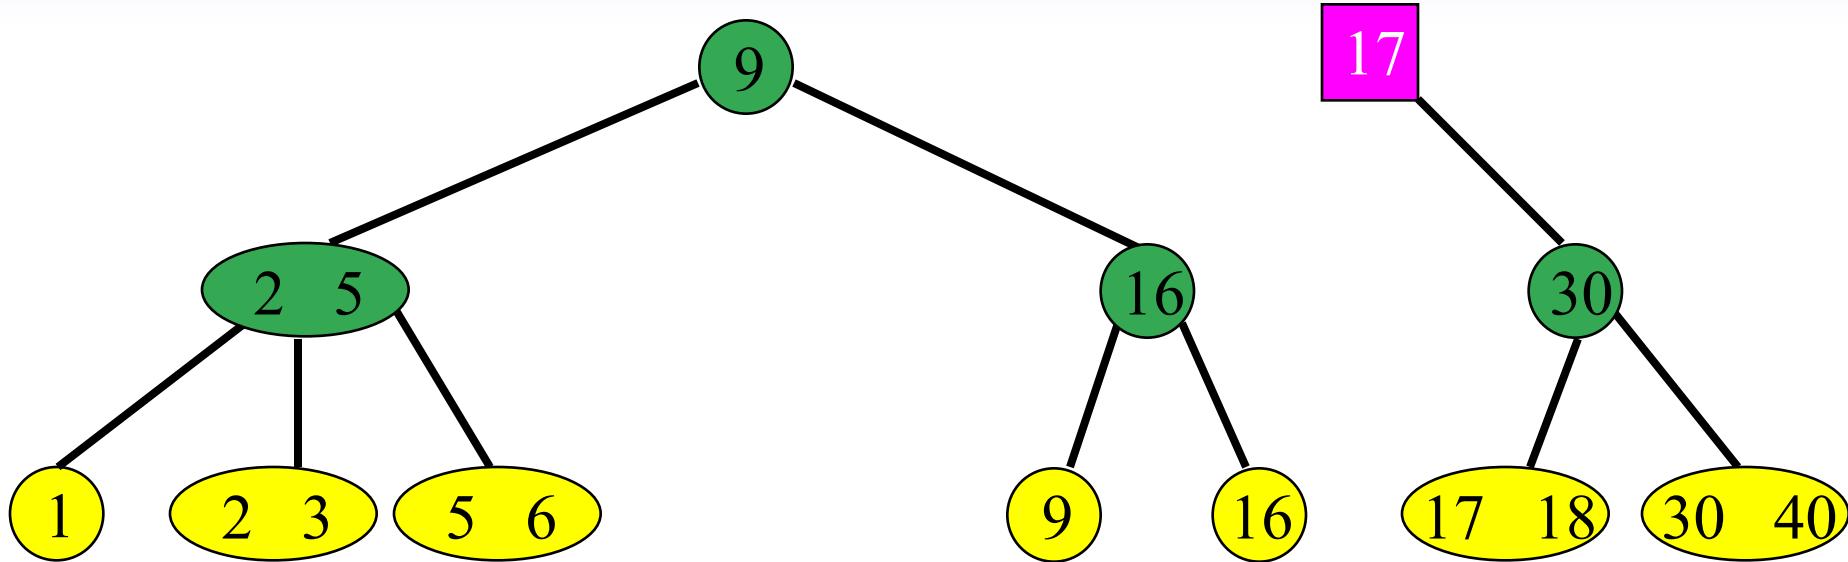
- Now, insert a pair with key = 18.

# Insert



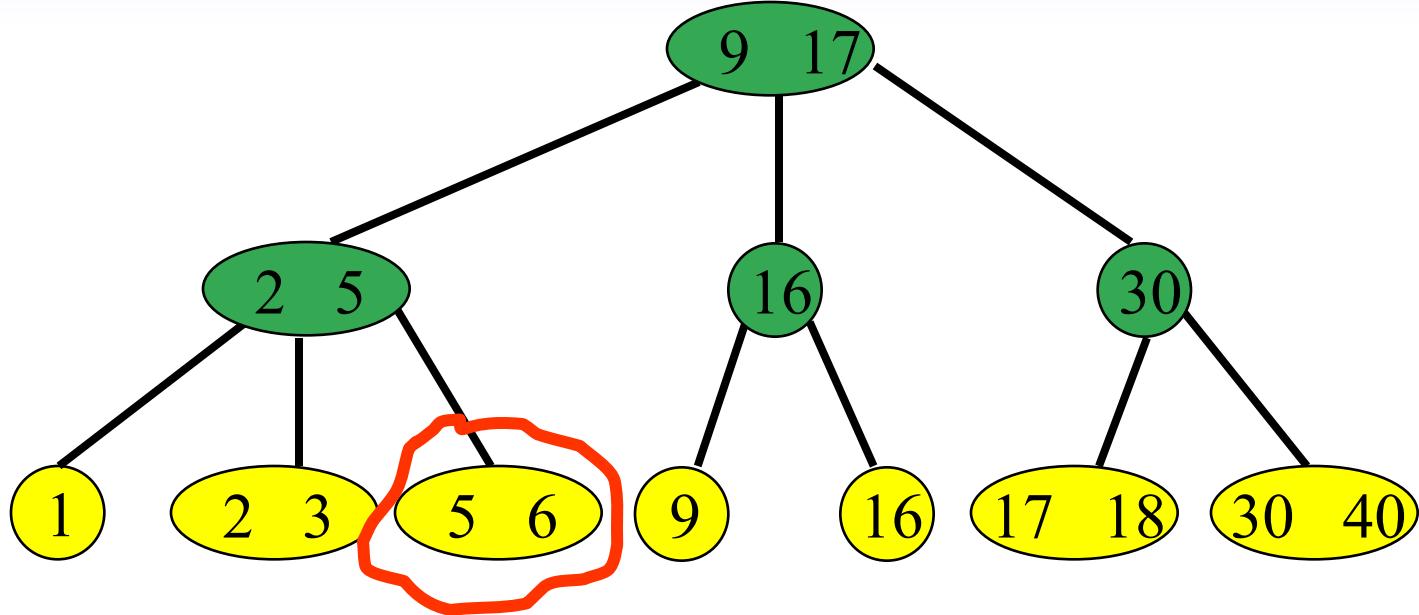
- Now, insert a pair with key = 18.
- Insert an index entry 17 plus a pointer into parent.

# Insert



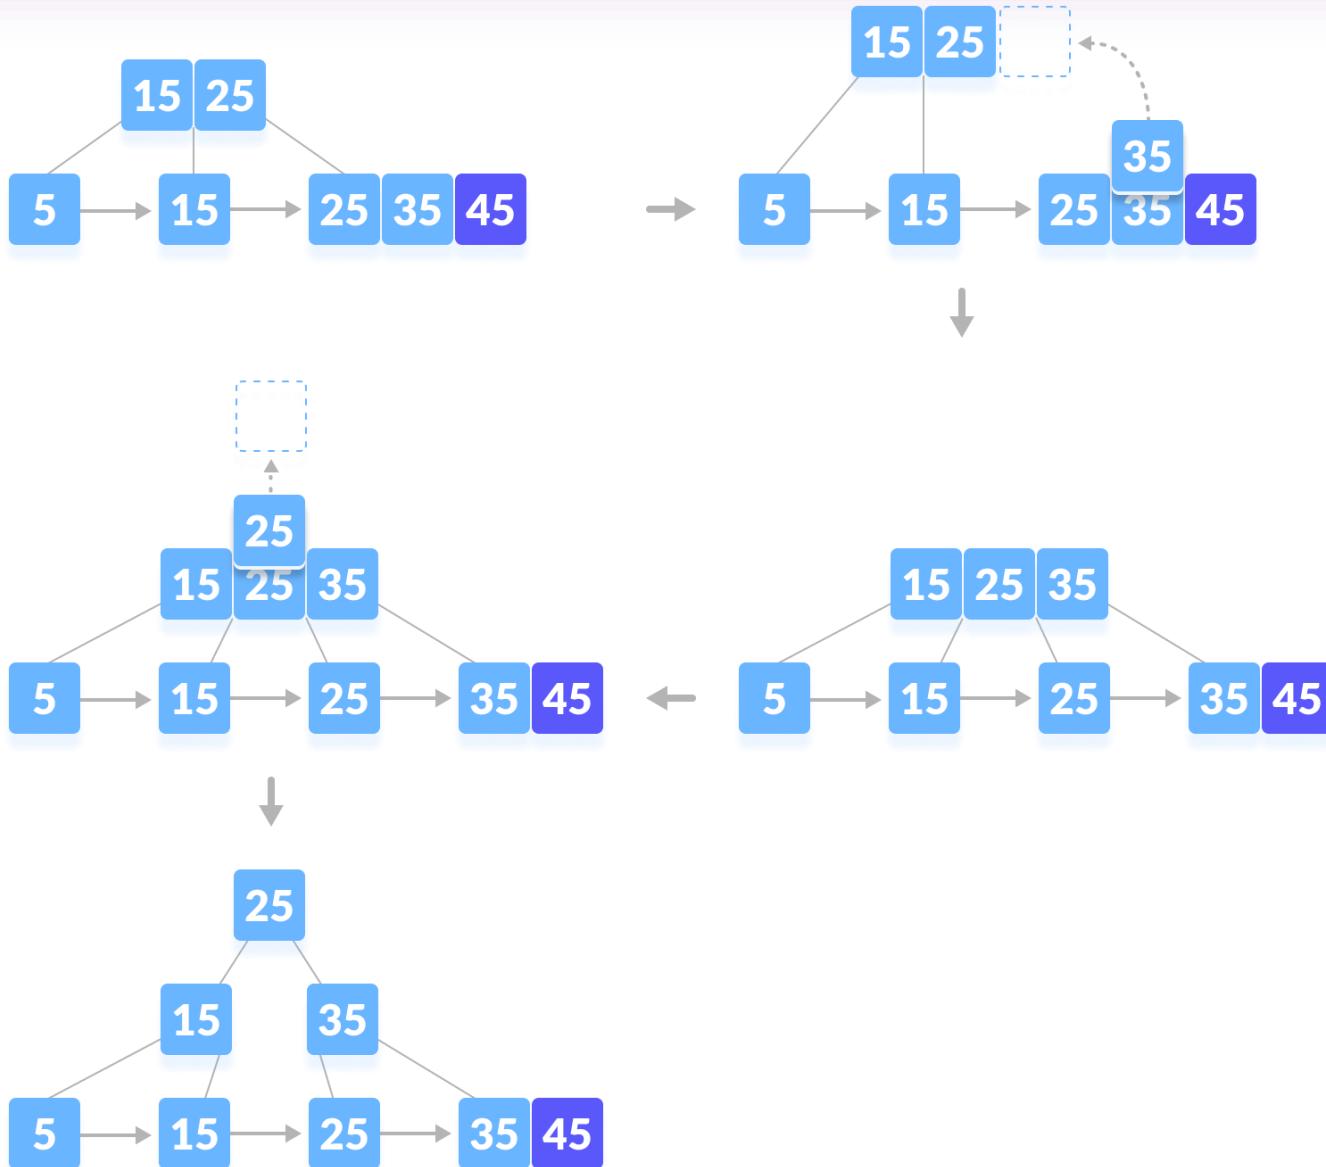
- Now, insert a pair with key = 18.
- Insert an index entry 17 plus a pointer into parent.

# Insert



- Now, insert a pair with key = 7.

# Insertion in B+Trees



# Deletion in B+Trees

- ◆ Deletion involves:
  - ★ 1. Finding and removing the key from the appropriate leaf node
  - ★ 2. Rebalancing the tree to maintain B+ Tree properties:
    - Borrowing keys from siblings or
    - Merging nodes if necessary

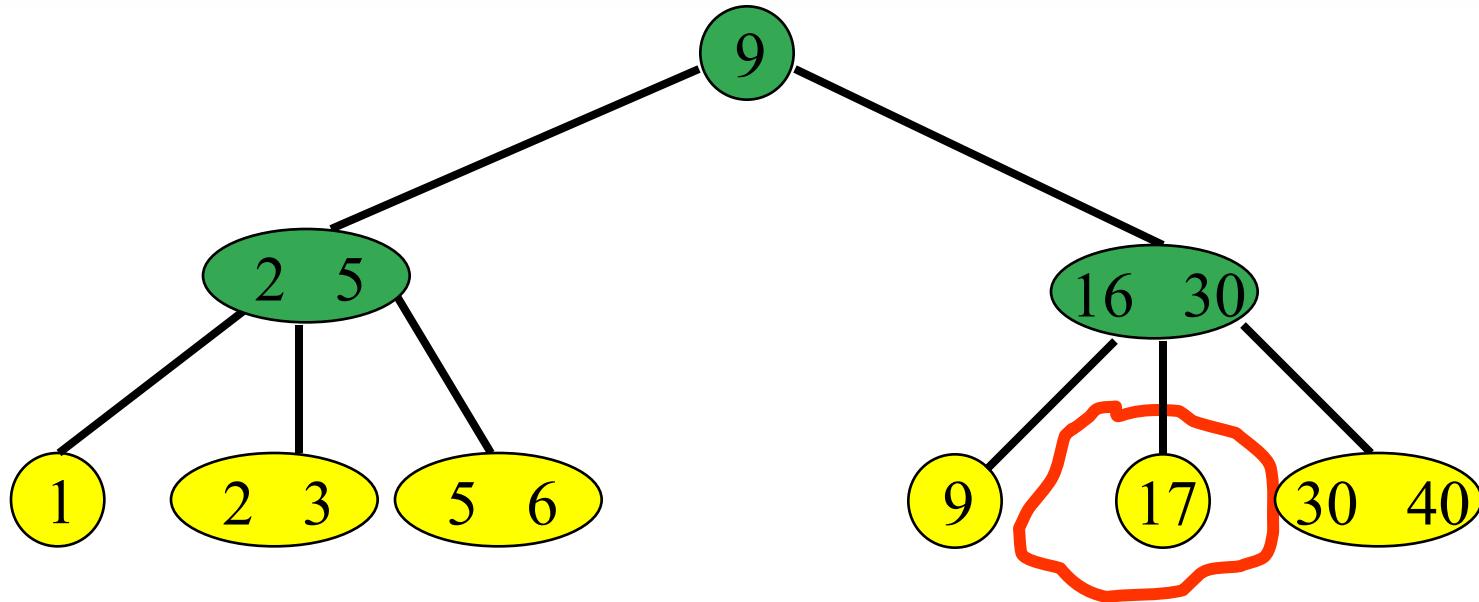
# B+TREE – DELETE

- ◆ Start at root, find leaf **L** where entry belongs.  
Remove the entry.
- ◆ If **L** is at least half-full, done! If **L** has only **m/2-1** entries,
  - ★ → Try to re-distribute, borrowing from sibling (adjacent node with same parent as **L**).  
★ → If re-distribution fails, merge **L** and sibling.
- ◆ If merge occurred, must delete entry (pointing to **L** or sibling) from parent of **L**.

# Deletions

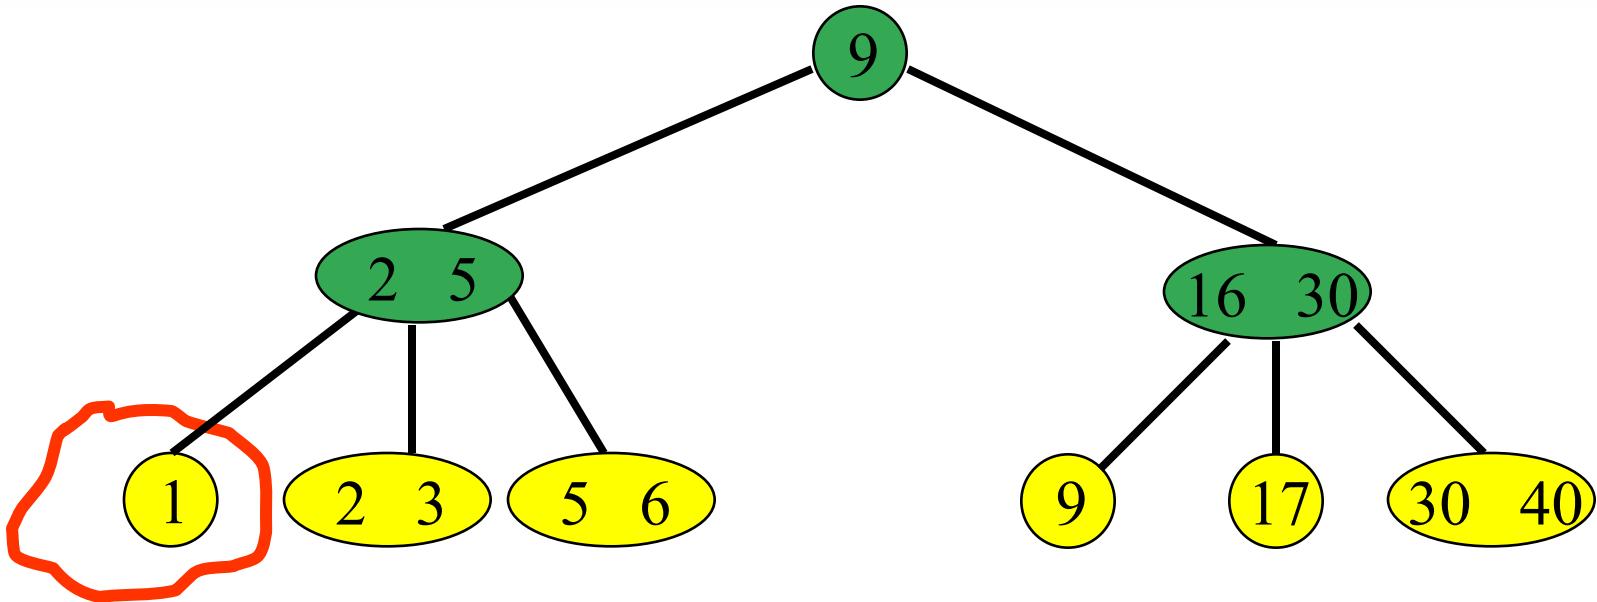
- ◆ **def delete (record) :**
  - ★ Locate target leaf and remove the entry
  - ★ If leaf is less than half full:
    - Try to re-distribute, taking from sibling (adjacent node with same parent)
    - If re-distribution fails:
      - Merge leaf and sibling
      - Delete entry to one of the two merged leaves
      - Merge could propagate to root

# Delete



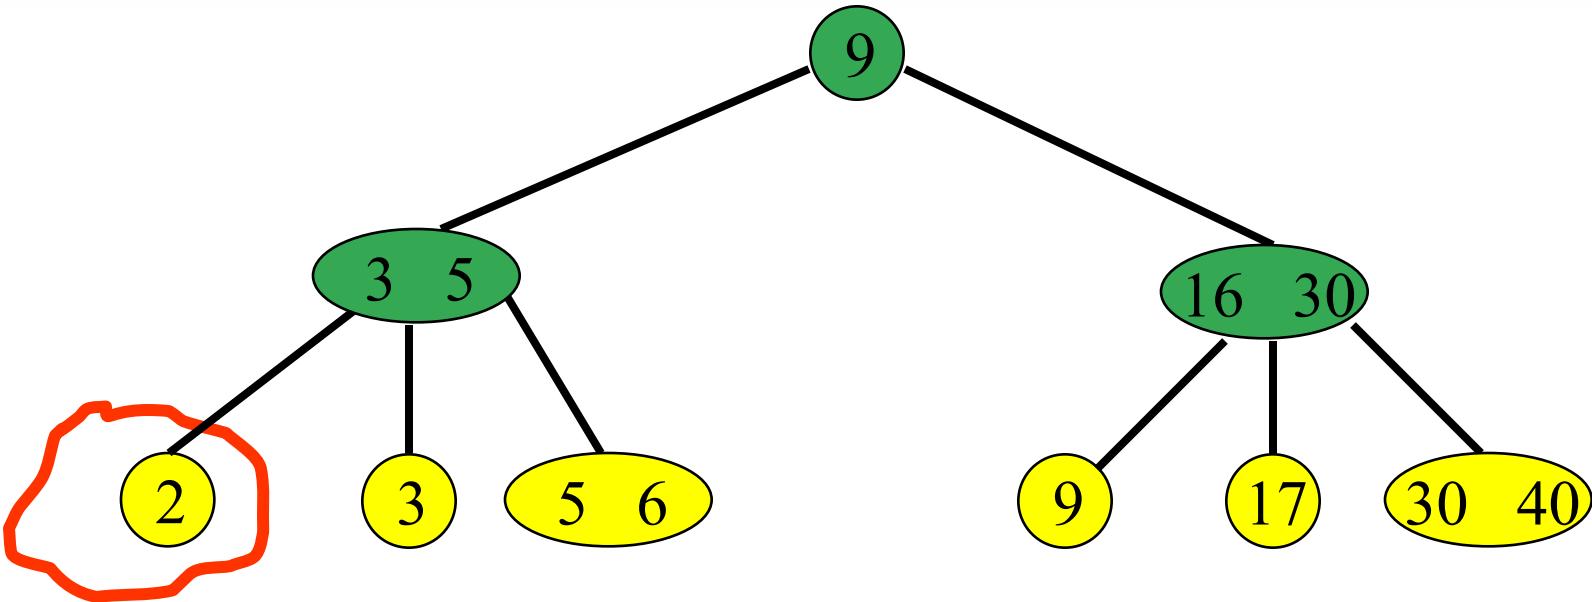
- Delete pair with key = 16.
- Note: delete pair is always in a leaf.

# Delete



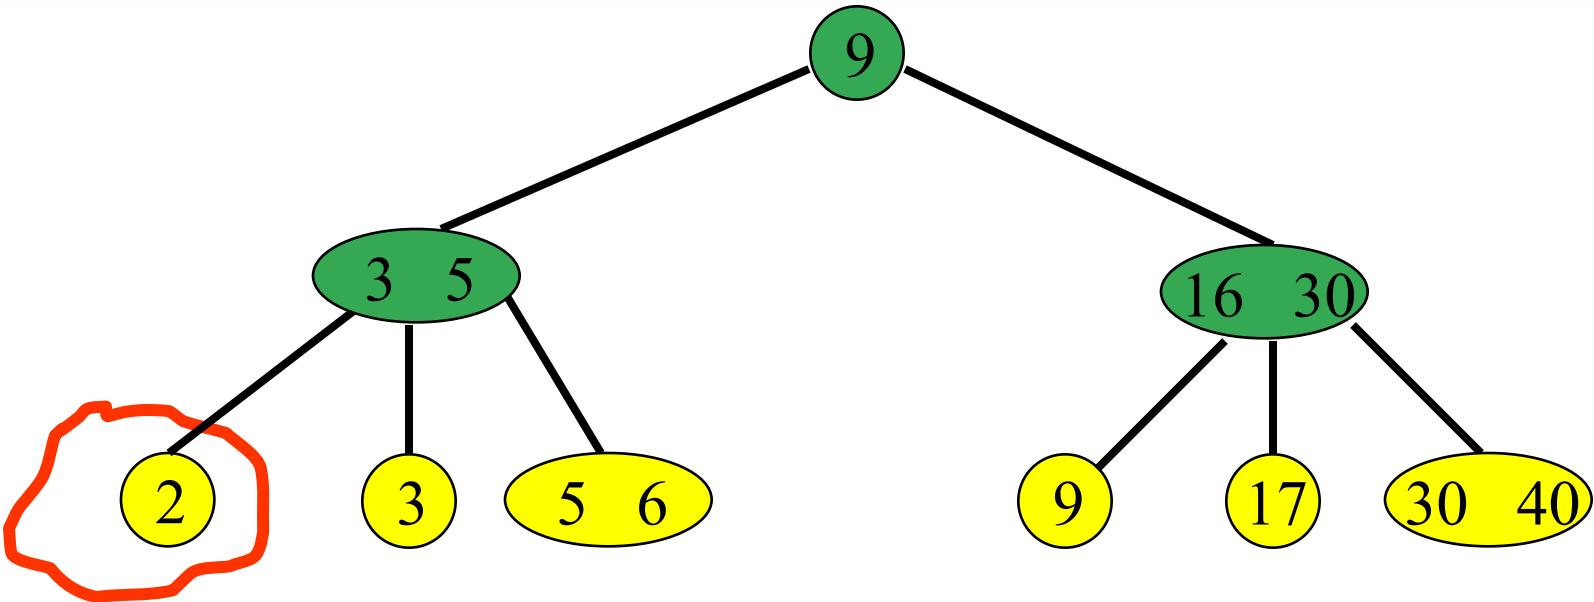
- Delete pair with key = 1.
- Get  $\geq 1$  from adjacent sibling and update parent key.

# Delete



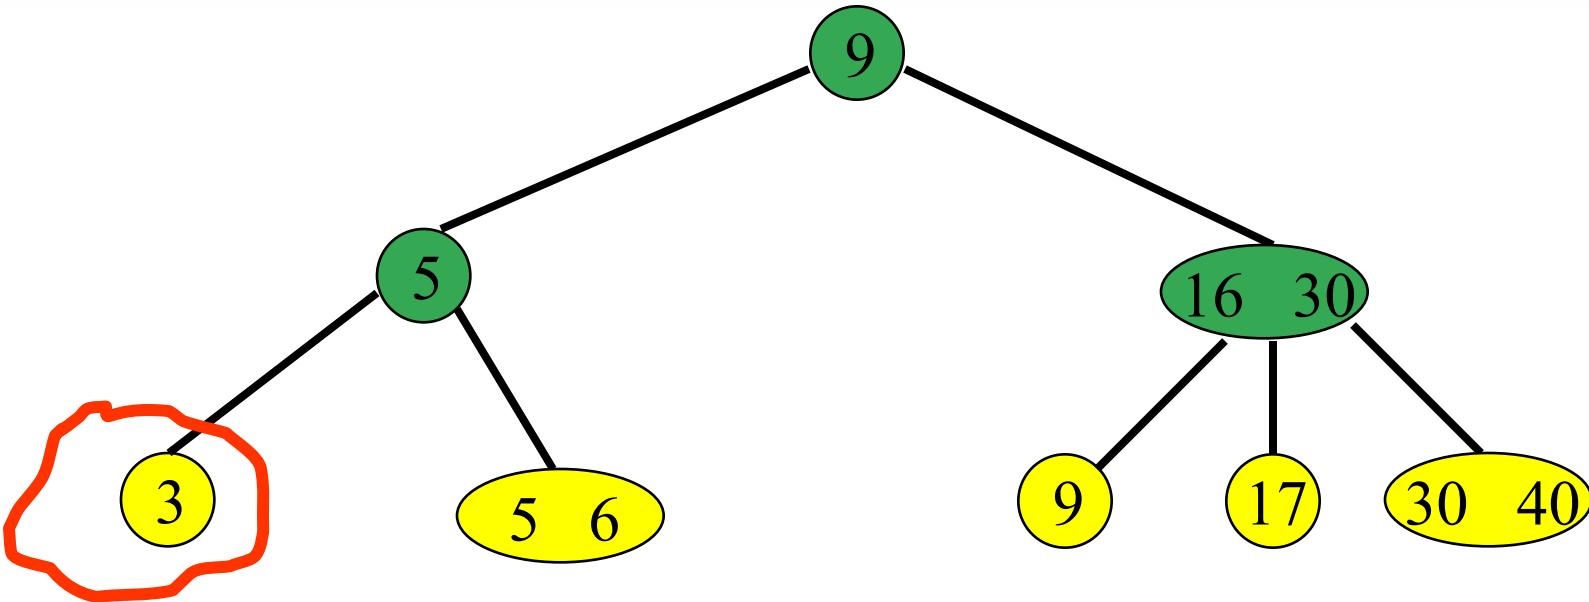
- Delete pair with key = 1.
- Get  $\geq 1$  from sibling and update parent key.

# Delete



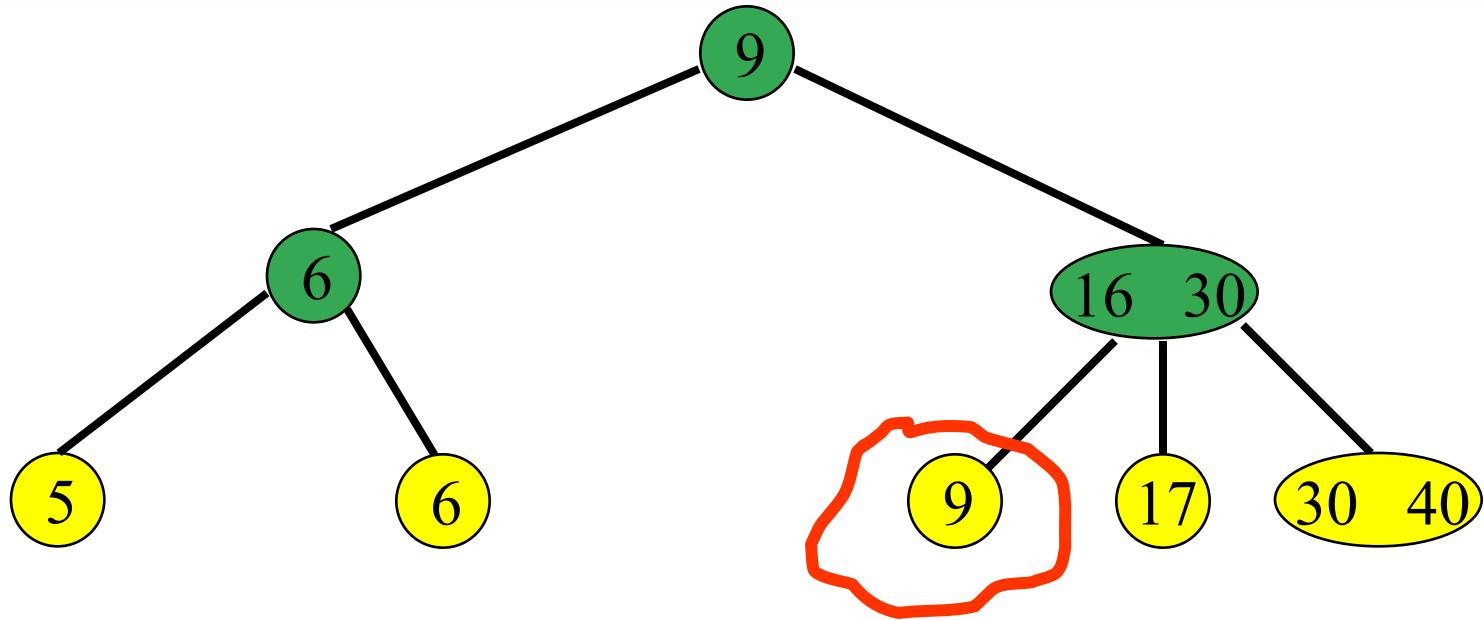
- Delete pair with key = 2.
- Merge with sibling, delete in-between key in parent.

# Delete



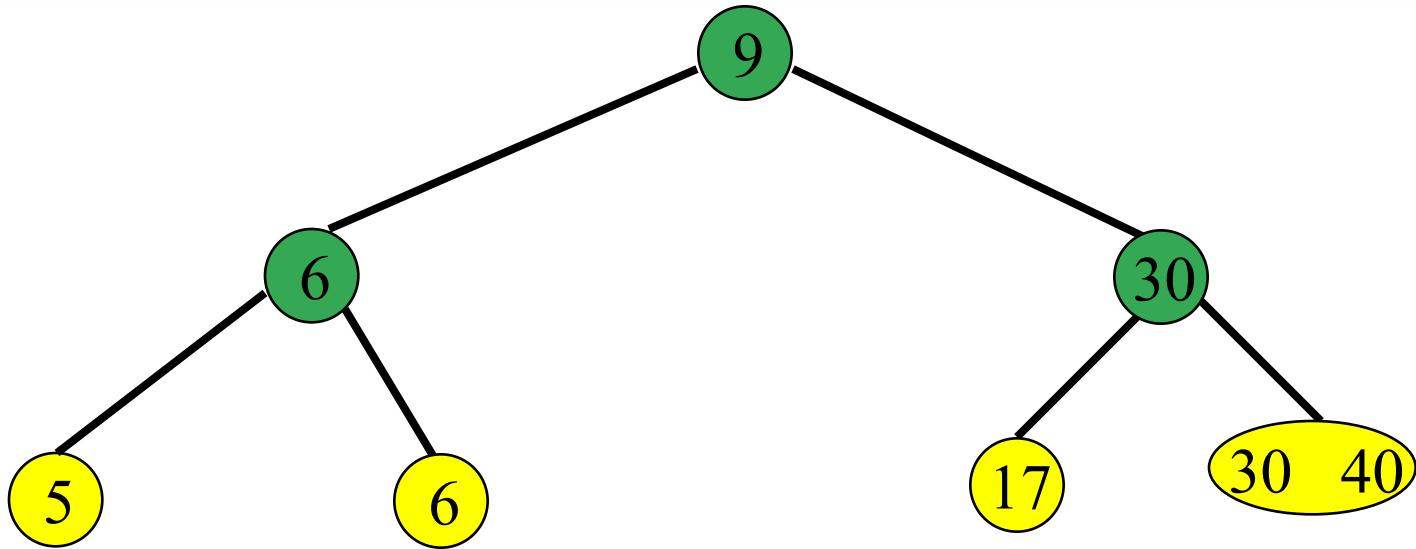
- Delete pair with key = 3.
- Get  $\geq 1$  from sibling and update parent key.

# Delete

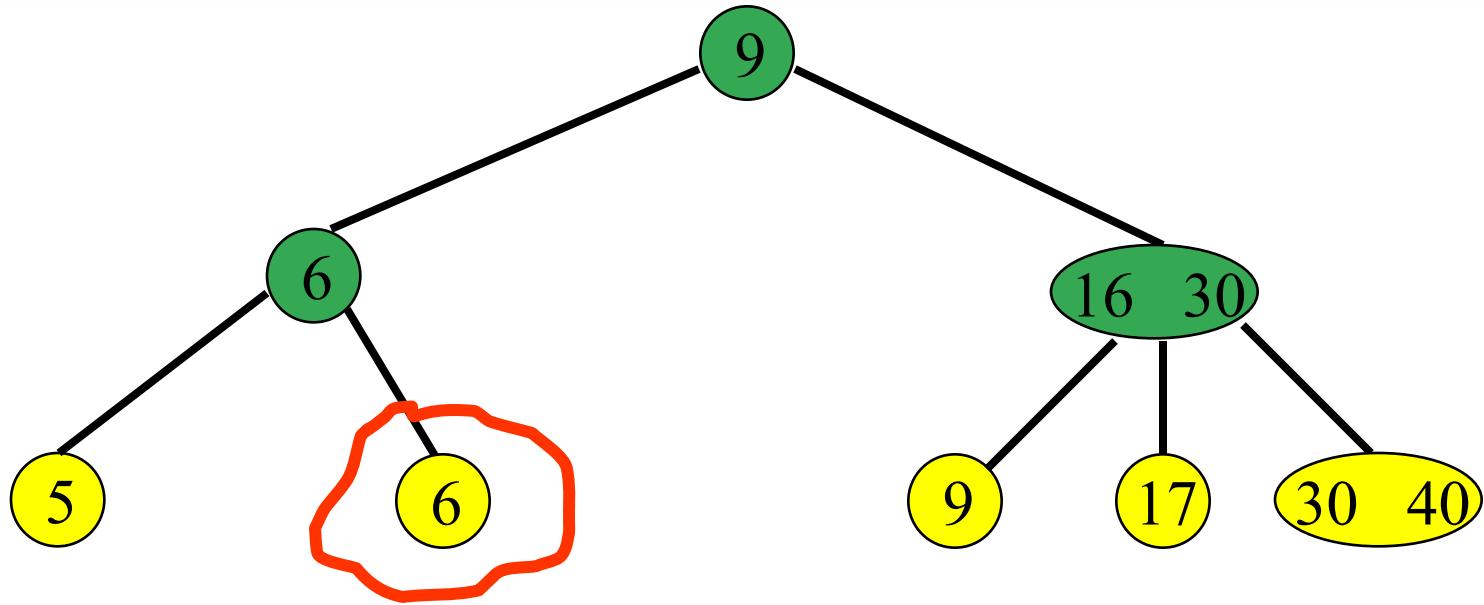


- Delete pair with key = 9.
- Merge with sibling, delete in-between key in parent.

# Delete

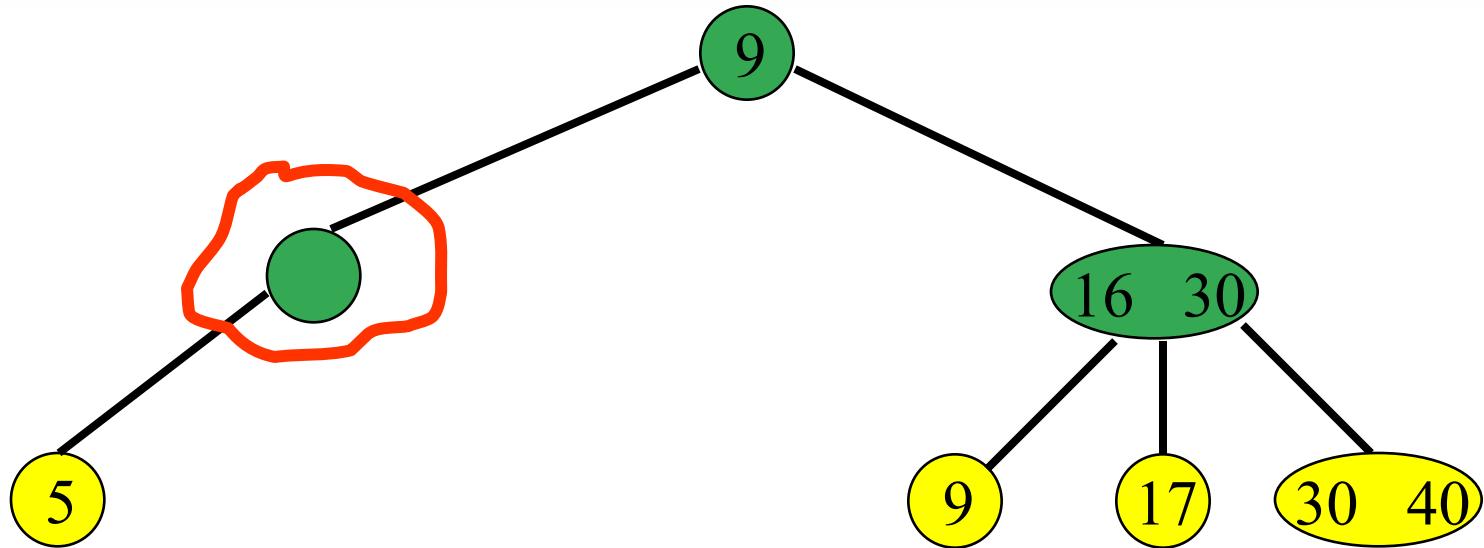


# Delete



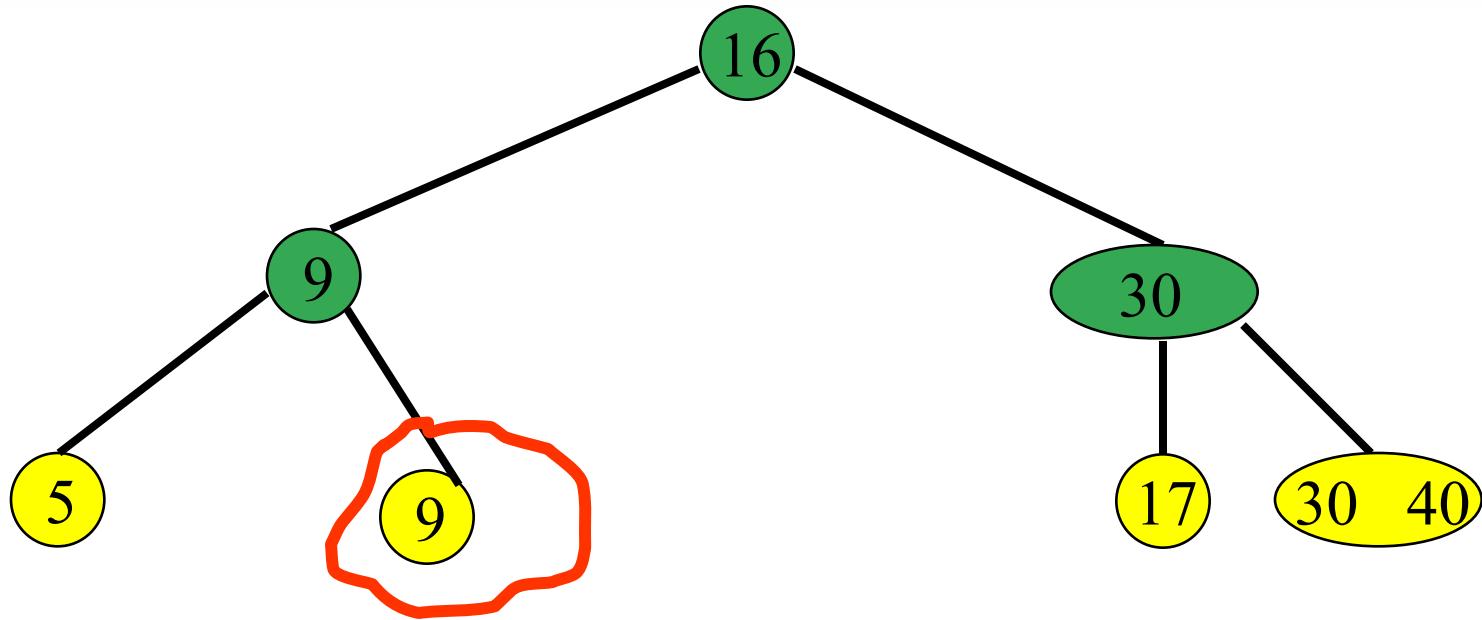
- Delete pair with key = 6.
- Merge with sibling, delete in-between key in parent.

# Delete



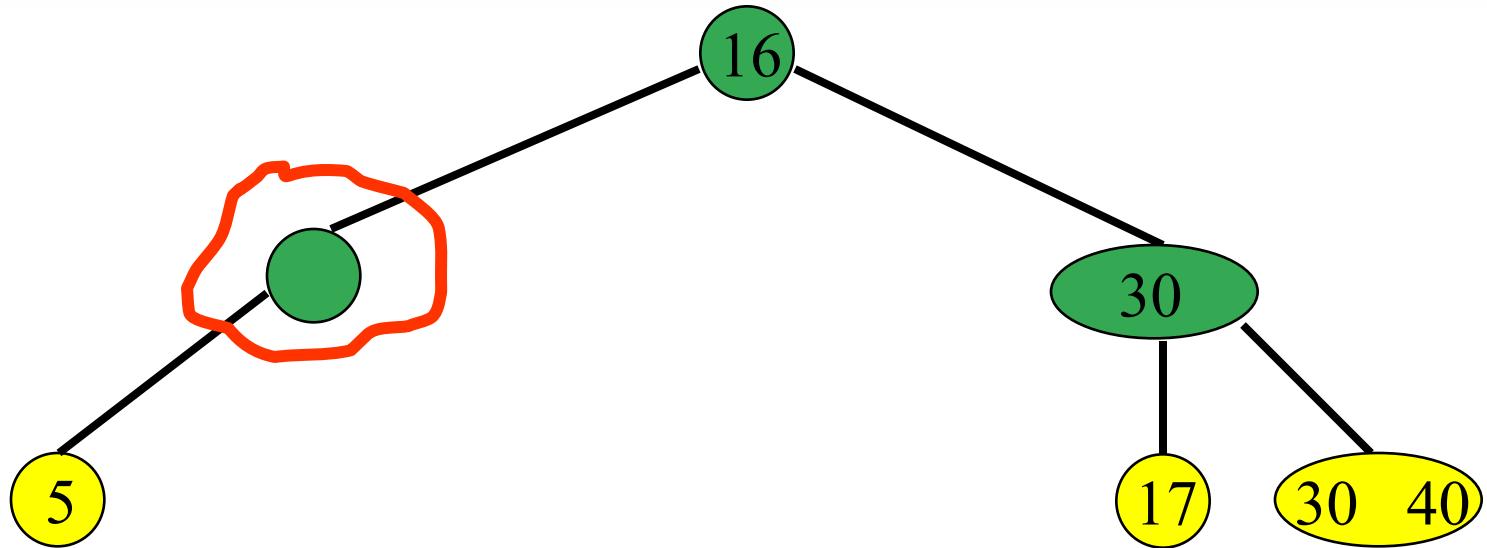
- Index node becomes deficient.
- Get  $\geq 1$  from sibling, move last one to parent, get parent key.

# Delete



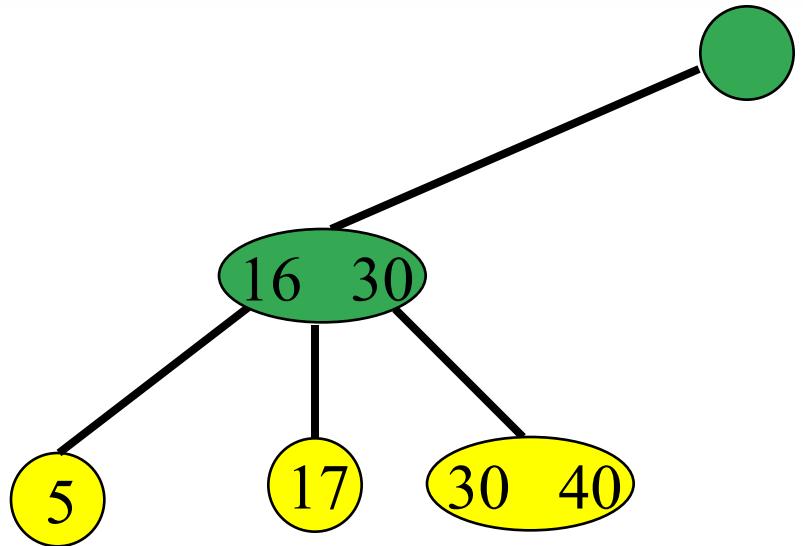
- Delete 9.
- Merge with sibling, delete in-between key in parent.

# Delete



- Index node becomes deficient.
- Merge with sibling and in-between key in parent.

# Delete



- Index node becomes deficient.
- It's the root; discard.

# Leaf node values

## ◆ 1: Record IDs



## ◆ 2: Tuple Data



★ Index-Organized Storage

★ Primary Key Index

○ Leaf nodes store the contents of the tuple.

★ Secondary Indexes

○ Leaf nodes store tuples' primary key as their values.

# B-TREE VS. B+TREE

- ◆ The original **B-Tree** from 1971 stored keys and values in all nodes in the tree.
  - ★ More space-efficient, since each key only appears once in the tree.
- ◆ A **B+Tree** only stores values in leaf nodes.
  - ★ Inner nodes only guide the search process.

# B+Tree design choices

Node Size

Merge Threshold

Variable-Length Keys

Intra-Node Search

# Node Size

- ◆ The slower the storage device, the larger the optimal node size for a B+Tree.
  - ★ HDD: ~1MB
  - ★ SSD: ~10KB
  - ★ In-Memory: ~512B
- ◆ Optimal sizes can vary depending on the workload
  - ★ Leaf Node Scans vs. Root-to-Leaf Traversals

# Merge Threshold

- ◆ Some DBMSs do not always merge nodes when they are half full.
  - ★ Average occupancy rate for B+Tree nodes is 69%.
- ◆ Delaying a merge operation may reduce the amount of reorganization.
- ◆ It may also be better to let underfilled nodes exist and then periodically rebuild entire tree.
- ◆ This is why PostgreSQL calls their B+Tree a "non-balanced" B+Tree (**nbtree**).

# Variable-length Keys

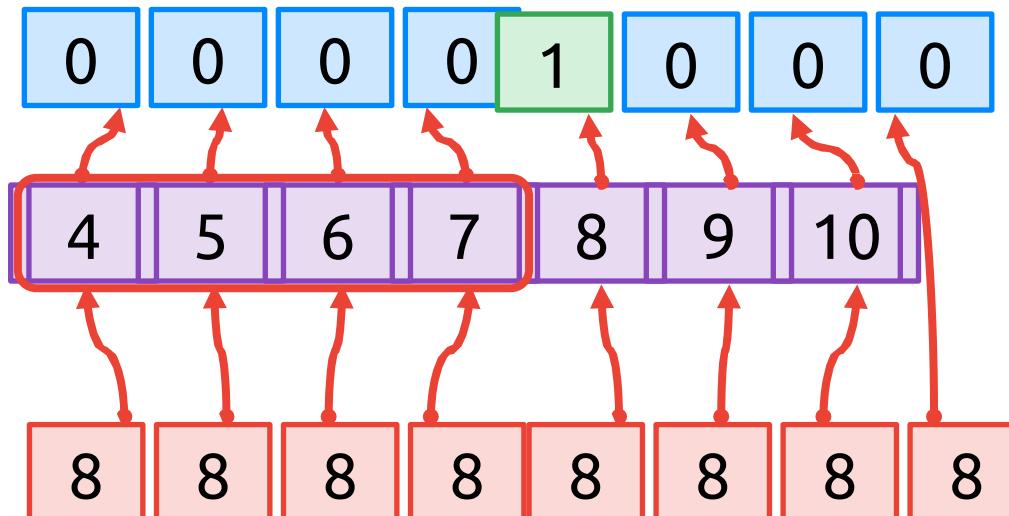
- ◆ Pointers
  - ★ Store the keys as pointers to the tuple's attribute.
- ◆ Variable-Length Nodes
  - ★ The size of each node in the index can vary.
  - ★ Requires careful memory management.
- ◆ Padding
  - ★ Always pad the key to be max length of the key type.
- ◆ Key Map / Indirection
  - ★ Embed an array of pointers that map to the key + value list within the node.

# Intra-node Search

## ◆ Linear

- ★ Scan node keys from beginning to end.
- ★ Use SIMD to vectorize comparisons.

Find Key=8

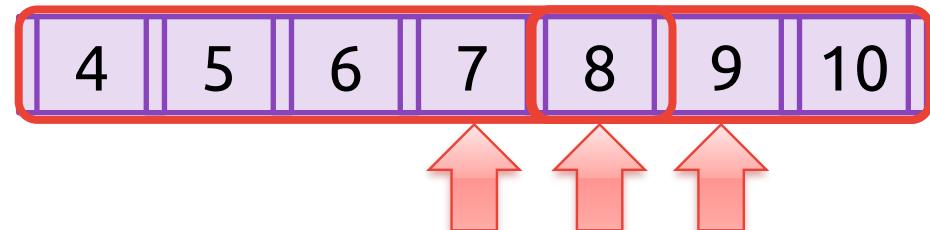


# Intra-node Search

## ◆ Binary

- ★ Jump to middle key
- ★ Pivot left/right depending on comparison.

Find Key=8



# Intra-node Search

## ◆ Interpolation

★ Approximate location of desired key based on known distribution of keys.

Find Key=8



$$\text{Offset: } \frac{8-4}{10-4} \times 7 = 4$$

# Optimizations

Prefix Compression

Suffix Truncation

Bulk Insert

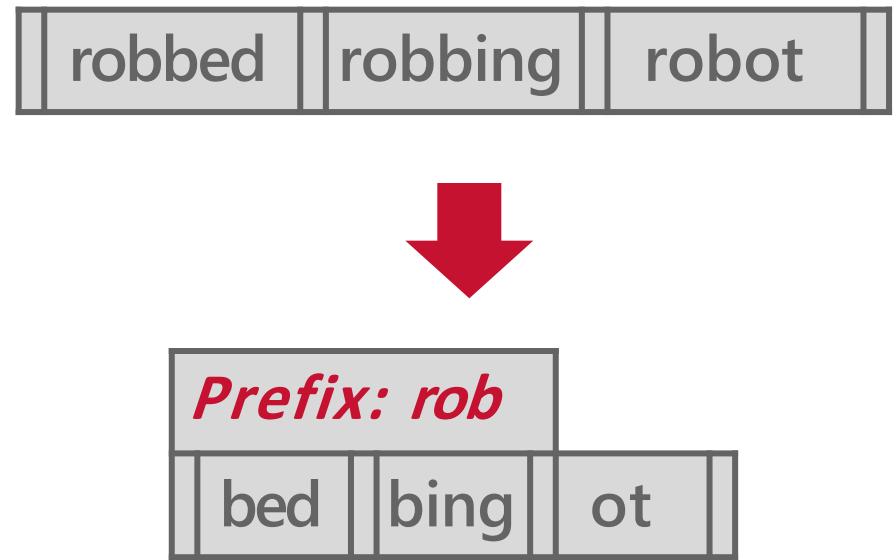
Deduplication

Pointer Swizzling

Buffered Updates

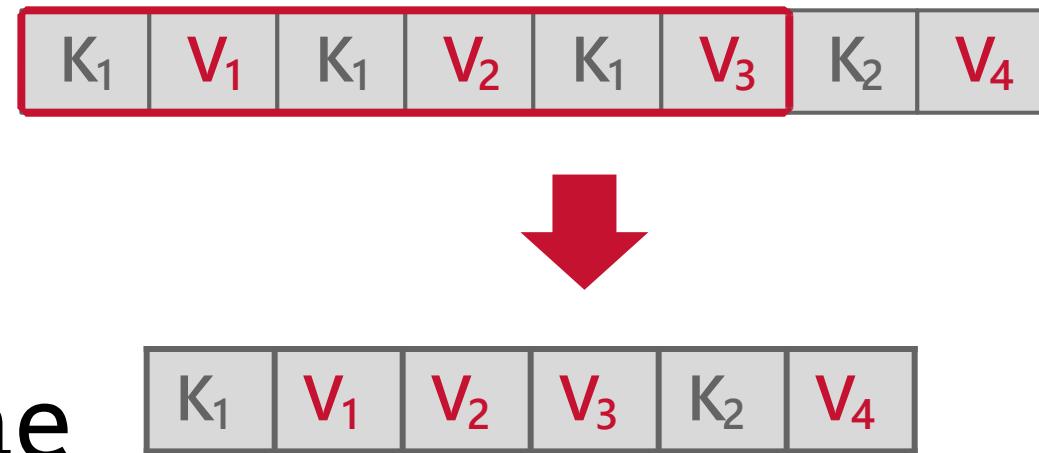
# Prefix compression

- ◆ Sorted keys in the same leaf node are likely to have the same prefix.
- ◆ Instead of storing the entire key each time, extract common prefix and store only unique suffix for each key.



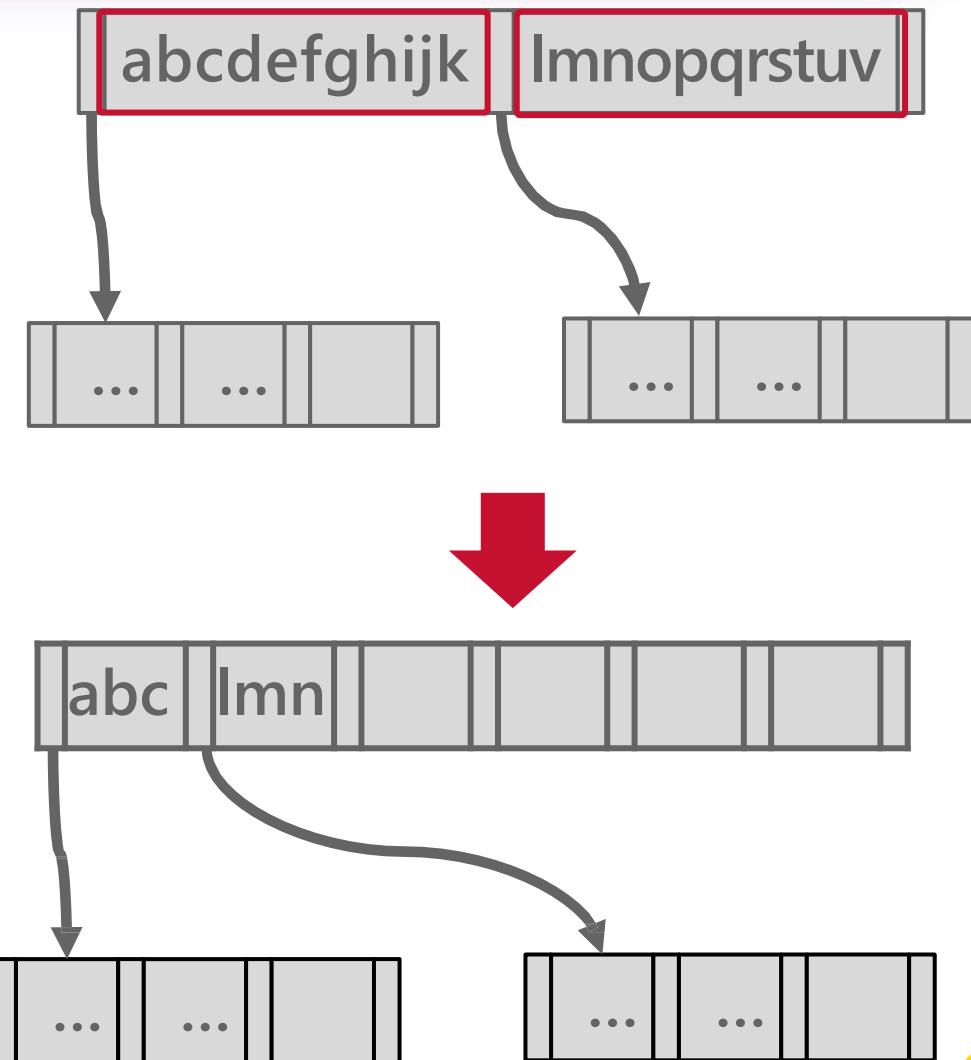
# Deduplication

- ◆ Non-unique indexes can end up storing multiple copies of the same key in leaf nodes.
- ◆ The leaf node can store the key once and then maintain a "posting list" of tuples with that key



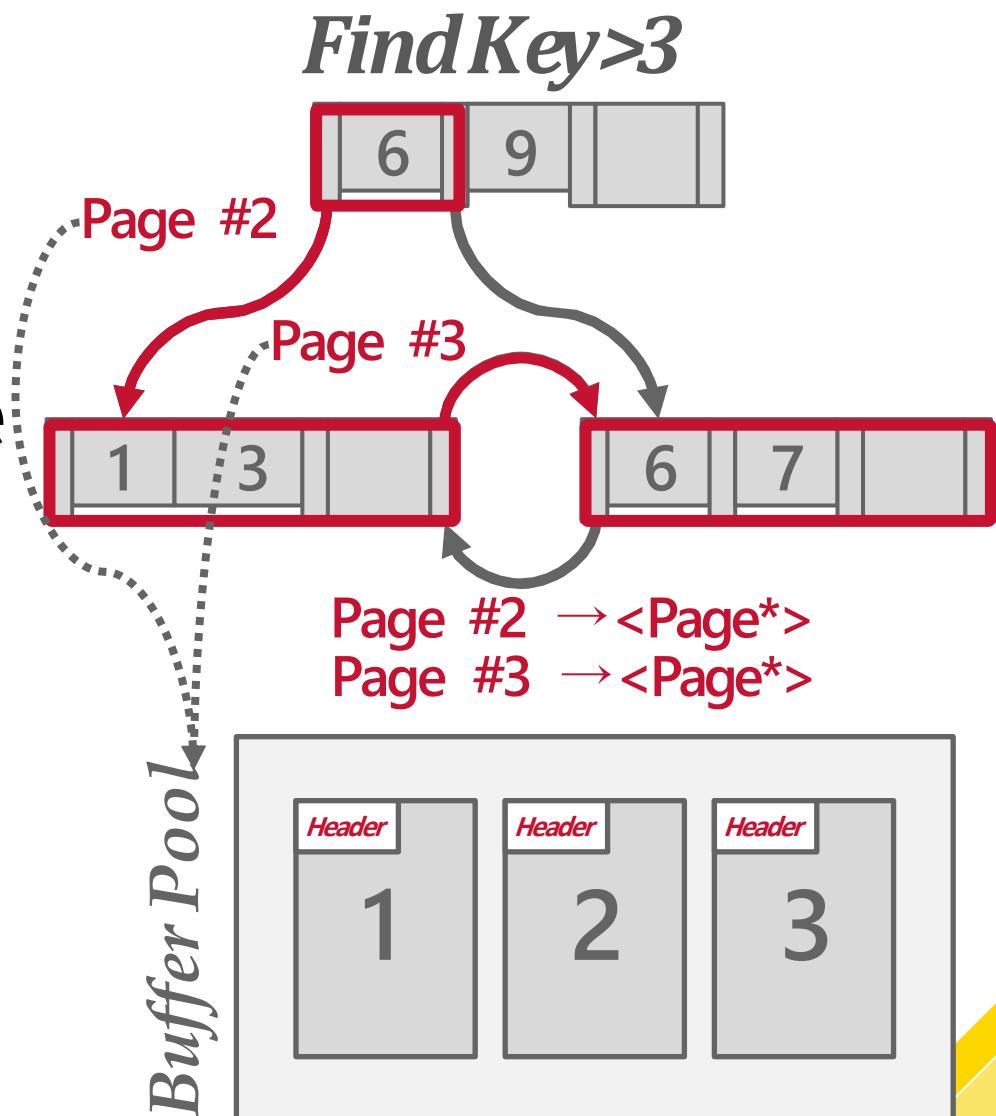
# Suffix truncation

- ◆ The keys in the inner nodes are only used to "direct traffic".
  - ★ We don't need the entire key.
- ◆ Store a minimum prefix that is needed to correctly route probes into the index.



# Pointer swizzling

- ◆ Nodes use page ids to reference other nodes in the index.
- ◆ The DBMS must get the memory location from the page table during traversal.
- ◆ If a page is pinned in the buffer pool, then we can store raw pointers instead of page ids.
- ◆ This avoids address lookups from the page table.

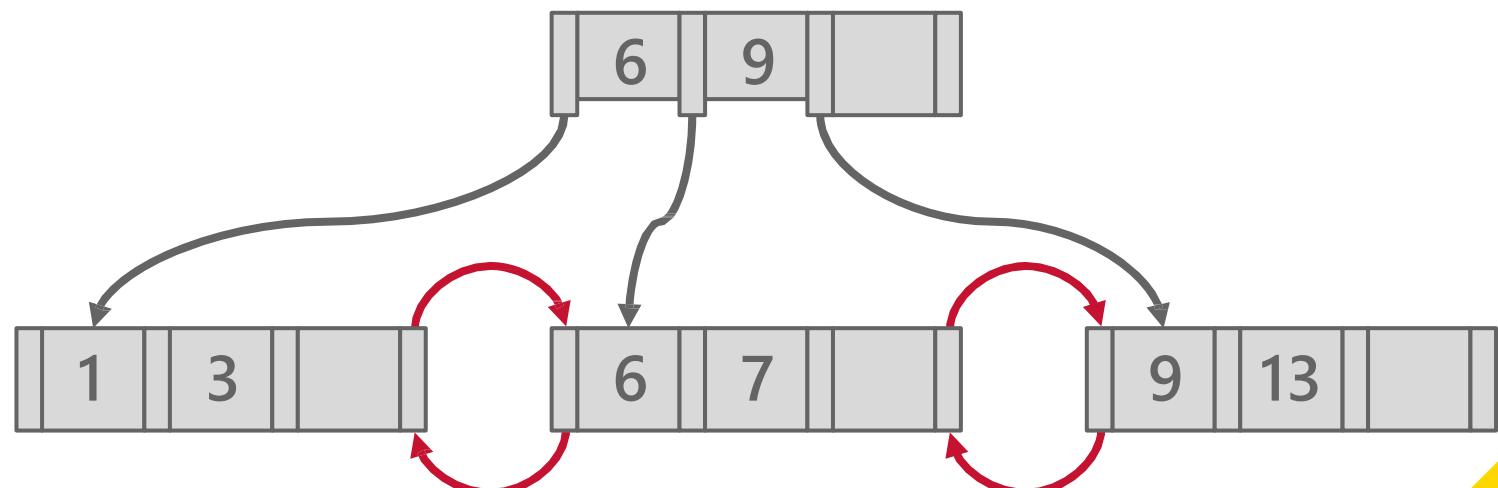


# Bulk insert

- ◆ The fastest way to build a new B+Tree for an existing table is to first sort the keys and then build the index from the bottom up.

Keys: 3, 7, 9, 13, 6, 1

Sorted Keys: 1, 3, 6, 7, 9, 13





THANK  
YOU ☺

