

**CSAI 302**

# **Advanced Database Systems**

**Lec 07**

## **Transaction Processing and Schedules**

# Introduction

- ◆ Transactions (informally):
  - ★ mechanism for managing logical units of data processing:  
**independent on others, all or nothing**
  - ★ Examples:
    - a single data retrieval query
    - a sequence of data manipulation queries that should be executed together
- ◆ Transaction management (or processing) systems:
  - ★ systems with **large databases** and many **concurrent users** require  
**high availability** and **fast response** time
  - ★ Examples:
    - banking, airline booking, online retail, stocks

# Fund Transfer Example

◆ Transfer \$50 from account A to account B:

- ★ 1. read(A)
- ★ 2. A := A - 50
- ★ 3. write(A)
  
- ★ 4. read(B)
- ★ 5. B := B + 50
- ★ 6. write(B)

**Failures, such as:**

- hardware failures
- system crashes

**Concurrent execution of  
multiple transactions**

# Fund Transfer Example

- ◆ Transfer \$50 from account A to account B:

- ★ 1. read(A)
- ★ 2.  $A := A - 50$
- ★ 3. write(A)
- ★ 4. read(B)
- ★ 5.  $B := B + 50$
- ★ 6. write(B)

- ◆ If the transaction fails after step 3 and before step 6, money will be “lost” leading to an inconsistent database state
- ◆ The system should ensure that updates of a partially executed transaction are not reflected in the database

**Atomicity requirement**

# Fund Transfer Example

- ◆ Transfer \$50 from account A to account B:

- ★ 1. read(A)
- ★ 2.  $A := A - 50$
- ★ 3. write(A)
- ★ 4. read(B)
- ★ 5.  $B := B + 50$
- ★ 6. write(B)

- ◆ Once the user has been **notified** that the transaction **has completed**
  - ★ The transfer of the \$50 has taken place
- ◆ the updates to the database by the transaction must persist even if there are software or hardware failures.

Durability requirement

# Fund Transfer Example

- ◆ Transfer \$50 from account A to account B:

- ★ 1. read(A)
- ★ 2.  $A := A - 50$
- ★ 3. write(A)
- ★ 4. read(B)
- ★ 5.  $B := B + 50$
- ★ 6. write(B)

- ◆ Sum of A and B **is unchanged** by the execution of the transaction
  - ★ A transaction must see a **consistent database**.
  - ★ During transaction execution the database may be temporarily inconsistent.
  - ★ When the transaction completes successfully the database must be consistent

**Consistency requirement**

# Fund Transfer Example

- ◆ Transfer \$50 from account A to account B:

- ★ 1. `read(A)`
- ★ 2. `A := A - 50`
- ★ 3. `write(A)`
- ★ 4. `read(B)`
- ★ 5. `B := B + 50`
- ★ 6. `write(B)`

- ◆ if between steps 3 and 6, another transaction T2 is allowed to access the partially updated database
  - ◆ T2
    - ★ `read(A), read(B), print(A+B)`
    - ◆ it will see an inconsistent database
    - ◆ (the sum  $A + B$  will be less than it should be).

Isolation requirement

# Single-user vs multi-user

## ◆ Single-user DBMS:

- ★ at most one user at a time (usually, personal computer)

## ◆ Multi-user DBMS:

- ★ many users (processes with own computation) with concurrent access to the same data

- (usually, servers with many CPUs, but may be handled by one CPU with interleaving concurrency)

## ◆ Both need transaction management, with:


- ★ concurrency control


- (transactions independent of each other)

- ★ fail recovery

- (each transaction executed all or nothing)

# The place of transaction management





# Transaction idea

- ◆ Transaction is a program that forms a logical unit of database processing
  - ★ has its own memory and computation ability
  - ★ includes one or more access operations to (shared) database (e.g., retrieval, insertion, deletion)
- ◆ boundaries can be specified by **begin** and **end** statements
  - ★ if something goes wrong on the way, effects **roll back**
- ◆ an application program may have several transactions
  - ★ may be **read-only** or **read-write**

# Main transaction operations

## ◆ **read(X)**

- ★ reads an item named X from the global database into a local program variable named X
  - includes finding the address of the block on the disk (or in cache) with X , and copying to a main memory buffer

## ◆ **write(X)**

- ★ writes the value of local program variable named X into the global database item named X
  - includes finding the address of the block on the disk (or in cache) with X , read it to the local memory buffer, modify it, and write it back (to the disk or cache)

## ◆ **Program local operations**

- ★ (for example, update  $X := X + 50$ )

# Example transactions

- ◆ Let X and Y be the numbers of reserved seats in two flights (stored in a database)
- ◆ Transaction T1 transfers N reservations from X to Y
- ◆ Transaction T2 reserves M seats in X
- ◆ Important:

| $T_1$  |
|--|
| read(X);<br>$X := X - N;$<br>write(X);<br>read(Y);<br>$Y := Y + N;$<br>write(Y); |

| $T_2$                                  |
|--|
| read(X);<br>$X := X + M;$<br>write(X); |

- ★ updates are local, the database is not updated until the new value is written
- ★ some commands are often omitted if they are not relevant (both local and transaction management)

# Types of failures

Computer failure  
(system crash)

- hardware, software, network errors

Transaction failure

- division by zero, integrity constraint violation, user interrupt, etc.

Local transaction errors

- no data found, programmed exception, etc.

Concurrency control enforcement

- serializability violation, deadlock resolving, etc.

Disk failure

- errors with disk reads or writes


Physical problems

- power cut, fire, catastrophe, etc.

# Transaction State

- ◆ Active
  - ★ the initial state; the transaction stays in this state while it is executing
- ◆ Partially committed
  - ★ after the final statement has been executed.
- ◆ Failed
  - ★ after the discovery that normal execution can no longer proceed.
- ◆ Aborted
  - ★ after the transaction has been rolled back and the database restored to its state prior to the start of the transaction.
  - ★ Two options after it has been aborted:
    - Restart the transaction (*only if no internal logical error*)
    - Kill the transaction
- ◆ Committed
  - ★ after successful completion.

# Transaction State (Cont.)





# **Desirable properties ACID principles**

# ACID properties

## Atomicity

Transaction performed in its entirety or not at all

ensured by transaction **recovery subsystem**

## Consistency

The database should always remain consistent

ensured by transaction **recovery subsystem**

## Isolation

Transaction should not interfere with other transactions

ensured by the **concurrency control subsystem**

## Durability

Changes of committed transactions must persist

ensured by transaction **recovery subsystem**


# Example

- ◆ Consider two transactions (Xacts):
  - ◆ **T1: BEGIN A=A+100, B=B-100 END**
  - ◆ **T2: BEGIN A=1.06\*A, B=1.06\*B END**
- ★ 1st xact transfers \$100 from B's account to A's
- ★ 2nd credits both accounts with 6% interest.
- ◆ Assume at first A and B each have \$1000. What are the legal outcomes of running T1 and T2?
  - ★ T1 ; T2 (A=1166,B=954)
  - ★ T2 ; T1 (A=1160,B=960)
  - ★ In either case,  $A+B = \$2000 * 1.06 = \$2120$
  - ★ There is no guarantee that T1 will execute before T2 or vice-versa, if both are submitted together.

# **Example (Contd.)**

- ◆ Consider a possible interleaved schedule:
  - ◆ T1: A=A+100, B=B-100
  - ◆ T2: A=1.06\*A, B=1.06\*B
    - ★ This is OK (same as T1;T2). But what about:
  - ◆ T1: A=A+100, B=B-100
  - ◆ T2: A=1.06\*A, B=1.06\*B
    - ★ Result: A=1166, B=960; A+B = 2126, bank loss
  - ◆ The DBMS's view of the second schedule:
  - ◆ T1: R(A), W(A),
  - ◆ T2: R(A), W(A), R(B), W(B)

# Interleaved vs. Parallel processing



# **Transactions and schedules**

# Transactions

- ◆ Transaction is a sequence of operations, an atomic unit of work
- ◆ Structure of a **committed** (successful) transaction:
  1. **begin** ← marks the beginning of transaction execution
  2. one or several **read(X), write(X)**, local operations (e.g., var updates), transaction control operations (e.g., locks), etc.
  3. **End** ← mark the end of transaction execution
  4. one or several operations checking consistency, serializability, etc.
  5. **commit** ← successful completion, changes cannot be undone
- ◆ Structure of an **aborted** (unsuccessful) transaction: same beginning, but ends at any point with
  - ★ N. **abort** ← unsuccessful completion, all changes must be undone
- ◆ A **partial** transaction:
  - ★ a beginning of one above ('waiting' for next operation)

# Schedules

- ◆ Schedule (or history, execution plan)  $S$  of transactions  $T_1, \dots, T_n$ : a (total) ordering of the operations of  $T_1, \dots, T_n$ 
  - ★ operations of different transactions can interleave
  - ★ operations of each  $T_i$  are in the same order as in  $T_i$

# Schedules

## ◆ Example:

★ Sa:  $r1(X)$ ,  $r2(X)$ ,  $w1(X)$ ,  $r1(Y)$ ,  $w2(X)$ ,  $w1(Y)$

| $T_1$  | $T_2$  |
|--|--|
| <p>Time ↓</p> <pre>read(X);<br/>X := X - N;<br/><br/>write(X);<br/>read(Y);<br/><br/>Y := Y + N;<br/>write(Y);</pre> | <pre>read(X);<br/>X := X + M;<br/><br/>write(X);</pre> |

## ◆ Example (complete):

★ Sb :  $r1(X)$ ,  $w1(X)$ ,  $r2(X)$ ,  $w2(X)$ ,  $r1(Y)$ ,  $a1$ ,  $c2$

# Complete Schedule

- ◆ all  $T_1, \dots, T_n$  are committed or aborted
  - ★ A transaction that successfully completes its execution will have a **commit** instructions as the last statement
    - By default, transaction assumed to execute commit instruction as its last step
  - ★ A transaction that fails to successfully complete its execution will have an **abort** instruction as the last statement

# Example

- ◆ Let
  - ★ T1 transfer \$50 from A to B,
- ◆ and
  - ★ T2 transfer 10% of the balance from A to B.

# Schedule 1

- ◆ A serial schedule in which T1 is followed by T2

| $T_1$  | $T_2$   |
|--|---|
| read ( $A$ )<br>$A := A - 50$<br>write ( $A$ )<br>read ( $B$ )<br>$B := B + 50$<br>write ( $B$ )<br>commit | read ( $A$ )<br>$temp := A * 0.1$<br>$A := A - temp$<br>write ( $A$ )<br>read ( $B$ )<br>$B := B + temp$<br>write ( $B$ )<br>commit |

# Schedule 2

- ◆ A serial schedule where T2 is followed by T1

| $T_1$   | $T_2$  |
|---|--|
| <pre>read (A) temp := A * 0.1 A := A - temp write (A)</pre> | <pre>read (B) B := B + temp write (B) commit</pre> |

# Schedule 3

- ◆ This schedule is not a serial schedule, but it is equivalent to Schedule 1
- ★ In Schedules 1, 2 and 3, the sum  $A + B$  is preserved.

| $T_1$  | $T_2$   |
|--|---|
| read ( $A$ )<br>$A := A - 50$<br>write ( $A$ )           | read ( $A$ )<br>$temp := A * 0.1$<br>$A := A - temp$<br>write ( $A$ ) |
| read ( $B$ )<br>$B := B + 50$<br>write ( $B$ )<br>commit | read ( $B$ )<br>$B := B + temp$<br>write ( $B$ )<br>commit            |

# Schedule 4

- ◆ The following concurrent schedule does not preserve the value of  $(A + B)$ .

| $T_1$   | $T_2$   |
|---|---|
| read ( $A$ )<br>$A := A - 50$   | read ( $A$ )<br>$temp := A * 0.1$<br>$A := A - temp$<br>write ( $A$ )<br>read ( $B$ ) |
| write ( $A$ )<br>read ( $B$ )<br>$B := B + 50$<br>write ( $B$ )<br>commit | $B := B + temp$<br>write ( $B$ )<br>commit  |

# **Serializability**

# Serializability

- ◆ Each transaction should preserve database **consistency**.
  - ★ ***Serial execution*** of a set of transactions preserves database consistency.
- ◆ A (possibly **concurrent**) schedule is **serializable** if it is **equivalent** to a serial schedule.

# Serializable schedules

Time ↓

| $T_1$  | $T_2$   |
|--|---|
| read( $X$ );<br>$X := X - N$ ;<br><br>write( $X$ );<br>read( $Y$ );<br><br>$Y := Y + N$ ;<br>write( $Y$ ); | read( $X$ );<br>$X := X + M$ ;<br><br>write( $X$ ); |

Schedule C

Time ↓

| $T_1$  | $T_2$   |
|--|---|
| read( $X$ );<br>$X := X - N$ ;<br>write( $X$ );<br><br>read( $Y$ );<br>$Y := Y + N$ ;<br>write( $Y$ ); | read( $X$ );<br>$X := X + M$ ;<br>write( $X$ ); |

Schedule D

◆ C not serializable

D serializable

# Conflicting Instructions

- ◆ Instructions  $l_i$  and  $l_j$  of transactions  $T_i$  and  $T_j$  respectively,
- ◆ **conflict**
  - ★ ***if and only if*** there exists some item  $Q$  accessed by both  $l_i$  and  $l_j$ ,
  - ★ and ***at least*** one of these instructions **wrote**  $Q$ .

# Conflicting Instructions

◆ **li = read(Q), lj = read(Q).** don't conflict.

◆ **li = read(Q), lj = write(Q).** They conflict

★ r1(X), r2(X), w1(X), r1(Y), **w2(X)**, w1(Y)

◆ **li = write(Q), lj = read(Q).** They conflict

★ r1(X), **w1(X)**, r2(X), w2(X), r1(Y), a1

◆ **li = write(Q), lj = write(Q).** They conflict

★ r1(X), **w1(X)**, r2(X), **w2(X)**, r1(Y), a1

# Conflicting Instructions

- ◆ A conflict between  $li$  and  $lj$  forces a (logical) temporal order between them.
- ◆ If  $li$  and  $lj$  are consecutive in a schedule and they do not conflict,
  - ★ their results would remain the same even if they had been interchanged in the schedule.

# **Serializability Types**

Different forms of schedule equivalence give the notions of:

**Conflict  
serializability**

**View  
serializability**

# Conflict Serializability

- ◆ If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, we say that  $S$  and  $S'$  are *conflict equivalent*.
- ◆ We say that a schedule  $S$  is **conflict serializable** if it is conflict equivalent to a serial schedule

# Conflict equivalent

- ◆ Two schedules are conflict equivalent if
  - ★ they are schedules of the same transactions
  - ★ if the **relative order** of every conflict (read-write, write-write) is the **same** in both schedules

# Conflict Serializability (Cont.)

- ◆ Schedule 3 can be transformed into Schedule 6, a serial schedule where  $T_2$  follows  $T_1$ , by series of swaps of non-conflicting instructions.

★ Therefore **Schedule 3** is conflict serializable.

| $T_1$                         | $T_2$                         |
|-------------------------------|-------------------------------|
| read ( $A$ )<br>write ( $A$ ) | read ( $A$ )<br>write ( $A$ ) |
| read ( $B$ )<br>write ( $B$ ) | read ( $B$ )<br>write ( $B$ ) |



| $T_1$  | $T_2$  |
|--|--|
| read ( $A$ )<br>write ( $A$ )<br>read ( $B$ )<br>write ( $B$ ) | read ( $A$ )<br>write ( $A$ )<br>read ( $B$ )<br>write ( $B$ ) |
|  | read ( $A$ )<br>write ( $A$ )<br>read ( $B$ )<br>write ( $B$ ) |

**Schedule 3**

**Schedule 6**


# Conflict Serializability (Cont.)

- ◆ Example of a schedule that is not conflict serializable:



- ◆ We are unable to swap instructions in the above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$ .

# Conflict Serializability (Cont.)



Schedule  $C$  is **not (conflict-)serializable**

Schedule  $D$  is **Serializable** (equivalent to  $T_1; T_2$ )

# Testing for Serializability


- ◆ Consider some schedule of a set of transactions

$T_1, T_2, \dots, T_n$

- ◆ **Precedence graph**




★ a direct graph where the vertices are the transactions.

- ◆ We draw an arc from  $T_i$  to  $T_j$  if the two transactions conflict, and  $T_i$  accessed the data item on which the conflict arose earlier.
- ◆ We may label the arc by the item that was accessed.



# Test for Conflict Serializability

- ◆ A schedule is conflict serializable if and only if its precedence graph is ***acyclic***.
- ◆ If precedence graph is acyclic, the serializability order can be obtained by a ***topological sorting*** of the graph.
  - ★ This is a linear order consistent with the partial order of the graph.
  - ★ For example, a serializability order for Schedule A would be  
 $T_5 \rightarrow T_1 \rightarrow T_3 \rightarrow T_2 \rightarrow T_4$ 
    - Are there others?



# conflict serializable

Time ↓



| $T_1$  | $T_2$   |
|--|---|
| read( $X$ );<br>$X := X - N$ ;<br>write( $X$ );<br>read( $Y$ );<br>$Y := Y + N$ ;<br>write( $Y$ ); | read( $X$ );<br>$X := X + M$ ;<br>write( $X$ ); |

Schedule A

Time ↓

| $T_1$ | $T_2$  |
|-------|--|
|       | read( $X$ );<br>$X := X + M$ ;<br>write( $X$ );<br>read( $Y$ );<br>$Y := Y + N$ ;<br>write( $Y$ ); |

Schedule B



# Example 1


| $T_1$                          | $T_2$  |
|--------------------------------|--|
| read( $X$ );<br>$X := X - N;$  |  |
| write( $X$ );<br>read( $Y$ );  | read( $X$ );<br>$X := X + M;$<br>write( $X$ ); |
| $Y := Y + N;$<br>write( $Y$ ); |  |

Schedule C


| $T_1$  | $T_2$  |
|--|--|
| read( $X$ );<br>$X := X - N;$<br>write( $X$ ); |  |
|  | read( $X$ );<br>$X := X + M;$<br>write( $X$ ); |
|  | read( $Y$ );<br>$Y := Y + N;$<br>write( $Y$ ); |

Schedule D



# Example 2

| Transaction $T_1$  | Transaction $T_2$   | Transaction $T_3$   |
|--|---|---|
| Time<br>↓<br><pre>read(X);<br/>write(X);<br/><br/>read(Y);<br/>write(Y);</pre> | <pre>read(Z);<br/>read(Y);<br/>write(Y);<br/><br/>read(X);<br/><br/>write(X);</pre> | <pre>read(Y);<br/>read(Z);<br/><br/>write(Y);<br/>write(Z);</pre> |



Equivalent serial schedules

None


Reason

Cycle  $X(T_1 \rightarrow T_2), Y(T_2 \rightarrow T_1)$

Cycle  $X(T_1 \rightarrow T_2), YZ(T_2 \rightarrow T_3), Y(T_3 \rightarrow T_1)$

# Example 3

| Transaction $T_1$  | Transaction $T_2$  | Transaction $T_3$  |
|--|--|--|
| read( $X$ );<br>write( $X$ );<br><br>read( $Y$ );<br>write( $Y$ ); | read( $Z$ );<br><br>read( $Y$ );<br>write( $Y$ );<br>read( $X$ );<br>write( $X$ ); | read( $Y$ );<br>read( $Z$ );<br><br>write( $Y$ );<br>write( $Z$ ); |



Equivalent serial schedules

$T_3 \rightarrow T_1 \rightarrow T_2$

# View Serializability

- ◆ Let **S** and **S'** be two schedules with the same set of transactions.
- ◆ **S** and **S'** are **view equivalent** if the following three conditions are met, for each data item **Q**.
  1. If in **S**, transaction **Ti** **reads** the initial value of **Q**, then in **S'** transaction **Ti** must **read** the initial value of **Q**.
  2. If in **S** transaction **Ti** executes **read(Q)**, and that value was produced by transaction **Tj**, then in **S'** transaction **Ti** must **read** the value of **Q** that was produced by the same **write(Q)** of **Tj**.
  3. The transaction that performs the final **write(Q)** operation in **S** must also perform the final **write(Q)** operation in **S'**.

# View-serializability

- ◆ serializability based on view equivalence
- ◆ Each conflict-serializable is view-serializable
- ◆ Each view-serializable is semantic-serializable (so result-serializable)
- ◆ Difficult to check view-equivalence (NP-complete)

# View Serializability (Cont.)

- ◆ A schedule  $S$  is **view serializable** if it is **view-equivalent** to a serial schedule.
- ◆ Below is a schedule which is view-serializable but not conflict serializable.

| $T_{27}$      | $T_{28}$      | $T_{29}$      |
|---------------|---------------|---------------|
| read ( $Q$ )  | write ( $Q$ ) |               |
| write ( $Q$ ) |               | write ( $Q$ ) |

# Test for View Serializability

- ◆ The precedence graph test for **conflict serializability** cannot be used directly to test for **view serializability**.
  - ★ Extension to test for view serializability has cost exponential in the size of the precedence graph.
- ◆ The problem of checking if a schedule is view-serializable falls in the class of NP-complete problems.
  - ★ Thus, existence of an efficient algorithm is extremely unlikely.
- ◆ However practical algorithms that just check some **sufficient conditions** for view serializability can still be used.

# Serializability in concurrency control

- ◆ Every **serial** schedule is **serializable**, but not other way round
- ◆ Serializable schedules give benefit of **concurrent execution** without giving up any **correctness**
- ◆ Difficult to **test** for serializability in practice (even conflict-serializability)
  - system load, time of transaction submission, and process priority affect ordering of operations
  - often, we need to ensure serializability **before** the transactions complete
- ◆ DBMSs enforce concurrency control protocols that ensure serializability

# Recoverability

- ◆ For **durability**, we need to be able to **recover** schedules from transaction and system failures
- ◆ Schedules with respect to recoverability, are classified into:
  - ★ impossible to recover
  - ★ possible to recover
  - ★ recoverable
  - ★ easy to recover
  - ★ cascadeless, strict

# Recoverable Schedules

- ◆ Need to address the effect of transaction failures on concurrently running transactions.
- ◆ Recoverable schedule
  - ★ if a transaction  $T_j$  reads a data item previously written by a transaction  $T_i$ ,
  - ★ then the **commit** operation of  $T_i$  appears before the **commit** operation of  $T_j$ .

# Recoverable Schedules

- ◆ The following schedule is **not recoverable**
- ◆ If T8 should **abort**, T9 would have **read** (and possibly shown to the user) an **inconsistent** database state.

| $T_8$         | $T_9$                  |
|---------------|------------------------|
| read ( $A$ )  |                        |
| write ( $A$ ) |                        |
| read ( $B$ )  | read ( $A$ )<br>commit |

# Examples

- ◆  $r1(X), r2(X), w1(X), r1(Y), w2(X), c2, w1(Y), c1$ 
  - ★ recoverable
- ◆  $r1(X), w1(X), r2(X), r1(Y), w2(X), w1(Y), c2, a1$ 
  - ★ non-recoverable
- ◆  $r1(X), w1(X), r2(X), r1(Y), w2(X), w1(Y), a1, c2$ 
  - ★ still non-recoverable
- ◆  $r1(X), w1(X), r2(X), r1(Y), w2(X), w1(Y), a1, a2$ 
  - ★ recoverable

# Cascading Rollbacks

- ◆ A single transaction failure leads to a series of transaction rollbacks.
  - ★ Can lead to the undoing of a significant amount of work
- ◆ Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable)
  - ★ If T<sub>10</sub> fails, T<sub>11</sub> and T<sub>12</sub> must also be rolled back.

| $T_{10}$  | $T_{11}$  | $T_{12}$ |
|-----------|-----------|----------|
| read (A)  |           |          |
| read (B)  |           |          |
| write (A) |           |          |
|           | read (A)  |          |
|           | write (A) |          |
| abort     |           | read (A) |

# Cascadeless schedules

- ◆ Recoverable schedules may need cascading aborts (cascading rollbacks) to recover
  - ★ an (non-committed) transaction has to abort since it has read from another failed transaction
  - ★ Example
  - ★ **r1(X), w1(X), r2(X), r1(Y), w2(X), w1(Y), a1, a2**
- ◆ This may affect many transactions, and rollbacks may be expensive

***Cascadeless schedule idea: no cascade rollbacks***

# Cascadeless schedules

- ◆ Every transaction in the schedule reads only from **committed** transactions
  - ★ (i.e., for each  $r(X)$ ,  $X$  can be previously written by the same transaction, written by another **committed** transaction or have not been written before)
- ◆ Example:
  - ★  $r1(X), w1(X), r1(Y), r2(Z), w1(Y), c1, w2(Z), r2(X), w2(X), c2$

# Strict schedules

- ◆ Cascadeless schedules may still have fairly complicated recovery protocols, because we cannot just restore the values of all writes of an aborted transaction
  - ★ Example:  $w1(X)$ ,  $w2(X)$ ,  $a1$
  - ★ This may all affect many transactions (no cascades though)
- ◆ Strict schedule: no transaction can read or write an item  $X$  until the previous write of  $X$  is committed or aborted
- ◆ Examples:
  - ★  $w1(X)$ ,  $c1$ ,  $w2(X)$
  - ★  $w1(X)$ ,  $a1$ ,  $w2(X)$

# Recovery hierarchy

- ◆ Every **strict** schedule is **cascadeless**
- ◆ Every **cascadeless** is **recoverable**



THANK  
YOU ☺

