

CSAI 302

Advanced Database Systems

Lec 08

Concurrency Control

Introduction


- ◆ A database must provide a mechanism that will ensure that all possible schedules are:
 - ★ either **conflict** or **view serializable**, and
 - ★ are **recoverable** and preferably **cascadeless**
- ◆ A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency
- ◆ Testing a schedule for serializability after it has executed is a little too late!

Concurrency control problems

- ◆ Dirty write (or lost update)
- ◆ Dirty read (or temporary update)
- ◆ Non-repeatable read
- ◆ Incorrect summary (or phantom phenomena)


Dirty write (lost update)

- ◆ Two transactions update (read and write) the same item, second update starts before the first is complete
 - ★ (updates are interleaved)
- ◆ **Result:** incorrect value




Dirty read (temporary update)

- ◆ A transaction **updates** an item,
- ◆ new value is **used** by another transaction,
- ◆ the first transaction **fails** and its update is **rolled back**
- ◆ **Result:** the second transaction relies on an incorrect value



Non-repeatable read

- ◆ One transaction reads the same item **twice**
- ◆ another transaction **changes** its value in between
- ◆ **Result:** the first transaction gets different values of the same item



Incorrect summary (Phantom phenomena)

- ◆ One transaction calculates an **aggregate** summary
- ◆ while other transactions **update** some involved items
- ◆ **Result:** the aggregate is inconsistent

T_1	T_3
<pre>read(X); X := X - N; write(X);</pre>	<pre>sum := 0; read(A); sum := sum + A; ⋮</pre>
<pre>read(Y); Y := Y + N; write(Y);</pre>	<pre>read(X); sum := sum + X; read(Y); sum := sum + Y;</pre>

T_3 reads X after N is subtracted and reads Y before N is added; a wrong summary is the result (off by N).

why recovery is needed?

- ◆ (if something goes **wrong**, effects **roll back**)
- ◆ Should be either
 - ★ **Committed:**
 - evaluated in full and all effects permanently recorded
 - ★ **Aborted:**
 - no effect on the database (everything done rolled back)

The System Log

- ◆ System log keeps track of transaction operations
 - ★ Sequential, append-only file
 - ★ Not affected by failure
- ◆ Log buffer
 - ★ Main memory buffer
 - ★ When full, appended to end of log file on disk
- ◆ Log file is backed up periodically
- ◆ Undo and redo operations based on log possible

Commit Point of a Transaction

- ◆ Occurs when all operations that access the database have **completed successfully**
- ◆ Transaction writes a commit record into the log
 - ★ If system failure occurs, can search for transactions with recorded **start_transaction** but **no commit** record

Concurrency Control

Concurrency Control (Cont.)

- ◆ Goal
 - ★ Develop concurrency control protocols that will assure serializability.
- ◆ Concurrency-control schemes **tradeoff** between the ***amount of concurrency*** they allow and the ***amount of overhead*** that they incur.
 - ★ Some schemes allow only **conflict-serializable** schedules to be generated, while others allow **view-serializable**.

Concurrency Control vs. Serializability Tests

- ◆ Concurrency-control protocols
 - ★ allow concurrent schedules,
 - ★ ensure that the schedules are conflict/view serializable,
 - ★ and are recoverable and cascadeless.
- ◆ Concurrency control protocols ***do not examine*** the precedence graph as it is being created
 - ★ Instead, a protocol imposes a discipline that avoids non-serializable schedules.
- ◆ **Tests for serializability help us understand why a concurrency control protocol is correct.**

Weak Levels of Consistency

- ◆ Some applications are willing to live with weak levels of consistency, allowing schedules that are not serializable
 - ★ E.g., a read-only transaction that wants to get an approximate total balance of all accounts
 - ★ E.g., database statistics computed for query optimization can be approximate.
 - ★ Such transactions need not be serializable with respect to other transactions
- ◆ Tradeoff accuracy for performance

Isolation levels

- ◆ It is not always necessary to ensure no problems of all types
 - ★ in some applications we can tolerate some problems
 - ★ for the sake of more concurrency (i.e., efficiency)
- ◆ Simple isolation level hierarchy of schedules (and transaction management protocols):
 - ★ **Level 0**: no dirty reads
 - ★ **Level 1**: no dirty writes
 - ★ **Level 2**: no dirty reads, no dirty writes
 - ★ **Level 3**: (true isolation): no dirty reads, no dirty writes, no non-repeatable reads
- ◆ *A protocol that ensures recoverable and serializable schedules is Level 3*

Implementation of Isolation Levels

◆ Locking

- ★ Lock on whole database vs lock on items
- ★ How long to hold lock?
- ★ Shared vs exclusive locks

◆ Timestamps

- ★ Transaction timestamp assigned e.g. when a transaction begins
- ★ Data items store two timestamps
 - Read timestamp
 - Write timestamp

- ★ Timestamps are used to detect out of order accesses

◆ Multiple versions of each data item

- ★ Allow transactions to read from a “snapshot” of the database



Locking

Locks

- ◆ **Lock:** a variable associated with a data item
 - ★ describes status for operations that can be applied
 - ★ unique **Lock(X)** for each data item **X**
- ◆ binary locks:
 - ★ **Lock(X)** takes values **locked** or **unlocked** (or 1 and 0)
- ◆ shared/exclusive locks:
 - ★ **Lock(X)** takes values **read-locked**, **write-locked**, or **unlocked**
- ◆ These variables are managed by concurrency control subsystem

Binary locks

- ◆ for each item X , variable $\text{Lock}(X)$ can have one of the two values:
 - ★ $\text{Lock}(X)$ is locked (or 1) means that X is locked by a transaction and cannot be accessed by any other transaction
 - ★ $\text{Lock}(X)$ is unlocked (or 0) means that X is available and can be accessed when requested
- ◆ schedules start with all unlocked

Binary locks in transactions

- ◆ Use two operations **lock(X)** and **unlock(X)** applying these rules:
 - ★ all **read(X)** and **write(X)** must be between the two
 - ★ well-formed: no unlocking without locking, etc.
- ◆ These operations are implemented such that
 - ★ when **lock(X)** is issued
 - if **Lock(X) = 1**, the transaction is forced to wait for **Lock(X) = 0** (in the schedule)
 - if **Lock(X) = 0**, it is set to 1 and the transaction can proceed
 - ★ when **unlock(X)** is issued, **Lock(X)** is set to 0

Ex1: Transactions with binary locks

T_1

```
lock(X);
read(X);
X := X - 1;
write(X);
unlock(X);
lock(Y);
read(Y);
Y := Y + 1;
write(Y);
unlock(Y);
```

T_2

```
lock(X);
lock(Y);
read(X);
unlock(X);
X := X + 1;
lock(X);
write(X);
unlock(Y);
unlock(X);
```

$T_1 : l1(X), r1(X), w1(X), u1(X), l1(Y), r1(Y), w1(Y), u1(Y)$: Ok!

$T_2 : l2(X), l2(Y), r2(X), u2(X), l2(X), w2(X), u2(Y), u2(Y)$: Ok!

Note: T_2 locks and unlocks X twice and locks Y without a need

Ex2: Transactions with binary locks

T_1

```
lock(X);
read(X)
X := X - 1;
write(X);
unlock(X);
read(Y);
lock(Y);
Y := Y + 1;
write(Y);
unlock(Y);
```

T_2

```
lock(X);
lock(Y);
read(X);
unlock(X);
X := X + 1;
lock(X);
write(X);

unlock(X);
```

$T1 : l1(X), r1(X), w1(X), u1(X), r1(Y), l1(Y), w1(Y), u1(Y)$: Not ok

$T2 : l2(X), l2(Y), r2(X), u2(X), l2(X), w2(X), u2(Y)$: Not ok

Note: $T1$ reads Y without locking, and $T2$ does not unlock Y

Ex: Schedules with binary locks

- ◆ $S_1 : l1(X), r1(X), u1(X), l2(X), r2(X), u2(X), l1(X), w1(X), u1(X)$
★ Ok!
- ◆ $S_2 : l1(X), r1(X), l2(X), u1(X), r2(X), u2(X), l1(X), w1(X), u1(X)$
★ Ok!
- ◆ $S_3 : l1(X), r1(X), l2(X), \text{r2}(X), u1(X), u2(X), l1(X), w1(X), u1(X)$
★ Not ok
 - (implementation of $l2(X)$ is 'broken': should not allow to continue T_2 with $r2(X)$ before $u1(X)$)
- ◆ Note: locks (i.e., the rules and implementations as above) do not guarantee serializability by themselves;

Lock table

- ◆ Concurrency control subsystem has a lock manager module that relies on a lock table:
 - ★ has schema [Data_item_name, Locking_transaction]
 - ★ keeps only locked items (others are unlocked)
- ◆ Lock manager keeps track of and controls access to locks by checking the rules for transactions and enforcing allowed schedules (waiting, etc.)
- ◆ Binary locking is too restrictive anyway

Shared/exclusive locks

- ◆ Allow sharing for several reading transactions (but exclusive writes)
- ◆ for each item X, **Lock(X)** is
 - ★ **read-locked** (or **share-locked**) means that X is read by a transaction
 - variable **#READS(X)** keeps the number of reading transactions
 - ★ **write-locked** (or exclusive-locked) means that X is exclusively locked for writing by some transaction and cannot be accessed by others
 - ★ **unlocked** means that X is available for access (with locking)
- ◆ schedules start with all **Lock(X) = unlocked** and **#READS(X) = 0**

Transaction with shared/exclusive locks

- ◆ use three operations, `read_lock(X)`, `write_lock(X)` and `unlock(X)` with rules:
 - ★ each `read(X)` is after `read_lock(X)` or `write_lock(X)` (with no `unlock(X)` in between)
 - ★ each `write(X)` is after `write_lock(X)` (with no `unlock(X)` in between)
 - ★ each `read_lock(X)` and `write_lock(X)` has `unlock(X)` at some point after
 - ★ no **changing** of lock type before unlocking;
 - that is, no `read_lock(X)` after `write_lock(X)` without `unlock(X)` between

Implementation of read_lock

read_lock(X):

B: if LKD(X) = “unlocked”
 then begin LKD(X) \leftarrow “read-locked”;
 #READS(X) \leftarrow 1
 end
 else if LKD(X) = “read-locked”
 then #READS(X) \leftarrow #READS(X) + 1
 else begin
 wait (until LKD(X) = “unlocked”
 and the lock manager wakes up the transaction);
 go to B
 end;

Implementation of write_lock

write_lock(X):

B: if LKD(X) = “unlocked”

 then LKD(X) \leftarrow “write-locked”

else begin

 wait (until LKD(X) = “unlocked”

 and the lock manager wakes up the transaction);

 go to B

end;

Implementation of unlock

unlock (X):

if LKD(X) = “write-locked”

then **begin** LKD(X) \leftarrow “unlocked”;

wakeup one of the waiting transactions, if any

end

else if LKD(X) = “read-locked”

then **begin**

#READS(X) \leftarrow #READS(X) -1;

if #READS (X) = 0

then **begin** LKD(X) = “unlocked”;

wakeup one of the waiting transactions, if any

end

end;

Examples

- ◆ $S1 : rl1(X), r1(X), u1(X), rl2(X), r2(X), u2(X), wl1(X), w1(X), u1(X)$
 - ★ Ok!
- ◆ $S1' : rl1(X), r1(X), u1(X), rl2(X), r2(X), u2(X), rl1(X), w1(X), u1(X)$
 - ★ Not ok (write with read lock)
- ◆ $S1'' : wl1(X), r1(X), u1(X), rl2(X), r2(X), u2(X), wl1(X), w1(X), u1(X)$
 - ★ Ok!
- ◆ $S3 : rl1(X), r1(X), rl2(X), r2(X), u1(X), u2(X), wl1(X), w1(X), u1(X)$
 - ★ Ok!
- ◆ $S3' : rl1(X), r1(X), wl2(X), w2(X), u1(X), u2(X), wl1(X), w1(X), u1(X)$
 - ★ Not ok (implementation of $wl2(X)$ is 'broken': should be exclusive)
- ◆ $S3'' : rl1(X), r1(X), rl2(X), r2(X), u2(X), wl1(X), w1(X), u1(X)$
 - ★ Not ok ($T1$ is not following the rules: a second lock without unlocking)

Two-phase locking

Two-phase locking

- ◆ A transaction (binary or shared/exclusive) follows **two-phase locking protocol** if:
 - ★ all **locking** operations are **before** all **unlocking** operations
- ◆ Two phases:
 - ★ locking (**first**, **expanding**, or **growing**) phase
 - ★ unlocking (**second** or **shrinking**) phase

Two-phase locking: negative example

T_1	T_2
read_lock(Y); read(Y); unlock(Y); write_lock(X); read(X); $X := X + Y$; write(X); unlock(X);	read_lock(X); read(X); unlock(X); write_lock(Y); read(Y); $Y := X + Y$; write(Y); unlock(Y);

T_1	T_2
read_lock(Y); read(Y); unlock(Y); write_lock(X); read(X); $X := X + Y$; write(X); unlock(X);	read_lock(X); read(X); unlock(X); write_lock(Y); read(Y); $Y := X + Y$; write(Y); unlock(Y);

Transactions T_1 and T_2 (left) do **not** follow the **two-phase locking protocol**, so they allow for a **non-serializable** schedule (right)

Two-phase locking: positive example

T_1'	T_2'
read_lock(Y); read(Y); write_lock(X); unlock(Y); read(X); $X := X + Y$; write(X); unlock(X);	read_lock(X); read(X); write_lock(Y); unlock(X); read(Y); $Y := X + Y$; write(Y); unlock(Y);

Time ↓

T_1'	T_2'
read_lock(Y); read(Y); write_lock(X); unlock(Y); read(X); $X := X + Y$; write(X); unlock(X); read_lock(X); read(X); write_lock(Y); unlock(X); read(Y); $Y := X + Y$; write(Y); unlock(Y);	


Versions T_1' and T_2' (left) follow the two-phase locking protocol, so no non-serializable schedule complies the rules

serializability by two-phase locking

- ◆ Two-phase locking guarantees serializability
 - ★ (i.e., every allowed schedule of transactions following the rules of the two-phase locking protocol is (conflict-serializable))
- ◆ As the result, we can create a serializable schedule operation by operation, without knowing the transactions in full

Drawback

- ◆ we may end up in a **deadlock**




Deadlocks

Deadlock problem

- Two (or more) transactions may be stuck in a deadlock, where each transaction is waiting for another transaction, making a 'waiting cycle'

Time ↓

T_1'	T_2'
read_lock(Y); read(Y);	read_lock(X); read(X);
write_lock(X);	write_lock(Y);



- This may happen for all types of locks,

Deadlock avoidance: Overview

- ◆ protocols that deal with deadlocks:
 - ★ (naive)
 - ★ deadlock prevention
 - ★ deadlock detection
 - ★ timeouts
- ◆ Most of protocols may roll back (i.e., abort and restart) some transactions
 - ★ **cascadeless** (i.e., **recoverable**) due to two-phase locking (all deadlocked transactions in the growing phase, so no one have read what they have written)
 - ★ may face and have to deal with **starvation**:
 - a transaction rolls back over and over again indefinitely

naive approaches

- ◆ conservative two-phase locking (not practical)
- ◆ OR
- ◆ locking according to a predefined order of all items
 - ★ all items are ordered in advance (e.g., by address)
 - ★ all locking of a transaction is in the increasing order
 - ★ no deadlocks by construction
 - ★ not practical due to
 - limited concurrency
 - unreasonable restrictions on transactions

Deadlock prevention protocols

- ◆ Timestamp-based (**wait-die** and **wound-wait**)
- ◆ Waiting-based (**no-waiting** and **cautious-waiting**)
 - ★ All four **avoid** deadlocks, but often **force** transaction aborts without a real deadlock
- ◆ Deadlock prevention
 - ★ is *not a popular* approach in practice due to **high overhead**,
 - ★ but may be useful when transactions are **long** and **use many items**

Timestamp-based deadlock prevention

- ◆ Transaction timestamp:
 - ★ a unique number $\text{TS}(T)$ assigned to each transaction T so that
 - ★ $\text{TS}(T') < \text{TS}(T)$ for each T' started before (for the first time)

Wait-die rule

- ◆ let transaction \mathbf{T} try to lock an item that is already locked by \mathbf{T}'
 - ★ if $\mathbf{TS}(\mathbf{T}) < \mathbf{TS}(\mathbf{T}')$ (i.e., \mathbf{T} is older than \mathbf{T}'), then \mathbf{T} waits for \mathbf{T}' to release
 - ★ otherwise (i.e., \mathbf{T} is younger) \mathbf{T} aborts (i.e., 'dies') and restarts with the same timestamp
- ◆ The older proceeds (maybe, after waiting) and the younger restarts, so the wait-die rule guarantees no **deadlocks**, no **starvation**

Wound-wait rule

- ◆ let transaction T try to lock an item that is already locked by T'
 - ★ if $TS(T) < TS(T')$ (i.e., T is older than T'), then T' aborts (i.e., is 'wounded') and restarts with the same timestamp
 - ★ otherwise (i.e., T is younger) T waits for T' to release
- ◆ The older proceeds and the younger restarts (maybe, after waiting), so the wound-wait rule guarantees no **deadlocks**, no **starvation**

Examples

- ◆ When T1 tries to lock X that is already locked by T2
- ◆ if T1 is older
 - ★ wait-die: T1 waits
 - ◆ b1, ..., b2, ..., l2(X), ..., l1(X), [T1 waits] ..., u2(X), ...
 - ★ wound-wait: T2 dies
 - ◆ b1, ..., b2, ..., l2(X), ..., l1(X), a2, ...
- ◆ if T1 is younger
 - ★ wait-die: T1 dies
 - ◆ b2, ..., b1, ..., l2(X), ..., l1(X), a1, ...
 - ★ wound-wait: T1 waits
 - ◆ b2, ..., b1, ..., l2(X), ..., l1(X), [T1 waits] ..., u2(X), ...

Waiting-based deadlock prevention

- ◆ **(No timestamps in these protocols)**
- ◆ Let transaction T try to lock an item that is already locked by T'
 - ★ no-waiting rule:
 - T aborts and restarts
 - ★ cautious-waiting rule:
 - if T' is not waiting for anyone, then T waits (regularly checked)
 - otherwise, T aborts and restarts
- ◆ Any of these rules guarantees **no deadlocks** (no cyclic waiting is possible)
- ◆ Special care should be taken to **avoid starvation**.
- ◆ **Easier**, but even **more unnecessary aborts**

Deadlock detection

- ◆ Deadlock happens when there is a cycle of transactions waiting for each other
- ◆ We can detect a cycle and break the cycle by aborting one transaction (the victim)
- ◆ Can be done by maintaining the wait-for graph:
 - ★ nodes: active transactions
 - ★ (directed) edges: transactions waiting for transactions
- ◆ A cycle in the wait-for graph means a deadlock
- ◆ May be resource-consuming to maintain
- ◆ Victim selection should be done cautiously to avoid starvation

Timeouts

- ◆ Timeout-based deadlock avoidance:
 - ★ fix a waiting timeout period in advance
 - ★ a transaction waits for the timeout period (*at most*) after this abort and restart
- ◆ Simple and easy to maintain
- ◆ No guarantees
 - ★ (unknown overhead, possible starvation)

Timestamps

Timestamp-based concurrency control

- ◆ We can use timestamps for concurrency control without locking:
 - ★ Run transactions until something goes wrong (i.e., a conflict is detected)
 - ★ Roll back one of the conflict transactions based on their timestamps
 - roll back is abort and restart with a new timestamp

Timestamp-based concurrency control

- ◆ Timestamp-based concurrency control protocols:

- ★ all ensure **serializability**
 - transactions arranged by their **timestamps**
 - ★ may not guarantee **recoverability**
 - ★ may not guarantee **no starvation**

- ◆ Several timestamp-based protocols

- ★ **basic**: ensures **serializability**, but **not recoverability**
 - ★ **strict**: ensures **serializability** and **strictness** (so **recoverability**)
 - ★ **basic/strict with Thomas write rule**: less roll-backs

Item timestamp variables


- ◆ Besides the timestamp of a **transaction**, the concurrency control subsystem maintains **two variables** for each item X :
 - ★ **read timestamp** $ReadTS(X)$:
 - the timestamp of the **youngest** transaction that **has read** X
 - ★ **write timestamp** $WriteTS(X)$:
 - the timestamp of the **youngest** transaction that **has written** X
- ◆ Maintained in appropriate **timestamp tables**
 - ★ (similar to the **lock table** in lock-based protocols)
 - ◆ (All **initialised** by 0)

Basic timestamp-ordering protocol

- ◆ If there is a **conflict** that violates the order (of the **serial schedule** imposed by the timestamps),
- ◆ then **roll back** the transaction of the **second** operation of the conflict

Read-write conflict 1


- ◆ a transaction may try to **read** data item X **too late**



- ◆ Should abort and restart **T1**

Read-write conflict 2


- ◆ a transaction may try to **write** data item X **too late**



- ◆ Should abort and restart **T1**

Write-write conflict

- ◆ a transaction may try to **write** data item X **too late**



- ◆ Should abort and restart **T1**

Basic timestamp-ordering Rules

◆ Rule 1. When transaction T issues $\text{read}(X)$:

- ★ if $\text{WriteTS}(X) > \text{TS}(T)$ (i.e., someone younger written X),
- ★ then roll-back T (i.e., abort and restart T with a new timestamp)
- ★ otherwise, execute the rest of $\text{read}(X)$ and set $\text{ReadTS}(X)$ to $\max(\text{ReadTS}(X), \text{TS}(T))$

◆ Rule 2. When transaction T issues $\text{write}(X)$:

- ★ if $\text{ReadTS}(X) > \text{TS}(T)$ or $\text{WriteTS}(X) > \text{TS}(T)$ (i.e., someone younger accessed X),
- ★ then roll-back T (i.e., abort and restart T with a new timestamp)
- ★ otherwise, execute the rest of $\text{write}(X)$ and set $\text{WriteTS}(X)$ to $\text{TS}(T)$

Basic timestamp-ordering properties

- ◆ Ensures serializable schedules, but they
 - ★ may be not recoverable,
 - (ensured by strict version)
 - ★ may have cascading roll-backs,
 - (avoided by strict version)
 - ★ may have starving transactions

Strict timestamp-ordering protocol

- ◆ Idea
- ★ In the rules the execution is **postponed** (using e.g., locking) until the previous write is committed (so ensuring **strict** schedules)

Rule 1


- ◆ When transaction T issues $\text{read}(X)$:
 - ◆ if $\text{WriteTS}(X) > \text{TS}(T)$ then roll-back T
 - ◆ otherwise
 - ★ wait until the transaction T' such that
 - ◆ $\text{WriteTS}(X) = \text{TS}(T')$ is committed (if there is such)
 - ★ execute the rest of $\text{read}(X)$ and
 - ◆ set $\text{ReadTS}(X)$ to $\max(\text{ReadTS}(X), \text{TS}(T))$

Rule 2

- ◆ When transaction T issues $\text{write}(X)$:
 - ◆ if $\text{ReadTS}(X) > TS(T)$ or $\text{WriteTS}(X) > TS(T)$ then roll-back T
 - ◆ otherwise
 - ★ wait until the transaction T' such that
 - ◆ $\text{WriteTS}(X) = TS(T')$ is committed (if there is such)
 - ★ execute the rest of $\text{write}(X)$ and set $\text{WriteTS}(X)$ to $TS(T')$

Thomas writing rule

- ◆ Improvement (less roll-backs) for both basic and strict protocols in case of write-write conflict




- ◆ Rule 1. When transaction T issues $\text{read}(X)$: ... (same)
- ◆ Rule 2. When transaction T issues $\text{write}(X)$:
 - ★ if $\text{ReadTS}(X) > TS(T)$ then roll-back T
 - ★ if $\text{WriteTS}(X) > TS(T)$ ignore the operation and proceed
 - ★ otherwise ... (same, i.e., possibly wait, execute, update)
- ◆ Same guarantees as the versions without the improvement

Multi-versioning

Multiversing concurrency control

- ◆ Multi-versioning idea:
 - ★ several versions of the (value of the) same item are kept by the system
 - ★ transactions are the same as usual (in particular, they request reads/writes for items, not versions)
 - ★ the versions of items are managed by the concurrency control subsystem
- ◆ Pros: increase concurrency
- ◆ Cons: more storage is needed to maintain multiple versions



Concurrency with Multiple granularity

Data item granularity

- ◆ Granularity: the size of data items
 - ★ field value of a record (fine)
 - ★ record
 - ★ disk block
 - ★ table file
 - ★ whole database (coarse)
- ◆ The tradeoff:
 - ★ more coarse allows for lower degree of concurrency
 - ★ more fine leads to more overhead (lock management, etc.)
- ◆ Best item size depends on transaction type:
 - ★ if a transaction accesses few records, a record as item is good
 - ★ if a transaction accesses many records, a block as item is good
- ◆ It may be nice to have different granularity for different transactions



THANK
YOU ☺

